

Assessing the statistical significance of association rules

Wilhelmiina Hämäläinen

Abstract An association rule is statistically significant, if it has a small probability to occur by chance. It is well-known that the traditional frequency-confidence framework does not produce statistically significant rules. It can both accept spurious rules (type 1 error) and reject significant rules (type 2 error). The same problem concerns other commonly used interestingness measures and pruning heuristics.

In this paper, we inspect the most common measure functions – frequency, confidence, degree of dependence, χ^2 , correlation coefficient, and J -measure – and redundancy reduction techniques. For each technique, we analyze whether it can make type 1 or type 2 error and the conditions under which the error occurs. In addition, we give new theoretical results which can be used to guide the search for statistically significant association rules.

Keywords Association rule · Statistical significance · Interestingness measure

1 Introduction

One of the most important tasks of data mining is the search for partial dependencies in data. A partial dependency between attributes A_1, \dots, A_l means that some values a_1, \dots, a_l occur together more often than expected, if the attributes were independent. When two attribute sets X and Y are partially dependent, the dependency can be expressed as a rule $X = \bar{x} \rightarrow Y = \bar{y}$, for the given value combinations \bar{x} and \bar{y} . If the rule is common and strong enough, it is called an *association rule* [2].

The commonness and strength of rule $X = \bar{x} \rightarrow Y = \bar{y}$ are measured by frequency $fr(X = \bar{x} \rightarrow Y = \bar{y}) = P(X = \bar{x}, Y = \bar{y})$ and confidence $cf(X = \bar{x} \rightarrow Y = \bar{y}) = P(Y = \bar{y} | X = \bar{x})$. It is required that $fr(X = \bar{x} \rightarrow Y = \bar{y}) \geq min_{fr}$ and $cf(X = \bar{x} \rightarrow Y = \bar{y}) \geq min_{cf}$ for some user-defined thresholds min_{fr} and min_{cf} .

The problem of traditional association rules is that they do not reflect the statistical significance of partial dependencies. Statistically, the significance of an association rule is defined by the probability that it has occurred by chance. In practice, the statistical significance depends on two things: frequency, $P(X = \bar{x}, Y = \bar{y})$, and the degree of dependence, $\frac{P(X = \bar{x}, Y = \bar{y})}{P(X = \bar{x})P(Y = \bar{y})}$. The stronger the dependence, the smaller the frequency can be, and vice versa.

This means that no absolute values can be given for the minimum frequency and minimum confidence.

This problem is well-known, and especially Webb [29,30] has criticized the frequency-confidence framework. He has shown that in the worst case, all discovered rules are spurious (statistically insignificant). Berzal et al. [4] and Aggarwal and Yu [1] have shown that the frequency-confidence framework is problematic, even if the absolute threshold values are not used.

Still the mainstream has ignored the problem, because the efficiency of the search algorithms lies on the frequency-based pruning. Nearly all search algorithms utilize the anti-monotonicity of the frequency: if set X is not frequent (given some \min_{fr}), then none of its supersets $Y \supset X$ can be frequent.

If the minimum frequency is set too high, several significant rules can be missed. On the other hand, if the minimum frequency is too low, a large number of spurious rules is accepted and the problem becomes computationally intractable. In statistics, these two error types – accepting spurious patterns and rejecting true patterns – are known as *type 1* and *type 2 errors*.

As a solution, statistical or other objective measures have been used to rank the discovered rules or to guide the search (e.g. [14, 18, 19]). These measures have generally two problems: either they are designed to assess dependencies between attributes (not attribute values) and can miss significant association rules, or they are statistically unsound.

In this paper, we examine how well the common interestingness measures and search heuristics capture significant association rules. For each method, we analyze whether it can make type 1 or type 2 error and the conditions under the errors they occur. We give several important results which can be used to select the most suitable search heuristics for the given mining task. On the other hand, the new theoretical results can give an insight for developing new, statistically sound search methods for partial dependencies.

The paper is structured as follows: In Section 2 the problem of finding statistically significant association rules is formalized. The basic definitions of full and partial dependencies, association rules, and statistical significance are given.

The effect of commonly used interest measures and search heuristics to the statistical significance is analyzed in Sections 3–5. In Section 3, we analyze how well the basic measures of association rules, frequency, confidence, and the degree of dependency, indicate the statistical significance. In Section 4, we proceed into more complex measures: χ^2 , correlation coefficient ϕ , and J -measure. The effect of common redundancy reduction techniques is analyzed in Section 5.

The final conclusions are drawn in Section 6.

2 Definitions

We begin by formalizing the problem and give exact definitions for full and partial dependencies, association rules, and statistical significance. The basic notations are defined in Table 1. When it is clear from the context, we use abbreviations A and $\neg A$ for single attribute values ($A = 1$) and ($A = 0$), and X or A_1, \dots, A_l for assignment $A_1 = 1, \dots, A_l = 1$.

Table 1 Basic notations.

Notation	Meaning
$A, B, C, A_1, A_2, A_3, \dots$	binary attributes (variables)
$a, b, c, a_i, a_2, a_3, \dots \in \{0, 1\}$	attribute values
$R = \{A_1, \dots, A_k\}$	set of all attributes (relational schema)
$ R = k$	number of attributes in R
$Dom(R) = \{0, 1\}^k$	attribute space (domain of R)
$X, Y, Z \subseteq R$	attribute sets
$ X = A_1, \dots, A_l = l$	number of attributes in set X
$Dom(X) = \{0, 1\}^l \subseteq Dom(R)$	domain of X , $ X = l$
$(X = \bar{x}) = \{(A_1 = a_1), \dots, (A_l = a_l)\}$	event; attribute value assignment for X , $ X = l$
$t = \{A_1 = t(A_1), \dots, A_k = t(A_k)\}$	row (tuple) according to R
$r = \{t_1, \dots, t_n \mid t_i \in Dom(R)\}$	relation according to R
$ r = n$	size of relation r (the number of rows)
$\sigma_{X=\bar{x}}(r) = \{t \in r \mid t[X] = \bar{x}\}$	set of rows for which $X = \bar{x}$ holds
$m(X = \bar{x}) = \sigma_{X=\bar{x}}(r) $	number of rows, for which $X = \bar{x}$ holds; ($X = \bar{x}$)'s absolute frequency or support
$P(X = \bar{x})$	($X = \bar{x}$)'s relative frequency (probability) in r
$P(Y = \bar{y} \mid X = \bar{x}) = \frac{P(X = \bar{x}, Y = \bar{y})}{P(X = \bar{x})}$	conditional probability of Y given X in r

2.1 Statistical dependence

Statistical dependence is classically defined through statistical independence (see e.g. [23, 16]). In the following, we will concentrate on *two-way dependencies*, i.e. dependencies between two attribute sets or events.

Definition 1 (Statistical independence and dependence) Let $X \subsetneq R$ and $Y \subseteq R \setminus X$ be sets of binary attributes.

Events $X = \bar{x}$ and $Y = \bar{y}$, $\bar{x} \in Dom(X)$, $\bar{y} \in Dom(Y)$, are *mutually independent*, if $P(X = \bar{x}, Y = \bar{y}) = P(X = \bar{x})P(Y = \bar{y})$.

Attribute sets X and Y are *mutually independent*, if $P(X = \bar{x}, Y = \bar{y}) = P(X = \bar{x})P(Y = \bar{y})$ for all value combinations $\bar{x} \in Dom(X)$ and $\bar{y} \in Dom(Y)$.

If the events or attribute sets are not independent, they are *dependent*.

The following example demonstrates that attribute sets can be dependent, even if some events are independent:

Example 1 Let $R = \{A, B, C\}$ be a set of binary attributes, where attribute C depends on attribute set $\{A, B\}$. Still it is possible that events $(C = 1)$ and $(A = 1, B = 1)$ are mutually independent. Table 2 gives an example of such a probability assignment.

However, it is also possible that all events are dependent. An example of such a probability assignment is given in Table 3.

When we analyze the distributions further, we observe that in Table 2, C is actually dependent on A and B separately: $P(A, C) = P(A)P(C) + d$, $P(B, C) = P(B)P(C) - d$. In Table 3, $\{A, B\}$ is the minimal set which has a dependency with C .

It is usually required that the dependency should be significant, before events or attribute sets are called dependent. In the latter case, this means that all value combinations ($X =$

Table 2 A probability assignment, where attribute C depends on set $\{A, B\}$, but event (A, B, C) is independent. $0 < d \leq \min\{P(A, \neg B)P(\neg C), (1 - P(A, \neg B))P(C)\}$.

X	$P(X)$
ABC	$P(A, B)P(C)$
$AB\neg C$	$P(A, B)P(\neg C)$
$A\neg BC$	$P(A, \neg B)P(C) + d$
$A\neg B\neg C$	$P(A, \neg B)P(\neg C) - d$
$\neg ABC$	$P(\neg A, B)P(C) - d$
$\neg AB\neg C$	$P(\neg A, B)P(\neg C) + d$
$\neg A\neg BC$	$P(\neg A, \neg B)P(C)$
$\neg A\neg B\neg C$	$P(\neg A, \neg B)P(\neg C)$

Table 3 A probability assignment, where attribute C depends on set $\{A, B\}$ and all events are dependent. $0 < d \leq \min\{P(A, \neg B)P(\neg C), (1 - P(A, \neg B))P(C)\}$.

X	$P(X)$
ABC	$P(A, B)P(C) - d$
$AB\neg C$	$P(A, B)P(\neg C) + d$
$A\neg BC$	$P(A, \neg B)P(C) + d$
$A\neg B\neg C$	$P(A, \neg B)P(\neg C) - d$
$\neg ABC$	$P(\neg A, B)P(C) - d$
$\neg AB\neg C$	$P(\neg A, B)P(\neg C) + d$
$\neg A\neg BC$	$P(\neg A, \neg B)P(C) + d$
$\neg A\neg B\neg C$	$P(\neg A, \neg B)P(\neg C) - d$

$\bar{x}, Y = \bar{y}$) should be represented in the data and the dependences should be sufficiently strong for most events (e.g. [7, 11]).

The strength of a statistical dependency between $(X = \bar{x})$ and $(Y = \bar{y})$ is defined by comparing $P(X = \bar{x}, Y = \bar{y})$ and $P(X = \bar{x})P(Y = \bar{y})$. The measure functions can be based on either the *absolute difference* (*dependence value* [16]), $d(X = \bar{x}, Y = \bar{y}) = P(X = \bar{x}, Y = \bar{y}) - P(X = \bar{x})P(Y = \bar{y})$, or the *relative difference*,

$$r(X = \bar{x}, Y = \bar{y}) = \frac{d(X = \bar{x}, Y = \bar{y})}{P(X = \bar{x})P(Y = \bar{y})}.$$

In the association rule literature, the relative difference is often defined via another measure, called the *degree of dependence* (*dependence* [31], *degree of independence* [32], or *interest* [7]):

$$\gamma(X = \bar{x}, Y = \bar{y}) = \frac{P(X = \bar{x}, Y = \bar{y})}{P(X = \bar{x})P(Y = \bar{y})} = 1 + \frac{d(X = \bar{x}, Y = \bar{y})}{P(X = \bar{x})P(Y = \bar{y})}. \quad (1)$$

In the real world data, it is quite common that some value combinations are overrepresented, while others are totally missing. In this situation, we cannot make any judgements concerning dependences between attribute sets, but still we can find significant dependencies between certain events. In this paper, these two kinds of significant dependencies are called partial and full dependencies:

Definition 2 (Partial and full dependence) Let X and Y be like before. Attribute sets X and Y are called *partially dependent*, if the dependency between events $(X = \bar{x})$ and $(Y = \bar{y})$ is significant for some $\bar{x} \in \text{Dom}(X)$ and $\bar{y} \in \text{Dom}(Y)$.

X and Y are called *fully dependent*, if the dependency between X and Y is significant.

Thus, full dependence implies partial dependence, but not vice versa. This means that the methods for assessing the significance of full dependencies do not necessarily capture all significant partial dependencies.

One trick is to turn a partial dependency into a full dependency by treating events $X = \bar{x}$ and $Y = \bar{y}$ as binary attributes. Table 4 gives a contingency table of the associated probabilities. Now it is more likely that all four value combinations are represented in the data and the methods for assessing full dependencies can be applied.

Table 4 A contingency table with probabilities of $P(X, Y)$, $P(X, \neg Y)$, $P(\neg X, Y)$ and $P(\neg X, \neg Y)$. If $d > 0$, $d \leq \min\{P(\neg X)P(Y), P(X)P(\neg Y)\}$, and if $d < 0$, $d \leq \min\{P(X)P(Y), P(\neg X)P(\neg Y)\}$.

	Y	$\neg Y$	Σ
X	$P(X, Y) = P(X)P(Y) + d$	$P(X, \neg Y) = P(X)P(Y) - d$	$P(X)$
$\neg X$	$P(\neg X, Y) = P(\neg X)P(Y) - d$	$P(\neg X, \neg Y) = P(\neg X)P(Y) + d$	$P(\neg X)$
Σ	$P(Y)$	$P(\neg Y)$	1

2.2 Association rules

Often, the dependency between events is expressed as rule $X = \bar{x} \rightarrow Y = \bar{y}$. Association rules [2] are a natural framework to express such rules. Traditionally, association rules are defined in the *frequency-confidence framework*:

Definition 3 (Association rule) Let R be a set of binary attributes and r a relation according to R . Let $X \subsetneq R$ and $Y \subseteq R \setminus X$, be attribute sets and $\bar{x} \in \text{Dom}(X)$ and $\bar{y} \in \text{Dom}(Y)$ their value combinations.

The *confidence* of rule $(X = \bar{x}) \rightarrow (Y = \bar{y})$ is

$$cf(X = \bar{x} \rightarrow Y = \bar{y}) = \frac{P(X = \bar{x}, Y = \bar{y})}{P(X = \bar{x})} = P(Y = \bar{y} | X = \bar{x})$$

and the *frequency* of the rule is

$$fr(X = \bar{x} \rightarrow Y = \bar{y}) = P(X = \bar{x}, Y = \bar{y}).$$

Given user-defined thresholds $\min_{cf}, \min_{fr} \in [0, 1]$, rule $(X = \bar{x}) \rightarrow (Y = \bar{y})$ is an *association rule* in r , if

- (i) $cf(X = \bar{x} \rightarrow Y = \bar{y}) \geq \min_{cf}$, and
- (ii) $fr(X = \bar{x} \rightarrow Y = \bar{y}) \geq \min_{fr}$.

The first condition requires that an association rule should be strong enough and the second condition requires that it should be common enough. In this paper, we call rules association rules, even if no thresholds \min_{fr} and \min_{cf} are specified.

Often it is assumed that the consequent $Y = \bar{y}$ contains just one attribute, $|Y| = 1$. When the consequent is a fixed class attribute C , rules $X = \bar{x} \rightarrow C = c$, $c \in \text{Dom}(C)$, are called *classification rules*.

Another common restriction is to allow only positive attribute values ($A_i = 1$). The reasons are mostly practical: in the traditional context of market-basket data, most of the items do not occur in a single basket. Thus, it is sensible to search only correlations between items that often occur together. On the other hand, the number of items is very large, typically > 1000 , and searching all association rules would be impossible. In the other contexts, negative attribute values cannot be excluded. For example, when we search dependencies in the demographic data, we cannot exclude all women, unmarried, employed, etc.

The main problem of the frequency-confidence framework is that the minimum frequency and confidence requirements do not guarantee any statistical dependence or significance [7, 1, 19]. However, most researchers have adopted Piatetsky-Shapiro's [21] argument that a rule cannot be interesting, if its antecedent and consequent are statistically independent. That is why it is often demanded that $\gamma(X = \bar{x} \rightarrow Y = \bar{y}) \neq 1$ (e.g. [7, 31, 26]). According to the sign of $\gamma - 1$, the rule or its type is called positive, negative or independent ("null association rule") [14, 9]. Usually, only positive dependencies are searched, since they can be used for prediction.

We note that from the statistical point of view, the direction of a rule ($X = \bar{x} \rightarrow Y = \bar{y}$ or $Y = \bar{y} \rightarrow X = \bar{x}$) is a matter of choice. In the worst case, the direction can be misleading, since rules are usually associated with causation and association rules (or correlations) do not necessarily represent any causality relationship [12].

Another important notice is that the association rules are not implications. Especially, rule $Y \rightarrow X$ is not the same as $\neg X \rightarrow \neg Y$. Unless $P(X) = P(Y) = 0.5$, rules $Y \rightarrow X$ and $\neg X \rightarrow \neg Y$ have different frequencies, confidences and degrees of dependence.

2.3 Statistical significance of partial dependencies

The idea of statistical significance tests is to estimate the probability of the observed or a rarer phenomenon, under some null hypothesis. When the objective is to test the significance of the dependency between $X = \bar{x}$ and $Y = \bar{y}$, the null hypothesis is the independence assumption: $P(X = \bar{x}, Y = \bar{y}) = P(X = \bar{x})P(Y = \bar{y})$. If the estimated probability p is very small, we can reject the independence assumption, and assume that the observed dependency is not due to chance, but significant at level p . The smaller p is, the more significant the observation is.

Usually the minimum requirement for any significance is $p \leq 0.05$. It means that there is 5% chance that a spurious rule passes the significance test ("type 1 error"). If we test 10 000 rules, it is likely that will find 500 spurious rules. This so called *multiple testing problem* is inherent in the knowledge discovery, where we often perform an exhaustive search over all possible patterns.

As a solution, the more patterns we test, the stricter bounds for the significance we should use. The most well-known method is *Bonferroni adjustment* [22], where the desired significance level p is divided by the number of tests. In the association rule discovery, we can give an upper bound for the number of rules to be tested. However, this rule is so strict that there is a risk that we do not recognize all significant patterns ("type 2 error"). Webb [30, 29] has argued that this is a less serious problem than finding spurious rules, because the number of rules is anyway large. He has also suggested another approach, where a part of the data is held as an evaluation set. Now the number of rules to be tested is known before testing, and higher significance levels can be used.

Let us now analyze the significance of partial dependency $X = \bar{x} \rightarrow Y = \bar{y}$. To simplify the notations, the sets are denoted by X and Y .

The significance of the observed frequency $m(X, Y)$ can be estimated exactly by the binomial distribution. Each row in relation r , $|r| = n$, corresponds to an independent Bernoulli trial, whose outcome is either 1 (XY occurs) or 0 (XY does not occur). All rows are mutually independent.

Assuming the independence of attributes X and Y , combination XY occurs on a row with probability $P(X)P(Y)$. Now the number of rows containing X, Y is a binomial random variable M with parameters $P(X)P(Y)$ and n . The mean of M is $\mu_M = nP(X)P(Y)$ and its variance is $\sigma_M^2 = nP(X)P(Y)(1 - P(X)P(Y))$. The probability that $M \geq m(X, Y)$ is

$$p = P(M \geq m(X, Y)) = \sum_{i=m(X, Y)}^n \binom{n}{i} (P(X)P(Y))^i (1 - P(X)P(Y))^{n-i}. \quad (2)$$

This can be approximated by the standard normal distribution

$$p \approx 1 - \Phi(t),$$

where $\Phi(t(X, Y)) = \frac{1}{\sqrt{2\pi}} \int_{-\infty}^{t(X, Y)} e^{-u^2/2} du$ is the standard normal cumulative distribution function and $t(X, Y)$ is standardized $m(X, Y)$:

$$t(X, Y) = \frac{m(X, Y) - \mu_M}{\sigma_M} = \frac{m(X, Y) - nP(X)P(Y)}{\sqrt{nP(X)P(Y)(1 - P(X)P(Y))}}. \quad (3)$$

The approximation is quite good for large n , but it should not be used, when the expected counts $nP(X)P(Y)$ and $n(1 - P(X)P(Y))$ are small. As a rule of thumb, it is often required that $nP(X)P(Y) > 5$ and $n(1 - P(X)P(Y)) > 5$ (e.g. [17, p. 121]).

The cumulative distribution function $\Phi(t)$ is quite difficult to calculate, but for the association rule mining it is enough to know $t(X, Y)$. Since $\Phi(t)$ is monotonically increasing, probability p is monotonically decreasing in the terms of $t(X, Y)$. Thus, we can use t as a measure function for ranking association rules according to their significance. On the other hand, we know that in the normal distribution $P(-2\sigma_M < M - \mu_M < 2\sigma_M) \approx 0.95$ or, equivalently,

$$P\left(-2 < \frac{M - \mu_M}{\sigma_M} < 2\right) \approx 0.95.$$

I.e. $P(t(X, Y) \geq 2) \approx 0.025$, which is a minimum requirement for any significance. Thus, we can prune all rules $X \rightarrow Y$ for which $t(X, Y) < 2$. Generally, we can set the threshold K according to Chebyshev's inequality (the proof is given e.g. in [17, pp. 780-781]):

$$P\left(-K < \frac{M - \mu_M}{\sigma_M} < K\right) \geq 1 - \frac{1}{K^2}.$$

I.e. $P(t \geq K) < \frac{1}{2K^2}$. Now the Bonferroni adjustment is achieved by using $\sqrt{m}K$ instead of K , where m is the number of tests.

Equations (2) and (3) can be directly generalized to attribute-value sets $\{A_1 = a_1, \dots, A_l = a_l\}$. Now the null hypothesis is that all attributes are mutually independent:

$$P(A_1 = a_1, \dots, A_l = a_l) = P(A_1 = a_1)P(A_2 = a_2) \dots P(A_l = a_l) = \prod_{i=1}^l P(A_i = a_i).$$

The significance of the dependence in set $\{A_1 = a_1, \dots, A_l = a_l\}$ is measured by

$$t(A_1 = a_1, \dots, A_l = a_l) = \frac{m(A_1 = a_1, \dots, A_l = a_l) - n \prod_{i=1}^l P(A_i = a_i)}{\sqrt{n \prod_{i=1}^l P(A_i = a_i) (1 - \prod_{i=1}^l P(A_i = a_i))}}.$$

3 Basic measures for association rules

The statistical significance of rule $X \rightarrow Y$ is a function of $P(X)$, $P(Y)$ and $P(X, Y)$ (Equation (3)). All basic measures, like frequency, confidence, and the degree of dependency, are composed from these elements. In the frequency-confidence framework, the assumption is that a high frequency ($P(X, Y)$) and a high confidence ($P(Y|X)$) indicate an interesting rule. In the following, we will analyze conditions under which this assumption fails. As an alternative, we analyze “frequency-dependence framework”, and show that a high frequency and a high degree of dependence, γ , indicate statistical significance.

3.1 Frequency and confidence

Figure 1 illustrates the significance of rule $X \rightarrow Y$ as a function of frequency $P(X, Y)$ and confidence $P(Y|X)$, when Y is fixed. The values of $P(Y)$ are 0.2, 0.4, 0.6 and 0.8. Now the significance measure t is expressed as

$$\hat{t} = \frac{\sqrt{P(X, Y)(P(Y|X) - P(Y))}}{\sqrt{P(Y)(P(Y|X) - P(X, Y)P(Y))}}.$$

Data size n is omitted, and the real significance is $t = \sqrt{n}\hat{t}$. The function is not defined when $P(Y|X) \leq P(X, Y)P(Y)$. For clarity, only areas where $t > 0$ are drawn. In addition, it holds that $P(X, Y) \leq P(Y)$.




Fig. 1 The significance of $X \rightarrow Y$ as a function of frequency $P(X, Y)$ and confidence $P(Y|X)$, when $P(Y) = 0.2$ (left top), $P(Y) = 0.4$ (right top), $P(Y) = 0.6$ (left bottom) and $P(Y) = 0.8$ (right bottom).

The contours are compatible with a common intuition that the significance is maximal, when both frequency and confidence are maximal. However, the significance of the rule depends on $P(Y)$. The higher $P(Y)$ is, the higher confidence the rule should have. The minimum requirement for the confidence is $P(Y|X) > P(Y)$, since otherwise $t \leq 0$. In the extreme case, when $P(Y) = 1$, the rule is totally insignificant ($P(Y|X) = 1$ for all X and $t = 0$). That is why rules with different consequents are not comparable, in the terms of their frequency and confidence. Often a rule with higher frequency and confidence may be less significant than a weaker and less frequent rule. A significant rule can be easily missed, when absolute min_{fr} and min_{cf} values are used.

Generally, the preference for high frequency and confidence can cause both type 1 and type 2 errors. Let us first analyze what kind of rules are accepted in the frequency-confidence framework. Let X and Y be like in Table 4. The frequency of rule $X = \bar{x} \rightarrow Y = \bar{y}$ is $P(X = \bar{x})P(Y = \bar{y}) + d$. Now any combination $X = \bar{x}, Y = \bar{y}$ can be frequent, if $P(X = \bar{x}, Y = \bar{y}) \geq min_{fr}$. If $P(X = \bar{x})P(Y = \bar{y}) \geq min_{fr}$, $X = \bar{x}$ and $Y = \bar{y}$ can be statistically independent ($d = 0$) or even negatively correlated ($d < 0$).

The confidence of rule $X = \bar{x} \rightarrow Y = \bar{y}$ is $P(Y = \bar{y}) + \frac{d}{P(X = \bar{x})}$. The highest confidence is achieved, when $P(Y = \bar{y})$ is large and $P(X = \bar{x})$ is small. If $P(Y = \bar{y}) \geq min_{cf}$, the rule is confident, even if $X = \bar{x}$ and $Y = \bar{y}$ are statistically independent.

On the other hand, the frequency-confidence framework can reject significant rules. Let us analyze what the minimum frequency and confidence should be for a rule to be significant.

Let $t(X = \bar{x} \rightarrow Y = \bar{y}) \geq K$. This holds, when the frequency is

$$P(X = \bar{x}, Y = \bar{y}) \geq P(X = \bar{x})P(Y = \bar{y}) + \frac{K\sqrt{P(X = \bar{x})P(Y = \bar{y})(1 - P(X = \bar{x})P(Y = \bar{y}))}}{n}$$

and the confidence is

$$P(Y = \bar{y}|X = \bar{x}) \geq P(Y = \bar{y}) + \frac{K\sqrt{P(Y = \bar{y})(1 - P(X = \bar{x})P(Y = \bar{y}))}}{nP(X = \bar{x})}.$$

We see that the larger n is, the smaller frequency and confidence suffice for significance. On the other hand, the larger significance level we require (expressed by K), the larger frequency and confidence should be. The problem is that both of them depend on $P(X = \bar{x})$ and $P(Y = \bar{y})$. The minimum frequency is an increasing function of $P(X = \bar{x})P(Y = \bar{y})$. The minimum confidence is obtained from the minimum frequency by dividing it by $P(X = \bar{x})$. Thus, the larger $P(X = \bar{x})$ is, the larger min_{cf} should be.

Example 2 Let $P(X) = P(Y) = 0.5$ and $n = 10000$. Now rule $X \rightarrow Y$ is significant, if $P(X, Y) = 0.25 + \frac{K\sqrt{3}}{400}$ and $P(Y|X) = 0.5 + \frac{K\sqrt{3}}{200}$. Especially the confidence is low and the rule is easily rejected with the typical min_{cf} settings. For example, if we require that $K = 10$ (indicating quite high significance), then confidence $0.5 + \frac{10\sqrt{3}}{200} < 0.60$ suffices.

3.2 Frequency and degree of dependence

The problems of the frequency-confidence framework could be easily corrected by using the degree of dependency, $\gamma(X = \bar{x}, Y = \bar{y})$, instead of confidence. This approach is adopted e.g.

in [25, 26, 1]. Since $\gamma(X = \bar{x}, Y = \bar{y}) = \frac{cf(X=\bar{x} \rightarrow Y=\bar{y})}{P(Y=\bar{y})}$, the frequency and degree of dependence alone determine the statistical significance t .

Figure 2 illustrates the significance of rule $X \rightarrow Y$ as a function of frequency $P(X, Y)$ and degree of dependence $\gamma = \gamma(X \rightarrow Y)$:

$$\hat{t}(X \rightarrow Y) = \frac{\sqrt{P(X, Y)(\gamma - 1)}}{\sqrt{\gamma - P(X, Y)}}.$$

Once again, the data size n is omitted and $t = \sqrt{n}\hat{t}$.




Fig. 2 The significance of $X \rightarrow Y$ as a function of frequency $P(X, Y)$ and degree of dependency γ (left) and the corresponding contours (right).

There is only one graph, because $P(Y)$ has no effect on the significance, when γ is given. $P(Y)$ determines only the maximal possible value for γ : $\gamma(X \rightarrow Y) \leq \frac{1}{P(Y)}$. (In the Figure, $\gamma \leq 2$ and $P(Y) \geq 0.5$.) The minimum value, $\gamma \geq 1$, is set, because we are interested in only positive correlations.

The function is defined when $P(X, Y)\gamma \leq 1$, because $\gamma \leq \frac{1}{\max\{P(X), P(Y)\}}$ and $P(X, Y)\gamma \leq P(X)\gamma \leq 1$.

From the contours we observe that t is nearly symmetric in the terms of $P(X, Y)$ and $(\gamma - 1)$. It means that the larger the frequency is, the smaller the degree of dependence can be, and vice versa. If rule R_1 has both higher frequency and higher degree of dependence than rule R_2 , it is more significant. If R_1 has only higher frequency, then the t -values of rules should be compared to decide the significance order.

The following theorem expresses the relationship between the frequency and the degree of dependence:

Theorem 1 When $t(X \rightarrow Y) = K$,

$$P(X, Y) = \frac{K^2 \gamma}{n(\gamma - 1)^2 + K^2}.$$

Proof By solving

$$t(X \rightarrow Y) = \frac{\sqrt{nP(X)(\gamma - 1)}}{\sqrt{\gamma - P(X)}} = K.$$

□

This result can be used for pruning areas in the search space, when an upper bound for γ is known. At least areas where $K < 2$ can be safely pruned, since $t \geq 2$ is a minimum requirement for any significance.

The simplest method to search all statistically significant rules is to search all frequent sets with sufficiently small \min_{fr} and then select from each frequent set the rules with sufficient t .

The following theorem gives a safe minimum frequency threshold for the whole data set. It guarantees that no significant rules are missed. For simplicity, we assume that $|Y| = 1$.

Theorem 2 *Let $p_{min} = \min\{P(A_i = a_i) \mid A_i \in R, a_i = \{0, 1\}\}$. Let $K \geq 2$ be the desired significance level. For all sets $X \subseteq R$ and any $A \in X$*

- (i) $\gamma(X \setminus A \rightarrow A) \leq \frac{1}{p_{min}}$ and
- (ii) $X \rightarrow A$ cannot be significant, unless

$$P(X) \geq \frac{K^2 p_{min}}{n(1 - p_{min})^2 + K^2 p_{min}^2}.$$

Proof By solving

$$t(X \setminus A \rightarrow A) = \frac{\sqrt{nP(X)}(\gamma - 1)}{\sqrt{\gamma - P(X)}} \geq K.$$

□

4 Measures for dependence detection

Next, we analyze the most common objective measures for dependence detection. We focus on the χ^2 -measure, which is the most common statistical method for assessing the significance of dependencies. It is often used in association rule mining, when the goal is to find statistically significant association rules. Pearson correlation coefficient ϕ is another statistical measure, which has been applied to association rule mining.

J -measure [24] is derived from the mutual information, which is an information-theoretic measure for assessing dependencies between attributes. It is especially designed for ranking decision rules, and often applied in the association rule discovery.

Empirical comparisons of these and other interestingness measures can be found in e.g. [28, 27, 26].

4.1 χ^2 -measure

4.1.1 Definition

The χ^2 -independence test is the most popular statistical test for detecting dependencies between attributes. The idea of the χ^2 test is to compare the observed frequencies $O(m(X))$ to the expected frequencies $E(m(X))$ by

$$\chi^2(X) = \sum_{\bar{x} \in \text{Dom}(X)} \frac{O((m(X = \bar{x})) - E(m(X = \bar{x})))^2}{E(m(X = \bar{x}))}.$$

When the test variable is approximately normally distributed, the test measure follows the χ^2 -distribution. Usually this assumption holds for large n . As a rule of thumb, it is suggested (e.g. [17, p. 630]) that all of the expected frequencies should be at least 5.

When we test a dependency between two attribute sets, X and Y , the contingency table contains only four cells (Table 4). Now the test metric is

$$\chi^2(X, Y) = \sum_{i=0}^1 \sum_{j=0}^1 \frac{(m(X=i, Y=j) - nP(X=i)P(Y=j))^2}{nP(X=i)P(Y=j)} = \frac{n(P(X=1, Y=1) - P(X=1)P(Y=1))^2}{P(X=1)P(X=0)P(Y=1)P(Y=0)}.$$

If $\chi^2(X, Y)$ is less than the critical χ^2 value at level p and 1 degree of freedom, X and Y are statistically independent with probability $1 - p$. Otherwise, the dependency is significant at level p .

The above equation can be generalized to measure dependencies between all variables in set $X = A_1, \dots, A_l$:

$$\chi^2(X) = \sum_{(a_1, \dots, a_l)} \frac{n(P(A_1 = a_1, \dots, A_l = a_l) - P(A_1 = a_1) \dots P(A_l = a_l))^2}{P(A_1 = a_1) \dots P(A_l = a_l)}.$$

4.1.2 Applying χ^2 in the association rule discovery

The simplest way to use χ^2 -measure in the association rule discovery is to generate rules from frequent sets based on their χ^2 -values. For each frequent set X , all rules of form $X \setminus Y \rightarrow Y$ with a sufficient χ^2 -value are selected (e.g. [10]).

This approach does not find all rules which are significant in the χ^2 sense. First, the rules are preselected according to their frequency. If the minimum frequency is set too high, some significant rules are missed.

Second, it is possible that a weak rule ($P(Y|X) \leq 0.5$) is selected, because its companion rules $X \rightarrow \neg Y$, $\neg X \rightarrow Y$, and/or $\neg X \rightarrow \neg Y$ are significant. The rule confidence can be used to check that $P(Y|X) > P(\neg Y|X)$, but it does not guarantee that $X \rightarrow Y$ is significant. As a solution, it is often required (e.g. [14]) that $P(X, Y) > P(X)P(Y)$. Unfortunately, it is still possible that the high χ^2 -value is due to $\neg X \rightarrow \neg Y$.

The first problem would be solved, if we could search the rules directly with the χ^2 -measure. Unfortunately, this is not feasible, since χ^2 -measure is not monotonic. For any rule $X \rightarrow Y$ and its generalization $Z \rightarrow Y$, $Z \subseteq X$, it is unknown, whether $\chi^2(X \rightarrow Y) > \chi^2(Z \rightarrow Y)$ or $\chi^2(X \rightarrow Y) \leq \chi^2(Z \rightarrow Y)$.

There are at least two solutions to this problem: First, χ^2 is used to find only the interesting attribute sets [7]. Second, the convexity of the χ^2 -measure can be utilized in searching optimal rules with a fixed consequent C [18, 19].

When χ^2 is calculated for attribute sets, it is upwards closed under set inclusion. This means that the χ^2 value can only increase, when attributes are added to a set. (Proof for the special case $|Z| = 2$, $|X| = 3$, $Z \subsetneq X$, is given in [7].)

Theorem 3 For all attribute sets X, Z , $Z \subsetneq X$, $\chi^2(Z) \leq \chi^2(X)$.

Proof Let $X = Z, A$, $|Z| = l$ and $|Z| = l + 1$. $\chi^2(Z)$ contains 2^l terms of form $\frac{n(P(A_1=a_1, \dots, A_l=a_l) - P(A_1=a_1) \dots P(A_l=a_l))^2}{P(A_1=a_1) \dots P(A_l=a_l)} = \frac{n(U-V)^2}{V^2}$. For each such term $\chi^2(X)$ contains two terms:

$$\begin{aligned} & \frac{n(P(A_1 = a_1, \dots, A_l = a_l, A_{l+1} = 1) - P(A_1 = a_1) \dots P(A_l = a_l)P(A_{l+1} = 1))^2}{P(A_1 = a_1) \dots P(A_l = a_l)P(A_{l+1} = 1)} + \\ & \frac{n(P(A_1 = a_1, \dots, A_l = a_l, A_{l+1} = 0) - P(A_1 = a_1) \dots P(A_l = a_l)P(A_{l+1} = 0))^2}{P(A_1 = a_1) \dots P(A_l = a_l)P(A_{l+1} = 0)} = \\ & \frac{n(UP(A|Z) - VP(A))^2}{P(A)} + \frac{n(UP(\neg A|Z) - VP(\neg A))^2}{P(\neg A)}. \end{aligned}$$

Now it is enough to show that

$$\begin{aligned} \frac{(U-V)^2}{V^2} & \leq \frac{(UP(A|Z) - VP(A))^2}{P(A)} + \frac{(UP(\neg A|Z) - VP(\neg A))^2}{P(\neg A)} \\ & \Leftrightarrow \frac{U^2 - 2UV + V^2}{V} \leq \frac{U^2(P(A|Z)^2 + P(A)P(\neg A|Z))}{VP(A)P(\neg A)} + \frac{-2UV + V^2}{V}. \end{aligned}$$

This is always true, because

$$U^2 \leq \frac{U^2(P(A|Z)^2 + p(A)P(\neg A|Z))}{P(A)P(\neg A)} \Leftrightarrow P(A|Z)^2 - P(A)P(A|Z) + P(A)^2 \geq 0.$$

□

Thus, the most significant sets are the most specific, containing all k attributes. That is why Brin et al. [7, 23] used χ^2 -test to find the "minimally correlated sets", i.e. the most general attribute sets X such that $\chi^2(X) \geq \min_{\chi^2}$ for some cutoff value \min_{χ^2} . The type of correlation in set $X = A_1, \dots, A_l$ was determined by the interest measure

$$\rho(A_1, \dots, A_l) = \frac{P(A_1, \dots, A_l)}{P(A_1) \dots P(A_l)}.$$

In addition, they used a new frequency measure for pruning:

$$fr(X) = \max \left\{ s \mid \left| \frac{\{(X = \bar{x}) \mid x \in Dom(X) \wedge P(X = \bar{x}) \geq s\}}{|Dom(X)|} \right| \geq p \right\},$$

where $p \in [0, 1]$. This measure demands that in the contingency table of $|Dom(X)|$ cells the frequency must be sufficient in at least $p|Dom(X)|$ cells. High p and \min_{fr} values produce effective pruning, but several significant attribute sets can be missed, if their absolute frequency is too low or the dependency is only a partial dependency. For example, this heuristic ignores a dependency if $P(X = \bar{x})$ is high for some $\bar{x} \in Dom(X)$, but $P(X = \bar{x}_1) = P(X = \bar{x}_2) = \dots = P(X = \bar{x}_l)$ for all $\bar{x}_i \neq \bar{x}$. In addition, we note that parameters p and \min_{fr} determine how many attributes X can contain, since $|X| \leq -\log(p \cdot \min_{fr})$. For example, if $p = 0.25$ and $\min_{fr} = 0.01$ (as suggested in [7]), $|X| \leq 8$.

Brin et al. did not generate any rules, even if the attribute sets were called "correlation rules" [7] or "dependence rules" [23]. A natural question is whether we could generate significant rules from the correlated sets. Unfortunately, the dependence in a set is only a sufficient and not necessary condition for two-way dependencies. In addition, it is possible

that none of the association rules generated from a correlated set is necessarily significant [18].

The second approach, introduced by Morishita et al. [18, 19], is to utilize the convexity of the χ^2 function, when the consequent C is fixed. The idea is to prune a branch containing rule $Z \rightarrow C$ and all its specialization rules $X \rightarrow C$, $Z \subseteq X$, if $\max\{\chi^2(X \rightarrow C)\} < \min_{\chi^2}$ for the given cutoff value \min_{χ^2} . Because χ^2 is convex, $\max\{\chi^2(X \rightarrow C)\} < \min_{\chi^2}$ can be bounded by equation

$$\chi^2(X \rightarrow C) \leq \max \left\{ \frac{nP(Z, C)P(\neg C)}{(1 - P(Z, C))P(C)}, \frac{nP(Z, \neg C)P(C)}{(1 - P(Z, \neg C))P(\neg C)} \right\}.$$

Now the frequency-based pruning is not necessary and it is possible to find all rules with a sufficient χ^2 -value or the best rules in the χ^2 sense. This approach works correctly, when the goal is to find full dependencies. Partial dependencies with fixed C could be searched similarly by applying the properties of the t -measure.

4.1.3 Analysis

The main problem of the χ^2 -independence test is that it designed to measure dependencies between attributes. That is why it can fail to detect significant partial dependencies. On the other hand, χ^2 -test can yield a high value, thus indicating a significant dependency, even if the tested events were nearly independent. Negative correlations can be pruned by an extra test, $P(X, Y) > P(X)P(Y)$, but it does not guarantee that the high χ^2 -value is due to $X \rightarrow Y$.

Let us analyze the χ^2 -value, when $P(X, Y) = P(X)P(Y) + d$ (Table 4). Now χ^2 can be defined in the terms of d :

$$\chi^2(X, Y) = \frac{nd^2}{P(X)P(\neg X)P(Y)P(\neg Y)}.$$

χ^2 is high, when n and $|d|$ are large and $P(X)P(\neg X)P(Y)P(\neg Y)$ is small. The minimum value ($16nd^2$) is achieved, when $P(X) = P(Y) = 0.5$, and the maximum, when $P(X)$ and $P(Y)$ approach either 0 or 1. For example, if $P(X) = P(Y) = 0.01$, $\chi^2 = 10000nd^2$, and even minimal d suffices. E.g. if $n = 1000$, $d \geq 0.8 \cdot 10^{-3}$ for level 0.01, and $P(X, Y) = 0.0009$.

The problem is that if $P(X)$ and/or $P(Y)$ are large, the relative difference $\frac{d}{P(X)P(Y)}$ is small and the partial dependency between X and Y is not significant. Still the χ^2 -value can be large, because $\frac{d}{P(\neg X)P(\neg Y)}$ is large. Thus, the high χ^2 -value is due to partial dependency $\neg X \rightarrow \neg Y$, and $X \rightarrow Y$ is a false discovery (type 1 error).

Example 3 Let $P(X) = P(Y) = 1 - \varepsilon$ for arbitrary small $\varepsilon > 0$. Let d be maximal i.e. $d = P(X)(1 - P(Y)) = (1 - P(X))P(Y) = \varepsilon(1 - \varepsilon) < \varepsilon$. (The relative difference is still very small, $\frac{d}{P(X)P(Y)} = \frac{\varepsilon}{1 - \varepsilon}$.) Now $\chi^2(X, Y)$ is very large, the same as the data size, n :

$$\chi^2 = \frac{nd^2}{P(X)P(Y)(1 - P(X))(1 - P(Y))} = \frac{n\varepsilon^2(1 - \varepsilon)^2}{\varepsilon^2(1 - \varepsilon)^2} = n.$$

Still, rule $X \rightarrow Y$ is insignificant, since

$$t(X \rightarrow Y) = \frac{\sqrt{n}(1 - \varepsilon)\varepsilon}{(1 - \varepsilon)\sqrt{1 - (1 - \varepsilon)^2}} = \frac{\sqrt{n\varepsilon}}{\sqrt{2 - \varepsilon}} \rightarrow 0,$$

when $\varepsilon \rightarrow 0$.

The high χ^2 -value is due to partial dependency $\neg X \rightarrow \neg Y$, which has a high t -value:

$$t(\neg X \rightarrow \neg Y) = \frac{\sqrt{n(1-\varepsilon)}}{\sqrt{1+\varepsilon}} \rightarrow \sqrt{n},$$

when $\varepsilon \rightarrow 0$.

Rules $X \rightarrow \neg Y$ and $\neg X \rightarrow Y$ are meaningless, with

$$t = \frac{\sqrt{n\varepsilon(1-\varepsilon)}}{\sqrt{1-\varepsilon+\varepsilon^2}} < \frac{\sqrt{n\varepsilon(1-\varepsilon)}}{\sqrt{1-\varepsilon}} = \sqrt{n\varepsilon} \rightarrow 0.$$

χ^2 -measure is less likely to cause type 2 errors, i.e. to reject significant partial dependencies. The reason is that the χ^2 -value of rule $X \rightarrow Y$ increases quadratically in the terms of its t -value:

Theorem 4 *If $t(X \rightarrow Y) = K$, then $\chi^2(X, Y) \geq K^2$.*

Proof Let $x = P(X)$ and $y = P(Y)$. If $t(X \rightarrow Y) = K$, then

$$nd^2 = K^2 xy(1-xy) \text{ and } \chi^2(X, Y) = \frac{nd^2}{xy(1-x)(1-y)} = \frac{K^2(1-xy)}{(1-x)(1-y)} \geq K^2,$$

since $(1-x)(1-y) \leq 1-xy$ for all $x, y \in [0, 1]$. \square

If an association rule is just sufficiently significant, it passes also the χ^2 -test. However, the relative order of rules according to their χ^2 -values does not reflect their actual significance. If only m best rules are selected, it is possible that all of them are spurious and all significant rules are rejected.

4.2 Correlation coefficient

Some authors (e.g. [26]) have suggested Pearson correlation coefficient ϕ to measure the significance of an association rule. Traditionally, the Pearson correlation coefficient is used to measure linear dependencies between numeric attributes. When the Pearson correlation coefficient is calculated for the binary attributes, it reduces to the square root of χ^2/n :

$$\phi(X, Y) = \frac{P(X, Y) - P(X)P(Y)}{\sqrt{P(X)P(\neg X)P(Y)P(\neg Y)}} = \sqrt{\frac{\chi^2(X, Y)}{n}}.$$

Like $\chi^2(X, Y)$, $\phi(X, Y) = 0$, when $P(X, Y) = P(X)P(Y)$, and the variables are mutually independent. Otherwise, the sign of ϕ tells whether the correlation is positive ($\phi > 0$) or negative ($\phi < 0$).

The problem is to decide when the correlation is significant. General guidelines are sometimes given for defining a weak, moderate, or strong correlation, but they are rather arbitrary, because the significance depends on the data size, n . The smaller n is, the larger ϕ should be, to be statistically significant. That is why the correlation coefficient can produce very misleading results when applied to the association rule discovery.

We will first show that a rule can be insignificant, even if correlation coefficient $\phi(X, Y) = 1$. This means that ϕ -measure can produce false discoveries (type 1 error).

Observation 1 When $P(X)$ and $P(Y)$ approach 1, it is possible that $\phi(X, Y) = 1$, even if $t(X, Y) < K$ for any $K > 0$.

Proof Let $P(X) = P(Y) = 1 - \varepsilon$ for arbitrary small $\varepsilon > 0$. Let d be maximal i.e. $d = P(X)(1 - P(Y)) = (1 - P(X))P(Y) = \varepsilon(1 - \varepsilon)$. Now the correlation coefficient is 1:

$$\phi(X, Y) = \frac{d}{P(X)(1 - P(X))P(Y)(1 - P(X))} = \frac{(1 - \varepsilon)\varepsilon}{(1 - \varepsilon)\varepsilon} = 1.$$

Still, for any $K > 0$, $t(X, Y) < K$:

$$t(X, Y) = \frac{\sqrt{n}(1 - \varepsilon)\varepsilon}{(1 - \varepsilon)\sqrt{1 - (1 - \varepsilon)^2}} = \frac{\sqrt{n\varepsilon}}{\sqrt{2 - \varepsilon}} < K \Leftrightarrow \varepsilon < \frac{2K^2}{n + K^2}.$$

□

On the other hand, it is possible that ϕ -measure rejects significant rules (type 2 error), especially when n is large. The following observation shows that this can happen, when $P(X)$ and $P(Y)$ are relatively small. The smaller they are, the smaller n suffices. Therefore, we recommend that the correlation coefficient should be totally avoided as an interestingness measure for association rules.

Observation 2 It is possible that $\phi(X, Y) \rightarrow 0$, when $n \rightarrow \infty$, even if rule $X \rightarrow Y$ is significant.

Proof Let $t(X, Y) = \frac{\sqrt{nd}}{\sqrt{P(X)P(Y)(1 - P(X)P(Y))}} = K$. Then

$$d = \frac{K\sqrt{P(X)P(Y)(1 - P(X)P(Y))}}{\sqrt{n}} \text{ and } \phi(X, Y) = \frac{K\sqrt{P(X)P(Y)(1 - P(X)P(Y))}}{\sqrt{nP(X)P(Y)(1 - P(X))(1 - P(Y))}} = \frac{K\sqrt{1 - P(X)P(Y)}}{\sqrt{n(1 - P(X))(1 - P(Y))}}.$$

When $P(X) \leq p$ and $P(Y) \leq p$ for some $p < 1$, $\phi(X, Y) = \frac{K\sqrt{1-p}}{\sqrt{n(1-p)}} \rightarrow 0$, when $n \rightarrow \infty$. □

4.3 J-measure

Several objective measures used in the association rule discovery are adopted from the decision tree learning. A decision tree can be represented as a set of decision rules $X = \bar{x} \rightarrow C = c$, where $c \in \text{Dom}(C)$ is a class value. The measure functions can be used both in the rule generation (tree expansion) and post-pruning phases. In both cases, the objective is to estimate the impact of a single attribute-value condition to the generalization accuracy (i.e. how well the classifier performs outside the training set).

In the pruning phase, the test is as follows: If $M(X = \bar{x} \rightarrow C = c) \geq M(X = \bar{x}, A = a \rightarrow C = c)$, for the given measure function M , then condition $A = a$ can be pruned. This test may look fully adequate for the association rule pruning, but there is one crucial difference: in the classification, both $X = \bar{x} \rightarrow C = c$ and $\neg(X = \bar{x}) \rightarrow \neg(C = c)$ should be accurate, while for association rules it is enough that $X = \bar{x} \rightarrow C = c$ is significant. This means that the measure functions for classification rules are too restrictive for association rules, and significant associations can be missed.

Table 5 Summary of measures M for assessing association rules. The occurrence of type 1 (accepting spurious rules) and type 2 (rejecting significant rules) errors is indicated by + (occurs) and – (does not occur). In addition, all rules which contribute to $M(X \rightarrow Y)$ are listed. For all measures except $fr\&cf$, the antecedent and consequent of each rule can be switched.

M	Type 1 error	Type 2 error	Rules
$fr\&cf$	+	+	$X \rightarrow Y$
$fr\&\gamma$	–	–	$X \rightarrow Y$
χ^2	+	–	$X \rightarrow Y, \neg X \rightarrow Y, X \rightarrow \neg Y, \neg X \rightarrow \neg Y$
ϕ	+	+	$X \rightarrow Y, \neg X \rightarrow \neg Y$
J	+	+	$X \rightarrow Y, X \rightarrow \neg Y$

As an example, we analyze J -measure [24], which is often used to assess the interestingness of association rules. J -measure is an information-theoretic measure derived from the mutual information. For decision rules $X \rightarrow C$, J -measure is defined as

$$J(C|X) = P(X, C) \log \frac{P(C|X)}{P(C)} + P(X, \neg C) \log \frac{P(\neg C|X)}{P(\neg C)} \in [0, \infty[.$$

The larger J is, the more interesting the rule should be. On the other hand, $J(X, C) = 0$, when the variables X and C are mutually independent (assuming that $P(X) > 0$).

J -measure contains two terms from the mutual information, MI , between variables X and C : $MI(X, C) = J(C|X) + J(C|\neg X)$. Thus, it measures the information gain in two rules, $X \rightarrow C$ and $X \rightarrow \neg C$. Rule $X \rightarrow C$ has a high J -value, if its complement rule $X \rightarrow \neg C$ has high confidence (type 1 error). In the extreme case, when $P(C|X) = 0$, $J(C|X) = P(X) \log \frac{1}{P(\neg C)}$.

Type 2 error (rejecting true discoveries) can also occur with a suitable distribution. One reason is that J -measure omits n , which is crucial for the statistical significance.

It can be easily shown that $J(C|X) \rightarrow 0$, when $P(X, C) \rightarrow 0$ or $P(C) \rightarrow 1$. In the latter case, rule $X \rightarrow C$ cannot be significant, but it is possible that a rule is significant, even if its frequency is relatively small:

Example 4 Let $P(C|X) = 0.75$ and $P(C) = 0.5$. Now $J(C|X) = P(X)(0.75 \log 3 - 0.25) \approx 0.94P(X)$ and $t(X \rightarrow C) = \frac{\sqrt{nP(X)}}{2\sqrt{2-P(X)}}$. For example, when $P(X) = 0.25$, $t = \frac{\sqrt{n}}{2\sqrt{7}}$, which high, when n is high. Still $J(C|X) \approx 0.23$, which indicates that the rule is uninteresting.

According to [5], other information-theoretic measures are equally problematic, since they are designed for classification rules. In addition, the values are difficult to interpret, unless they express absolute independence.

In Table 5, we give a summary of the analyzed measures. For each measure, we report, whether it can produce type 1 or type 2 error and all rules which affect the measure in addition to the actually measured rule.

5 Effect of redundancy reduction

A common goal in association rule discovery is to find the most general rules (containing the minimal number of attributes) which satisfy the given search criteria. There is no sense

to output complex rules $X \rightarrow Y$, if their generalizations $Z \rightarrow Y$, $Z \subsetneq X$ are at least equally significant. Generally, the goal is to find *minimal* (or most general) *interesting rules*, and prune out *redundant rules* [3].

5.1 General definition

Generally, redundancy can be defined as follows:

Definition 4 (Minimal and redundant rules) Given some interestingness measure M , rule $X \rightarrow Y$ is a minimal rule, if there does not exist any rule $X' \rightarrow Y'$ such that $X' \subseteq X$, $Y \subseteq Y'$ and $M(X' \rightarrow Y') \geq M(X \rightarrow Y)$. If the rule is not minimal, then it is redundant.

Measure M can be t -measure, J -measure, χ^2 -measure, or any function which increases with the interestingness. In the traditional frequency-confidence-framework with minimum frequency and confidence thresholds, M can be defined as

$$M(X \rightarrow Y) = \begin{cases} 1 & \text{when } fr(X \rightarrow Y) \geq \min_{fr} \text{ and } cf(X \rightarrow Y) \geq \min_{cf}, \\ 0 & \text{otherwise.} \end{cases}$$

The motivation for the redundancy reduction is two-fold: First, a smaller set of general rules is easier to interpret than a large set of complex and often overlapping rules. Second, the problem complexity is reduced, because it is enough to find a small subset of all interesting rules. Thus, it is possible at least in principle to perform the search more efficiently.

In the previous research, redundancy has been defined in various ways. An important distinction is whether the redundancy refers to the *interestingness* of a rule or the *representation* of rules. In the first case, a rule is considered redundant, if there are more general rules which are at least equally interesting. Such a redundant rule contains no new information and it can be pruned out. In the second case, even an interesting rule is considered redundant, if it (or its frequency and confidence) can be derived from the other rules in the representation. Now the rule is not pruned out, but it is not represented explicitly. Examples of such *condensed representations* [15] are closed [20], free [6], and non-derivable sets [8].

We will briefly analyze the effect of two common pruning techniques on discovering statistically significant rules. The important question is, whether a statistically significant rule can be pruned out as “redundant” causing type 2 error.

5.2 Redundant rules

According to a classical definition (e.g. [1]), rule $X \rightarrow Y$ is redundant, if there exists $Z \subsetneq X$ such that $fr(X \rightarrow Y) = fr(Z \rightarrow Y)$. The aim of this definition is to achieve a compact representation of all frequent and strong association rules. The justification is sensible in the traditional frequency-confidence-framework with fixed thresholds \min_{fr} and \min_{cf} : If rule $Z \rightarrow Y$ is frequent and strong enough, then all its specializations $X \rightarrow Y$ with $P(X, Y) = P(Z, Y)$ are also frequent and strong.

However, this definition is not adequate, if the goal is to find the most significant rules. In fact, it causes always type 2 error (rejects the most significant rules), unless $P(X) = P(Z)$. If $P(X) < P(Z)$, then rule $X \rightarrow Y$ has higher confidence and is more significant than $Z \rightarrow Y$:

Theorem 5 If $fr(X \rightarrow Y) = fr(Z \rightarrow Y)$ for some $Z \subsetneq X$, then

- (i) $cf(X \rightarrow Y) \geq cf(Z \rightarrow Y)$, and
- (ii) $cf(X \rightarrow Y) = cf(Z \rightarrow Y)$ only if $P(X) = P(Z)$.

Proof Let $X \rightarrow Y$ redundant, i.e. $\exists Z \subsetneq X$ such that $fr(X \rightarrow Y) = fr(Z \rightarrow Y)$. Let $X = ZQ$, $P(X) = P(Z, Q)$ and $P(X, Y) = P(Z, Y, Q)$. According to the redundancy condition $P(Z, Y, Q) = P(Z, Y)$.

Now $cf(X \rightarrow Y) - cf(Z \rightarrow Y) = \frac{P(Z, Y, Q)}{P(Z, Q)} - \frac{P(Z, Y)}{P(Z)} = \frac{P(Z, Y)}{P(Z, Q)} - \frac{P(Z, Y)P(Q|Z)}{P(Z, Q)} \geq 0$, because $P(Q|Z) \leq 1$. $P(Q|Z) = 1$ iff $P(Z) = P(Z, Q) = P(X)$. \square

Type 1 error (accepting spurious rules) is also likely, because the least significant rules are output. So, in the worst case all significant rules are pruned and only spurious rules are presented.

In the context of closed sets, the definition of redundancy is similar (e.g. [33]). However, now it is required that there exists more general rule $X' \rightarrow Y'$, $X' \subsetneq X$ and $Y' \subseteq Y$, such that $P(X, Y) = P(X', Y')$ and $P(X) = P(X')$. This means that $X \rightarrow Y$ and $X' \rightarrow Y'$ have the same frequency and confidence. Still it is possible that $\gamma(X \rightarrow Y) > \gamma(X' \rightarrow Y')$ (i.e. $P(Y) < P(Y')$) and the more significant rule is pruned.

5.3 Productive rules

According to another common interpretation, rule $X \rightarrow Y$ is considered redundant or uninteresting, if there exists more general rule $Z \rightarrow Y$, $Z \subsetneq X$, such that $P(Y|Z) \geq P(Y|X)$. Following [30] we call these rules *non-productive*. If $P(Y|X) > P(Y|Z)$ for all $Z \subsetneq X$, rule $X \rightarrow Y$ is *productive*. The aim of this definition is to prune out rules which are less interesting than their generalizations.

The heuristic works correctly and avoids type 2 error. For non-productive rule $X \rightarrow Y$, $\gamma(X \rightarrow Y) \leq \gamma(Z \rightarrow Y)$. In addition, we know that $P(X, Y) \leq P(Z, Y)$ and $X \rightarrow Y$ cannot be more significant than $Z \rightarrow Y$. In practice, this means that X contains some attributes which are either independent from Y or negatively correlated with Y .

Generally, it is required that the improvement of rule $X \rightarrow Y$ is sufficient [13]:

$$imp(X \rightarrow Y) = cf(X \rightarrow Y) - \max_{Z \subsetneq X} \{cf(Z \rightarrow Y)\} \geq min_{imp}. \quad (4)$$

In practice, each rule is compared only to its immediate generalizations ($|Z| = |X| - 1$). If Equation (4) does not hold for some Z , then rule $X \rightarrow Y$ and all its specializations are pruned. The problem is that now there could be $X' \supsetneq X$ such that $cf(X' \rightarrow Y) > cf(Z \rightarrow Y)$ and which is statistically more significant than $Z \rightarrow Y$. This rule is not discovered, because the whole branch was pruned. Thus, the pruning condition should not be used to restrict the search space.

Instead, the pruning condition can be used in the post-processing phase, where a rule is compared to all its generalizations. We will show that requirement $min_{imp} = 0$ is a necessary but not sufficient condition for the superiority of $X \rightarrow Y$ over $Z \rightarrow Y$. This means that type 2 error does not occur, but type 1 error (accepting spurious rules) is possible. However, when $min_{imp} > 0$, also type 2 error is possible, and non-redundant significant rules can be missed.

The following theorem gives a necessary and sufficient condition for the superiority of $X \rightarrow Y$:

Theorem 6 Let $X = Z, Q$ for some $Z, Q \subseteq R$. Rule $X \rightarrow Y$ is more significant than $Z \rightarrow Y$, if and only if

$$\frac{P(Y|X) - P(Y)}{P(Y|Z) - P(Y)} > \frac{\sqrt{1 - P(X)P(Y)}}{\sqrt{P(Q|Z)(1 - P(Z)P(Y))}}.$$

Proof

$$\begin{aligned} t(X \rightarrow Y) > t(Z \rightarrow Y) &\Leftrightarrow \\ \frac{\sqrt{n}P(X)(P(Y|X) - P(Y))}{\sqrt{P(X)P(Y)(1 - P(X)P(Y))}} &> \frac{\sqrt{n}P(Z)(P(Y|Z) - P(Y))}{\sqrt{P(Z)P(Y)(1 - P(Z)P(Y))}} \Leftrightarrow \\ \frac{P(Z)P(Q|Z)(P(Y|X) - P(Y))}{\sqrt{P(Z)P(Q|Z)P(Y)(1 - P(X)P(Y))}} &> \frac{P(Z)(P(Y|Z) - P(Y))}{\sqrt{P(Z)P(Y)(1 - P(Z)P(Y))}} \Leftrightarrow \\ \frac{\sqrt{P(Q|Z)}(P(Y|X) - P(Y))}{\sqrt{(1 - P(X)P(Y))}} &> \frac{(P(Y|Z) - P(Y))}{\sqrt{(1 - P(Z)P(Y))}} \Leftrightarrow \\ \frac{(P(Y|X) - P(Y))}{(P(Y|Z) - P(Y))} &> \frac{\sqrt{(1 - P(X)P(Y))}}{\sqrt{P(Q|Z)(1 - P(Z)P(Y))}} \end{aligned}$$

□

Since $\frac{\sqrt{(1 - P(X)P(Y))}}{\sqrt{P(Q|Z)(1 - P(Z)P(Y))}} \geq 1$, it follows that

Corollary 1 If $t(X \rightarrow Y) > t(Z \rightarrow Y)$, then $P(Y|X) > P(Y|Z)$ and $\text{imp}(X \rightarrow Y) > 0$.

Now we can give a better pruning condition than Equation (4):

Corollary 2 If

$$\frac{P(Y|X) - P(Y)}{P(Y|Z) - P(Y)} \leq \frac{1}{\sqrt{P(Q|Z)}},$$

then $t(X \rightarrow Y) < t(Z \rightarrow Y)$.

The condition can be expressed equivalently as

$$\text{imp}(X \rightarrow Y) \leq \frac{(P(Y|Z) - P(Y))(1 - \sqrt{P(Q|Z)})}{P(Q|Z)}.$$

This pruning condition is more efficient than $\text{min}_{\text{imp}} = 0$, but still it does not prune out any non-redundant significant rules. Generally, the correct threshold min_{imp} depends on the rules considered and $P(Q|Z)$, and no absolute thresholds (other than $\text{min}_{\text{imp}} = 0$) can be used.

6 Conclusions

In this paper, we have formalized an important problem: how to find statistically significant association rules. We have inspected the most common interest measures and search techniques from the statistical point of view. For all methods, we have analyzed, whether they can cause type 1 error (accept spurious rules) or type 2 error (reject significant rules) and the conditions under which the errors can occur.

The conclusions are the following: The traditional frequency-confidence framework should be abandoned, because it can cause both type 1 and type 2 errors. The simplest correction is to adopt the so called frequency-dependence framework, where the degree of dependence is used instead of confidence. If the minimum frequency is set carefully (Theorem 2), no significant rules are missed. On the other hand, all insignificant rules with the desired level of significance can be pruned, using the t -measure.

The χ^2 -measure works correctly only if all significant partial dependencies in the data are actually full dependencies. When it is used to assess association rules, several spurious rules can be accepted (type 1 error). Type 2 error does not occur, if the partial dependencies are sufficiently significant, but the ranking order of association rules can be incorrect.

Pearson correlation coefficient ϕ is not recommendable for assessing association rules. It can easily cause both type 1 and type 2 errors. J -measure can also cause both error types, although type 1 error (accepting spurious rules) is more likely. Both ϕ and J omit the data size, n , and it can be hard to decide proper cut-off values for significant dependencies.

Finally, we analyzed two common redundancy reduction techniques, which compare rule $X \rightarrow Y$ to its generalizations $Z \rightarrow Y$, $Z \subsetneq X$. We showed that the minimum improvement condition, $\text{imp}(X \rightarrow Y) = cf(X \rightarrow Y) - \max_{Z \subsetneq X} \{cf(Z \rightarrow Y)\} \geq \text{min}_{\text{imp}}$, works correctly, if $\text{min}_{\text{imp}} = 0$. However, it cannot be used to restrict the search space, but only for post-processing. If $\text{min}_{\text{imp}} > 0$, significant rules can be missed. We gave also a more efficient pruning condition, which can be used to prune redundant rules without type 2 error.

The second redundancy condition, $fr(X \rightarrow Y) = fr(Z \rightarrow Y)$, does just the opposite and prunes out the more significant, specific rules. I.e. it causes always type 2 error, unless $P(X) = P(Z)$.

In the future research, these new insights should be utilized in the search algorithms for discovering the statistically most significant rules. The computational efficiency of such algorithms is a potential bottle-neck, but the starting point looks promising: while the small min_{fr} increases the number of frequent sets, we can use γ -based pruning to restrict the search space, and the effects may compensate each other.

Acknowledgements I thank professor O. Kolehmainen for checking the validity of statistical arguments and professor M. Nykänen for his valuable comments.

References

1. Aggarwal, C., Yu, P.: A new framework for itemset generation. In: Proceedings of the Seventeenth ACM SIGACT-SIGMOD-SIGART Symposium on Principles of Database Systems (PODS 1998), pp. 18–24. ACM Press, New York, USA (1998)
2. Agrawal, R., Imielinski, T., Swami, A.: Mining association rules between sets of items in large databases. In: P. Buneman, S. Jajodia (eds.) Proceedings of the 1993 ACM SIGMOD International Conference on Management of Data, pp. 207–216. Washington, D.C. (1993)
3. Bastide, Y., Pasquier, N., Taouil, R., Stumme, G., Lakhal, L.: Mining minimal non-redundant association rules using frequent closed itemsets. In: Proceedings of the First International Conference on Compu-

- tational Logic (CL'00), *Lecture Notes in Computer Science*, vol. 1861, pp. 972–986. Springer-Verlag, London, UK (2000)
4. Berzal, F., Blanco, I., Sánchez, D., Miranda, M.A.V.: A new framework to assess association rules. In: Proceedings of the 4th International Conference on Advances in Intelligent Data Analysis (IDA'01), *Lecture Notes in Computer Science*, vol. 2189, pp. 95–104. Springer-Verlag, London, UK (2001)
 5. Blanchard, J., Guillet, F., Gras, R., Briand, H.: Using information-theoretic measures to assess association rule interestingness. In: Proceedings of the Fifth IEEE International Conference on Data Mining (ICDM'05), pp. 66–73. IEEE Computer Society, Washington, USA (2005)
 6. Boulicaut, J.F., Bykowski, A., Rigotti, C.: Approximation of frequency queries by means of free-sets. In: Proceedings of the 4th European Conference Principles of Data Mining and Knowledge Discovery (PKDD'00), *Lecture Notes in Computer Science*, vol. 1910, pp. 75–85. Springer-Verlag (2000)
 7. Brin, S., Motwani, R., Silverstein, C.: Beyond market baskets: Generalizing association rules to correlations. In: J. Peckham (ed.) Proceedings ACM SIGMOD International Conference on Management of Data, pp. 265–276. ACM Press (1997)
 8. Calders, T., Goethals, B.: Mining all non-derivable frequent itemsets. In: Proceedings of the 6th European Conference on Principles of Data Mining and Knowledge Discovery (PKDD'02), *Lecture Notes in Computer Science*, vol. 2431, pp. 74–85. Springer-Verlag, London, UK (2002)
 9. Casas-Garriga, G.: Statistical strategies for pruning all the uninteresting association rules. In: R.L. de Mántaras, L. Saitta (eds.) Proceedings of the 16th European Conference on Artificial Intelligence (ECAI'2004), pp. 430–434. IOS Press (2004)
 10. Dong, X., Sun, F., Han, X., Hou, R.: Study of positive and negative association rules based on multi-confidence and chi-squared test. In: X. Li, O. Zaiane, Z. Li (eds.) Proceedings of the Second International Conference on Advanced Data Mining and Applications (ADMA), *Lecture Notes in Computer Science*, vol. 4093, pp. 100–109. Springer-Verlag, Berlin / Heidelberg (2006)
 11. Jaroszewicz, S., Simovici, D.A.: A general measure of rule interestingness. In: Proceedings of 5th European Conference on Principles of Data Mining and Knowledge Discovery (PKDD 2001), *Lecture Notes in Computer Science*, vol. 2168, pp. 253–265. Springer (2001)
 12. Jermaine, C.: Finding the most interesting correlations in a database: how hard can it be? *Information Systems* **30**(1), 21–46 (2005)
 13. Jr., R.B., Agrawal, R., Gunopulos, D.: Constraint-based rule mining in large, dense databases. *Data Mining and Knowledge Discovery* **4**(2/3), 217–240 (2000)
 14. Liu, B., Hsu, W., Ma, Y.: Pruning and summarizing the discovered associations. In: Proceedings of the fifth ACM SIGKDD international conference on Knowledge discovery and data mining (KDD'99), pp. 125–134. ACM Press, New York, USA (1999)
 15. Mannila, H., Toivonen, H.: Multiple uses of frequent sets and condensed representations (extended abstract). In: Proceedings of the Second International Conference on Knowledge Discovery and Data Mining (KDD'96), pp. 189–194 (1996)
 16. Meo, R.: Theory of dependence values. *ACM Transactions on Database Systems* **25**(3), 380–406 (2000)
 17. Milton, J., Arnold, J.: Introduction to Probability and Statistics: Principles and Applications for Engineering and the Computing Sciences, 4th edn. McGraw-Hill, New York (2003)
 18. Morishita, S., Nakaya, A.: Parallel branch-and-bound graph search for correlated association rules. In: M. Zaki, C.T. Ho (eds.) Revised Papers from Large-Scale Parallel Data Mining, Workshop on Large-Scale Parallel KDD Systems, SIGKDD, *Lecture Notes in Computer Science*, vol. 1759, pp. 127–144. Springer-Verlag, London, UK (2000)
 19. Morishita, S., Sese, J.: Transversing itemset lattices with statistical metric pruning. In: Proceedings of the nineteenth ACM SIGMOD-SIGACT-SIGART symposium on Principles of database systems (PODS'00), pp. 226–236. ACM Press, New York, USA (2000)
 20. Pasquier, N., Bastide, Y., Taouil, R., Lakhal, L.: Discovering frequent closed itemsets for association rules. In: Proceedings of the 7th International Conference on Database Theory (ICDT'99), *Lecture Notes in Computer Science*, vol. 1540, pp. 398–416. Springer-Verlag (1999)
 21. Piatetsky-Shapiro, G.: Discovery, analysis, and presentation of strong rules. In: G. Piatetsky-Shapiro, W. Frawley (eds.) *Knowledge Discovery in Databases*, pp. 229–248. AAAI/MIT Press (1991)
 22. Shaffer, J.: Multiple hypothesis testing. *Annual Review of Psychology* **46**, 561–584 (1995)
 23. Silverstein, C., Brin, S., Motwani, R.: Beyond market baskets: Generalizing association rules to dependence rules. *Data Mining and Knowledge Discovery* **2**(1), 39–68 (1998)
 24. Smyth, P., Goodman, R.: An information theoretic approach to rule induction from databases. *IEEE Transactions on Knowledge and Data Engineering* **4**(4), 301–316 (1992)
 25. Srikant, R., Agrawal, R.: Mining quantitative association rules in large relational tables. *SIGMOD Record* **25**(2), 1–12 (1996)
 26. Tan, P., Kumar, V.: Interestingness measures for association patterns: A perspective. Tech. Rep. TR00-036, Department of Computer Science, University of Minnesota (2000). URL citeseer.ist.psu.edu/tan00interestingness.html

27. Tan, P.N., Kumar, V., Srivastava, J.: Selecting the right interestingness measure for association patterns. In: Proceedings of the eighth ACM SIGKDD international conference on Knowledge discovery and data mining (KDD'02), pp. 32–41. ACM Press, New York, USA (2002)
28. Vilalta, R., Oblinger, D.: A quantification of distance bias between evaluation metrics in classification. In: Proceedings of the Seventeenth International Conference on Machine Learning (ICML'00), pp. 1087–1094. Morgan Kaufmann Publishers Inc., San Francisco, USA (2000)
29. Webb, G.: Discovering significant rules. In: Proceedings of the 12th ACM SIGKDD international conference on Knowledge discovery and data mining (KDD'06), pp. 434–443. ACM Press, New York, USA (2006)
30. Webb, G.I.: Discovering significant patterns. *Machine Learning* **68**(1), 1–33 (2007)
31. Wu, X., Zhang, C., Zhang, S.: Efficient mining of both positive and negative association rules. *ACM Transactions on Information Systems* **22**(3), 381–405 (2004)
32. Yao, Y., Zhong, N.: An analysis of quantitative measures associated with rules. In: Proceedings of the Third Pacific-Asia Conference on Methodologies for Knowledge Discovery and Data Mining (PAKDD'99), pp. 479–488. Springer-Verlag, London, UK (1999)
33. Zaki, M.: Generating non-redundant association rules. In: Proceedings of the sixth ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD'00), pp. 34–43. ACM, New York, USA (2000)