# The Sample Complexity of Learning Linear Predictors with the Squared Loss

# Ohad Shamir Weizmann Institute of Science ohad.shamir@weizmann.ac.il

#### Abstract

In this short note, we provide tight sample complexity bounds for learning linear predictors with respect to the squared loss. Our focus is on an agnostic setting, where no assumptions are made on the data distribution. This contrasts with standard results in the literature, which either make distributional assumptions, refer to specific parameter settings, or use other performance measures.

# 1 Introduction

In machine learning and statistics, the squared loss is the most commonly used loss for measuring real-valued predictions: Given a prediction p and actual target value y, it is defined as  $\ell(p,y) = (p-y)^2$ . It is intuitive, has a convenient analytical form, and has been extremely well-studied.

In this note, we concern ourselves with learning linear predictors with respect to the squared loss, in a standard agnostic learning framework. Formally, for some fixed parameters X,Y,B, we assume the existence of an unknown distribution over  $\{\mathbf{x} \in \mathbb{R}^d : \|\mathbf{x}\| \leq 1\} \times \{y \in \mathbb{R} : |y| \leq Y\}$ , from which we are given a training set  $S = \{\mathbf{x}_i, y_i\}_{i=1}^m$  of m i.i.d. examples, consisting of pairs of instances  $\mathbf{x}$  and target values y. Given a linear predictor  $\mathbf{x} \mapsto \langle \mathbf{w}, \mathbf{x} \rangle$ , its risk with respect to the squared loss is defined as

$$R(\mathbf{w}) = \mathbb{E}_{(\mathbf{x},y)}[(\langle \mathbf{w}, \mathbf{x} \rangle - y)^2].$$

Our goal is to find a linear predictor w from the hypothesis class of norm-bounded linear predictors,

$$\mathcal{W} = \{ \mathbf{w} : ||\mathbf{w}|| \le B \},$$

such that its excess risk

$$R(\mathbf{w}) - \min_{\mathbf{w} \in \mathcal{W}} R(\mathbf{w})$$

with respect to the best possible predictor in W is as small as possible. We focus here on the expected excess risk (over the randomness of the training set and algorithm), and consider how it is affected by the problem parameters Y, B, d and the sample size m, uniformly over any distribution.

Despite a huge literature on learning with the squared loss, we were unable to locate an explicit and self-contained analysis for this question. The existing results (some examples include [4, 5, 6, 10, 1, 7]) all appear to differ from our setting in one or more of the following manners:

• Distributional Assumptions: In our agnostic setting, we assume nothing whatsoever about the data distribution, other than boundedness (as specified by X, Y). In contrast, most existing works rely on additional assumptions. Perhaps the most common assumption is a well-specified model, under which there exists a fixed  $\mathbf{w} \in \mathbb{R}^d$  such that  $y = \langle \mathbf{w}, \mathbf{x} \rangle + \xi$ , where  $\xi$  is a zero-mean noise term. Other works impose some moment or other conditions on the distribution of  $\mathbf{x}$ , or consider a fixed design setting where the data instances are not sampled i.i.d.. These assumptions usually lead to excess risk bounds which scale (at least in finite dimensions) as  $dY^2/m$ , independent of the norm bound B. However, as we will see later, this is not the behavior in the distribution-free setting.

- Bounds not on the excess risk: Many of the existing results are not on the excess risk, but rather on  $\mathbb{E}[\|\mathbf{w}-\mathbf{w}^*\|^2]$  or  $\mathbb{E}[(\langle \mathbf{w}, \mathbf{x} \rangle \langle \mathbf{w}^*, \mathbf{x} \rangle)^2]$ , where  $\mathbf{w}^* = \arg\min_{\mathbf{w} \in \mathcal{W}} R(\mathbf{w})$ . The former measure is relevant for parameter estimation, while the latter measure can be shown to equal the excess risk when  $\mathbf{w}^* = \arg\min_{\mathbf{w} \in \mathbb{R}^d} R(\mathbf{w})$  (i.e.  $B = \infty$  see Lemma 1 below). However, when we deal with the hypothesis class of norm-bounded predictors, then the excess risk can be larger by an arbitrary factor<sup>1</sup>. Therefore, upper bounds on these measures do not imply upper bounds on the excess risk in our setting. We remark that in our distribution-free setting, we must constrain the hypothesis class, since if our hypothesis class contains all linear predictors ( $B = \infty$ ), then the lower bounds below imply that non-trivial learning is impossible with any sample size (regardless of the dimension d).
- Bounded Functions: Many learning theory results for the squared loss (such as thosed based on fat-shattering techniques) assume that the predictor functions and target values are bounded in some fixed interval (such as [-1,+1]). In our setting, this would correspond to assuming  $B,Y \leq 1$ . Other results assume Lipschitz loss functions, which is not satisfied for the squared loss. One notable exception is [9], which analyze smooth and strongly-convex losses (such as the squared loss) and provide tight bounds. However, their results apply either when the functions are bounded by 1, or when d is extremely large or infinite dimensional. In contrast, we provide more general results which hold for any d and when the functions are not necessarily bounded by 1.
- Collapsing Problem Parameters Together: Many results implicitly take Y to equal the largest possible prediction,  $\sup_{\mathbf{w},\mathbf{x}} |\langle \mathbf{w},\mathbf{x} \rangle| = B$ , and give results only in terms of B. However, we will see that B and Y affect the excess risk in a different manner, and it is thus important to discern between them. Moreover, B and B can often have very different magnitudes. For example, in learning problems where the instances  $\mathbf{x}$  tend to be sparse, we may want to have the norm bound B of the predictor to scale with the dimension  $\mathbf{w}$ , while the bound on the target values B remain a fixed constant.

## 2 Main Result

Our main result is the following lower bound on the attainable excess risk:

**Theorem 1.** There exists a universal constant c, such that for any dimension d, sample size m, target value bound Y, predictor norm bound  $B \ge 2Y$ , and for any algorithm returning a linear predictor  $\hat{\mathbf{w}}$ , there exists a valid data distribution such that

$$\mathbb{E}[R(\hat{\mathbf{w}}) - R(\mathbf{w}^*)] \ge c \min\left\{Y^2, \frac{B^2 + dY^2}{m}, \frac{BY}{\sqrt{m}}\right\},\,$$

where  $\mathbf{w}^* = \arg\min_{\mathbf{w}: ||\mathbf{w}|| \le B} R(\mathbf{w}).$ 

Based on existing results in the literature, this bound has essentially matching upper bounds, up to logarithmic factors:

- Using the trivial zero predictor  $\hat{\mathbf{w}} = \mathbf{0}$ , we are guaranteed that  $\mathbb{E}[R(\hat{\mathbf{w}}) R(\mathbf{w}^*)] \leq \mathbb{E}[R(\hat{\mathbf{w}})] = \mathbb{E}[\langle \mathbf{0}, \mathbf{x} \rangle y)^2] = \mathbb{E}[y^2] \leq Y^2$ .
- Using the Vovk-Azoury-Warmuth forecaster and a standard online-to-batch conversion ([11, 2, 3]), we have an algorithm for which  $\mathbb{E}[R(\hat{\mathbf{w}}) R(\mathbf{w}^*)] \leq \mathcal{O}\left(\frac{B^2 + dY^2 \log(1 + m/d)}{m}\right)$ .
- Alternatively, by corollary 3 in [9]  $^2$ , using mirror descent with an online-to-batch conversion gives us an algorithm for which  $\mathbb{E}[R(\hat{\mathbf{w}}) R(\mathbf{w}^*)] \leq \mathcal{O}\left(\frac{BY}{\sqrt{m}} + \frac{B^2}{m}\right)$ . In the regime where this bound is smaller than  $Y^2$ , it can be verified that  $BY/\sqrt{m}$  is the dominant term, in which case we get an  $\mathcal{O}(BY/\sqrt{m})$  bound.

For example, consider a distribution on (x,y) such that (x,y)=(1,1) with probability 1, and  $\mathcal{W}=\{w:w\in[-1/2,1/2]\}$ . Then clearly,  $w^*=1/2$ , and  $\mathbb{E}[(wx-w^*x)^2]=\mathbb{E}[(w-w^*)^2]=(1/2-w)^2$ . However, the excess risk equals  $(w-1)^2-(1/2-1)^2=w^2-2w+3/4=(1/2-w)^2+(1/2-w)$ . This is larger than the excess risk by an additive factor of (1/2-w), and a multiplicative factor of  $\frac{1}{1/2-w}$  – arbitrarily large if w is close to  $w^*=1/2$ .

Where  $\bar{L^*} \leq Y^2$  and H = 2 for the squared loss.

Taking the best of these algorithmic approaches, we get the minimum of these upper bounds, i.e. we can find a predictor  $\hat{\mathbf{w}}$  for which

$$\mathbb{E}[R(\hat{\mathbf{w}}) - R(\mathbf{w}^*)] \le \mathcal{O}\left(\min\left\{Y^2, \frac{B^2 + dY^2 \log\left(1 + \frac{m}{d}\right)}{m}, \frac{BY}{\sqrt{m}}\right\}\right).$$

We conjecture that the same bound can be shown for empirical risk minimization (i.e. given a training set  $\{(\mathbf{x}_i, y_i)\}_{i=1}^m$ , return  $\hat{\mathbf{w}} = \min_{\mathbf{w}: \|\mathbf{w}\| \le B} \frac{1}{m} \sum_{i=1}^m (\langle \mathbf{w}, \mathbf{x}_i \rangle - y_i)^2$ ).

This result has some interesting consequences: First, it implies that even when d=1 (i.e. a one-dimensional problem), there is a non-trivial dependence on the norm bound B. This is in contrast to results under the well-specified model or other common distributional assumptions, which lead to upper bounds independent of B. Second, it shows that in a finite-dimensional setting, although the squared loss  $(\langle \mathbf{w}, \mathbf{x} \rangle - y)^2$  may appear symmetric with respect to y and  $\langle \mathbf{w}, \mathbf{x} \rangle$ , the attainable excess risk is actually much more sensitive to the bound Y on |y| than to the bound B on  $|\langle \mathbf{w}, \mathbf{x} \rangle|$ , due to the d factor. For example, if Y is a constant, then B can be as large as the dimension d without affecting the leading term of the excess risk. Third, in the context of online learning, it implies that the Vovk-Azoury-Warmuth forecaster is essentially optimal in our setting and for a finite-dimensional regime, in terms of its dependence on both d and B (the lower bounds in [11, 8] do not show an explicit dependence on B).

### 3 Proof of Thm. 1

The proof of our main result consist of two separate lower bounds, each of which uses a different construction. The theorem follows by combining them and performing a few simplifications.

We begin by recalling the following result, which follows from the well-known orthogonality principle:

**Lemma 1.** Let  $R(\mathbf{w}) = \mathbb{E}[(\langle \mathbf{w}, \mathbf{x} \rangle - y)^2]$ , and  $\mathbf{w}^* = \arg\min_{\mathbf{w}: ||\mathbf{w}|| \le B} R(\mathbf{w})$ . Then for any  $\mathbf{w} \in \mathbb{R}^d$ , it holds that

$$R(\mathbf{w}) - R(\mathbf{w}^*) \ge \mathbb{E}[(\langle \mathbf{w}, \mathbf{x} \rangle - \langle \mathbf{w}^*, \mathbf{x} \rangle)^2],$$

with equality when  $B = \infty$ 

*Proof Sketch.* For any  $\mathbf{w} \in \mathbb{R}^d$ , define the linear function  $f_{\mathbf{w}} : \mathbb{R}^d \mapsto \mathbb{R}$  by  $f_{\mathbf{w}}(\mathbf{x}) = \langle \mathbf{w}, \mathbf{x} \rangle$ . Then  $\{f_{\mathbf{w}}(\cdot) : \|\mathbf{w}\| \le B\}$  corresponds to a closed convex set in the  $L^2$  function space defined via the inner product  $\langle f, g \rangle = \mathbb{E}_{\mathbf{x}}[f(\mathbf{x})g(\mathbf{x})]$  and norm  $\|f\|^2 = \mathbb{E}_{\mathbf{x}}[f^2(\mathbf{x})]$ . Moreover, letting  $\eta(\mathbf{x}) = \mathbb{E}[y|\mathbf{x}]$ , we have

$$R(\mathbf{w}) - R(\mathbf{w}^*) = \mathbb{E}[(\langle \mathbf{w}, \mathbf{x} \rangle - y)^2] - \mathbb{E}[(\langle \mathbf{w}, \mathbf{x} \rangle - y)^2] = \mathbb{E}[(f_{\mathbf{w}}(\mathbf{x}) - \eta(\mathbf{x}))^2] - \mathbb{E}[(f_{\mathbf{w}^*}(\mathbf{x}) - \eta(\mathbf{x}))^2] = \|f_{\mathbf{w}} - \eta\|^2 - \|f_{\mathbf{w}^*} - \eta\|^2.$$

In this representation, the inequality in the lemma reduces to

$$||f_{\mathbf{w}} - f_{\mathbf{w}^*}||^2 + ||f_{\mathbf{w}^*} - \eta||^2 \le ||f_{\mathbf{w}} - \eta||^2.$$

When  $B=\infty$ , then  $f_{\mathbf{w}^*}$  is the projection of  $\eta$  on the linear sub-space of linear functionals, hence the inequality above holds with equality by the pythagorean theorem. When  $B<\infty$ , then  $f_{\mathbf{w}^*}$  is the projection of  $\eta$  on a constrained subset of this linear space, and we only have an inequality.

Our first construction provides an excess risk lower bound even when we deal with one-dimensional problems:

**Theorem 2.** There exists a universal constant c, such that for any sample size m, target value bound Y, predictor norm bound  $B \ge 2Y$ , and any algorithm returning a linear predictor  $\hat{\mathbf{w}}$ , there exists a data distribution in d=1 dimensions such that

$$\mathbb{E}[R(\hat{w}) - R(w^*)] \ge c \min\left\{Y^2, \frac{B^2}{m}\right\}.$$

*Proof.* Let  $\alpha, \gamma$  be small positive parameters in (0,1] to be chosen later, such that  $\alpha > \gamma$ , and consider the following two distributions over (x,y):

• Distribution 
$$\mathcal{D}_0$$
:  $y = Y$  w.p. 1;  $x = \begin{cases} Y/B & \text{w.p. } \alpha \\ 0 & \text{w.p. } 1 - \alpha \end{cases}$ 

• Distribution 
$$\mathcal{D}_1$$
:  $y = Y$  w.p. 1;  $x = \begin{cases} 1 & \text{w.p. } \gamma \\ Y/B & \text{w.p. } \alpha - \gamma . \\ 0 & \text{w.p. } 1 - \alpha \end{cases}$ 

Note that since  $B \ge 2Y$ ,  $|x| \le 1$ , so these are indeed valid distributions. Intuitively, in both distributions x is small most of the time, but under  $\mathcal{D}_1$  it can occasionally have a "large" value of 1. Unless the sample size is large enough, it is not possible to distinguish between these two distributions, and this will lead to an excess risk lower bound.

Let  $\mathbb{E}_0$  and  $\mathbb{E}_1$  denote expectations with respect to  $\mathcal{D}_0$  and  $\mathcal{D}_1$  respectively. Let

$$w_0^* = B$$

denote the optimal predictor under  $\mathcal{D}_0$ , and let

$$w_1^* = \frac{\mathbb{E}_1[yx]}{\mathbb{E}_1[x^2]} = \frac{(Y^2/B)(\alpha - \gamma) + Y\gamma}{(Y^2/B^2)(\alpha - \gamma) + \gamma} = B\frac{Y^2(\alpha - \gamma) + BY\gamma}{Y^2(\alpha - \gamma) + B^2\gamma}$$

denote the optimal predictor under  $\mathcal{D}_1$ . Note that  $w_1^* \geq w_0^*$ , and moreover,

$$(w_1^* - w_0^*)^2 = B^2 \left( \frac{Y^2(\alpha - \gamma) + BY\gamma}{Y^2(\alpha - \gamma) + B^2\gamma} - 1 \right)^2 = B^4\gamma^2 \left( \frac{Y - B}{Y^2\alpha + (B^2 - Y^2)\gamma} \right)^2 \ge B^4\gamma^2 \left( \frac{Y - B}{Y^2\alpha + B^2\gamma} \right)^2 \tag{1}$$

By Yao's minimax principle, it is sufficient to show that when choosing either  $\mathcal{D}_0$  or  $\mathcal{D}_1$  uniformly at random, and generating a dataset according to that distribution, any deterministic algorithm attains the lower bound in the theorem. Using Lemma 1, and the notation  $\Pr_0$  (respectively  $\Pr_1$ ) to denote probabilities with respect to  $\mathcal{D}_0$  (respectively  $\mathcal{D}_1$ ), we have

$$\begin{split} \mathbb{E}\left[R(\hat{w}) - R(w^*)\right] &= \frac{1}{2} \left( \mathbb{E}_0[(\hat{w}x - w_0^*x)^2] + \mathbb{E}_1[(\hat{w}x - w_1^*x)^2] \right) \\ &\geq \frac{1}{2} \frac{Y^2 \alpha}{B^2} \left( \mathbb{E}_0[(\hat{w} - w_0^*)^2] + \mathbb{E}_1[(\hat{w} - w_1^*)^2] \right) \\ &\geq \frac{1}{2} \frac{Y^2 \alpha}{B^2} \left( \frac{w_1^* - w_0^*}{2} \right)^2 \left( \Pr_0\left( \hat{w} < \frac{w_0^* + w_1^*}{2} \right) + \Pr_1\left( \hat{w} \ge \frac{w_0^* + w_1^*}{2} \right) \right) \\ &= \frac{1}{2} \frac{Y^2 \alpha}{B^2} \left( \frac{w_1^* - w_0^*}{2} \right)^2 \left( 1 - \left( \Pr_0\left( \hat{w} \ge \frac{w_0^* + w_1^*}{2} \right) - \Pr_1\left( \hat{w} \ge \frac{w_0^* + w_1^*}{2} \right) \right) \right) \\ &\geq \frac{1}{2} \frac{Y^2 \alpha}{B^2} \left( \frac{w_1^* - w_0^*}{2} \right)^2 \left( 1 - \left| \Pr_0\left( \hat{w} \ge \frac{w_0^* + w_1^*}{2} \right) - \Pr_1\left( \hat{w} \ge \frac{w_0^* + w_1^*}{2} \right) \right| \right). \end{split}$$

By Pinsker's inequality, since  $\hat{w}$  is a deterministic function of the training set S, this is at least

$$\frac{1}{8} \frac{Y^2 \alpha}{B^2} \left( w_1^* - w_0^* \right)^2 \left( 1 - \sqrt{\frac{1}{2} D_{kl}(\Pr_0(S) || \Pr_1(S))} \right),$$

where  $D_{kl}$  is the Kullback-Leibler divergence. Since S is composed of m i.i.d. instances, and the target value y is fixed under both distributions, we can invoke the chain rule and rewrite this as

$$\frac{1}{8} \frac{Y^2 \alpha}{B^2} \left( w_1^* - w_0^* \right)^2 \left( 1 - \sqrt{\frac{m}{2} D_{kl}(\Pr_0(x)||\Pr_1(x))} \right),\,$$

To simplify the bound, note that the Kullback-Leibler divergence between two distributions p,q can be upper bounded by their  $\chi^2$  divergence, which equals  $\sum_a \frac{(p(a)-q(a))^2}{q(a)}$ . Therefore,

$$D_{kl}(\Pr_0(x)||\Pr_1(x)) \le \frac{\gamma^2}{\gamma} + \frac{\gamma^2}{\alpha - \gamma} = \gamma \left(1 + \frac{\gamma}{\alpha - \gamma}\right).$$

Plugging this back, as well as the value of  $(w_1^* - w_0^*)^2$  from Eq. (1), we get an excess loss lower bound on the form

$$\frac{1}{8}Y^2\alpha B^2\gamma^2\left(\frac{Y-B}{Y^2\alpha+B^2\gamma}\right)^2\left(1-\sqrt{\frac{m}{2}\gamma\left(1+\frac{\gamma}{\alpha-\gamma}\right)}\right),$$

We now consider two cases:

• If  $m \le B^2/Y^2$ , we pick  $\alpha = 1$  and  $\gamma = 1/3m$ , and get that the expression above is at least

$$\begin{split} &\frac{Y^2}{72} \frac{B^2}{m^2} \left( \frac{B - Y}{Y^2 + B^2/3m} \right)^2 \left( 1 - \sqrt{\frac{1}{6} \left( 1 + \frac{1/3m}{1 - 1/3m} \right)} \right) \\ &= \frac{Y^2}{72} \left( \frac{B(B - Y)}{mY^2 + B^2/3} \right)^2 \left( 1 - \sqrt{\frac{1}{6} \left( 1 + \frac{1}{3m - 1} \right)} \right) \\ &\geq \frac{Y^2}{72} \left( \frac{B(B - Y)}{(B^2/Y^2)Y^2 + B^2/3} \right)^2 \left( 1 - \sqrt{\frac{1}{6} \left( 1 + \frac{1}{3m - 1} \right)} \right) \\ &\geq \frac{Y^2}{72} \left( \frac{B(B - Y)}{(1 + 1/3)B^2} \right)^2 \left( 1 - \sqrt{\frac{1}{6} \left( 1 + \frac{1}{2} \right)} \right) \\ &\geq 0.003 \, Y^2 \left( \frac{B - Y}{B} \right)^2 = 0.003 \, Y^2 \left( 1 - \frac{Y}{B} \right)^2 \geq 0.003 \, Y^2 \left( 1 - \frac{1}{2} \right)^2, \end{split}$$

where we used the assumption that  $B \geq 2Y$ .

• If  $m > B^2/Y^2$ , we pick  $\alpha = B^2/(Y^2m)$  and  $\gamma = 1/3m$  and get that the expression above is at least

$$\frac{1}{8} \frac{B^4}{m} \frac{1}{9m^2} \left( \frac{B-Y}{B^2/m + B^2/3m} \right)^2 \left( 1 - \sqrt{\frac{1}{6}} \left( 1 + \frac{1/3m}{(B^2/Y^2 - 1/3)/m} \right) \right) 
\ge \frac{1}{72} \frac{(B-Y)^2}{m(1+1/3)^2} \left( 1 - \sqrt{\frac{1}{6}} \left( 1 + \frac{1/3}{4-1/3} \right) \right) 
\ge 0.004 \frac{(B-Y)^2}{m} \ge 0.004 \frac{(B-B/2)^2}{m} = 0.001 \frac{B^2}{m},$$

where we used the assumption that  $B \geq 2Y$ .

Combining the two cases, we get an excess risk lower bound of  $c \min \left\{ Y^2, \frac{B^2}{m} \right\}$  for some universal constant c.

Our second construction provides a different type of bound, which quantifies a dependence on the dimension d. The construction is similar to standard dimension-dependent lower bounds for learning with the squared loss, but we are careful to analyze the dependence on all relevant parameters.

**Theorem 3.** There exists a universal constant c, such that for any dimension d, sample size m, target value bound Y, predictor norm bound B and any algorithm returning a linear predictor  $\hat{\mathbf{w}}$ , there exists a data distribution in d dimensions such that

$$\mathbb{E}[R(\hat{\mathbf{w}}) - R(\mathbf{w}^*)] \ge c \min\left\{Y^2, B^2, \frac{dY^2}{m}, \frac{BY}{\sqrt{m}}\right\}.$$

*Proof.* By Yao's minimax principle, it is sufficient to display a randomized choice of data distributions, with respect to which the expected excess error of any deterministic algorithm attains the lower bound in the theorem.

In particular, fix some  $d' \le d$  to be chosen later, let  $\sigma \in \{-1, +1\}^{d'}$  be chosen uniformly at random, and consider the distribution  $\mathcal{D}_{\sigma}$  (indexed by  $\sigma$ ) over examples  $(\mathbf{x}, y)$ , defined as follows:  $\mathbf{x}$  is chosen uniformly at random among

the first d' standard basis vectors, and y = Y with probability  $\frac{1}{2}(1 + \sigma_i b)$ , where  $b = \min\{1/2, \sqrt{d'/6m}\}$ , and y = -Y otherwise.

A simple calculation shows that the optimum  $\mathbf{w}^* = \arg\min_{\mathbf{w}: \|\mathbf{w}\| \leq B} R(\mathbf{w})$  is such that

$$\forall i, \quad w_i^* = \sigma_i \min\{Yb, B/\sqrt{d}\}.$$

Therefore, using Lemma 1 and the notation  $1_A$  to be the indicator function for the event A:

$$\mathbb{E}\left[R(\hat{\mathbf{w}}) - R(\mathbf{w}^*)\right] = \mathbb{E}\left[\left(\langle \hat{\mathbf{w}}, \mathbf{x} \rangle - \langle \mathbf{w}^*, \mathbf{x} \rangle\right)^2\right]$$

$$= \mathbb{E}\left[\frac{1}{d'} \sum_{i=1}^{d'} (\hat{\mathbf{w}}_i - \mathbf{w}_i^*)^2\right]$$

$$= \frac{1}{d'} \sum_{i=1}^{d'} \mathbb{E}\left[(\hat{\mathbf{w}}_i - \mathbf{w}_i^*)^2\right]$$

$$\geq \frac{1}{d'} \sum_{i=1}^{d'} \mathbb{E}\left[(\mathbf{w}_i^*)^2 \mathbf{1}_{\hat{\mathbf{w}}_i \mathbf{w}_i^* \leq 0}\right]$$

$$= \frac{1}{d'} \left(\min\{Yb, B/\sqrt{d'}\}\right)^2 \sum_{i=1}^{d'} \Pr(\hat{\mathbf{w}}_i \mathbf{w}_i^* \leq 0).$$

Since  $\sigma_i$  is uniformly distributed on  $\{-1, +1\}$ , and has the same sign as  $w_i^*$ , this equals

$$\frac{1}{d'} \left( \min\{Yb, B/\sqrt{d'}\} \right)^{2} \sum_{i=1}^{d'} \frac{1}{2} \left( \Pr(\hat{\mathbf{w}}_{i} \ge 0 | \sigma_{i} < 0) + \Pr(\hat{\mathbf{w}}_{i} \le 0 | \sigma_{i} > 0) \right) \\
\ge \frac{1}{2d'} \left( \min\{Yb, B/\sqrt{d'}\} \right)^{2} \sum_{i=1}^{d'} \left( 1 - \Pr(\hat{\mathbf{w}}_{i} \le 0 | \sigma_{i} < 0) + \Pr(\hat{\mathbf{w}}_{i} \le 0 | \sigma_{i} > 0) \right) \\
\ge \frac{1}{2d'} \left( \min\{Yb, B/\sqrt{d'}\} \right)^{2} \sum_{i=1}^{d'} \left( 1 - |\Pr(\hat{\mathbf{w}}_{i} \le 0 | \sigma_{i} < 0) - \Pr(\hat{\mathbf{w}}_{i} \le 0 | \sigma_{i} > 0) | \right)$$

Using Pinsker's inequality and the fact that  $\hat{\mathbf{w}}$  is a deterministic function of the training set S, this is at least

$$\frac{1}{2d'} \left( \min\{Yb, B/\sqrt{d'}\} \right)^2 \sum_{i=1}^{d'} \left( 1 - \sqrt{\frac{1}{2} D_{kl} \left( \Pr(S|\sigma_i < 0) || \Pr(S|\sigma_i > 0) \right)} \right), \tag{2}$$

where  $D_{kl}$  is the Kullback-Leibler (KL) divergence. Since the training set is composed of m i.i.d. instances, we can use the chain rule and get that this divergence equals  $mD_{kl}\left(\Pr((\mathbf{x},y)|\sigma_i<0)||\Pr((\mathbf{x},y)|\sigma_i>0)\right)$ . Moreover, we note that

$$Pr((\mathbf{x}, y)|\sigma_i) = Pr(\mathbf{x} = \mathbf{e}_i) Pr((\mathbf{x}, y)|\sigma_i, \mathbf{x}_i = \mathbf{e}_i) + Pr(\mathbf{x} \neq \mathbf{e}_i) Pr((\mathbf{x}, y)|\sigma_i, \mathbf{x} \neq \mathbf{e}_i)$$

$$= \frac{1}{d'} Pr((\mathbf{x}, y)|\sigma_i, \mathbf{x} = \mathbf{e}_i) + \left(1 - \frac{1}{d'}\right) Pr((\mathbf{x}, y)|\sigma_i, \mathbf{x} \neq \mathbf{e}_i),$$

and therefore, by joint convexity of the KL-divergence, we get

$$D_{kl}(\Pr((\mathbf{x}, y)|\sigma_i > 0)||\Pr((\mathbf{x}, y)|\sigma_i < 0)) = \frac{1}{d'}D_{kl}(\Pr((\mathbf{x}, y)|\sigma_i < 0, \mathbf{x} = \mathbf{e}_i)||\Pr((\mathbf{x}, y)|\sigma_i > 0, \mathbf{x} = \mathbf{e}_i))$$

$$+ \left(1 - \frac{1}{d'}\right)D_{kl}(\Pr((\mathbf{x}, y)|\sigma_i < 0, \mathbf{x} \neq \mathbf{e}_i)||\Pr((\mathbf{x}, y)|\sigma_i > 0, \mathbf{x} \neq \mathbf{e}_i)).$$

Since the distribution of y is independent of  $\sigma_i$ , conditioned on  $\mathbf{x} \neq \mathbf{e}_i$ , this equals

$$\frac{1}{d'}D_{kl}\left(\Pr(y|\sigma_i > 0, \mathbf{x} = \mathbf{e}_i)||\Pr(y|\sigma_i < 0, \mathbf{x} = \mathbf{e}_i)\right). \tag{3}$$

The divergence in this equation is simply the KL divergence between two Bernoulli random variables, one with parameter  $\frac{1}{2}(1+b)$ , and the other with parameter  $\frac{1}{2}(1-b)$ . To get a simple upper bound, note that the KL divergence between two distributions p,q can be upper bounded by their  $\chi^2$  divergence, which equals  $\sum_a \frac{(p(a)-q(a))^2}{q(a)}$ . Therefore, we can upper bound Eq. (3) by

$$\frac{b^2}{d'}\left(\frac{1}{\frac{1}{2}(1+b)} + \frac{1}{\frac{1}{2}(1-b)}\right) = \frac{2b^2}{d'}\left(\frac{1}{1+b} + \frac{1}{1-b}\right) \leq \frac{2b^2}{d'}\left(1 + \frac{1}{1/2}\right) = \frac{6b^2}{d'},$$

where we used the fact that  $b \in [0, 1/2]$ . Summarizing the discussion so far, we showed that

$$D_{kl}\left(\Pr(S|\sigma_i < 0)||\Pr(S|\sigma_i > 0)\right) = m D_{kl}\left(\Pr((\mathbf{x}, y)|\sigma_i < 0)||\Pr((\mathbf{x}, y)|\sigma_i > 0)\right) = \frac{6mb^2}{d'}.$$

Plugging this back into Eq. (2), we get that the excess risk is lower bounded by

$$\begin{split} \frac{1}{2d'} \left( \min\{Yb, B/\sqrt{d'}\} \right)^2 \sum_{i=1}^{d'} \left( 1 - \sqrt{\frac{3mb^2}{d'}} \right) &= \left( \min\{Yb, B/\sqrt{d'}\} \right)^2 \frac{1}{2} \left( 1 - \sqrt{\frac{3mb^2}{d'}} \right) \\ &\geq \left( \min\{Yb, B/\sqrt{d'}\} \right)^2 \frac{1}{2} \left( 1 - \sqrt{\frac{3m(d'/6m)}{d'}} \right) \\ &\geq 0.14 \left( \min\{Yb, B/\sqrt{d'}\} \right)^2 \\ &= 0.14 \left( \min\left\{ Y \min\left\{ \frac{1}{2}, \sqrt{\frac{d'}{6m}} \right\}, \frac{B}{\sqrt{d'}} \right\} \right)^2 \\ &= 0.14 \min\left\{ \frac{1}{4}Y^2, \frac{d'Y^2}{6m}, \frac{B^2}{d'} \right\}. \end{split}$$

Now, recall that d' is a free parameter of value at most d. We now distinguish between two cases:

• If  $d > \sqrt{6m}B/Y$ , then we pick  $d' = \lceil \sqrt{6m}B/Y \rceil$ , and get that the expression above is at least

$$0.14 \min \left\{ \frac{1}{4} Y^2, \frac{B^2}{d'} \right\} \ \geq \ 0.14 \min \left\{ \frac{1}{4} Y^2, \frac{B^2}{\max \left\{ 1, 2 \sqrt{6m} \frac{B}{Y} \right\}} \right\} \ = \ 0.14 \min \left\{ \frac{1}{4} Y^2, B^2, \frac{BY}{2\sqrt{6m}} \right\}.$$

• If  $d \le \sqrt{6m}B/Y$ , we pick d' = d, and note that  $\frac{d'Y^2}{6m} \le \frac{B^2}{d}$  in this case. Therefore, the expression above is at least

$$0.14\min\left\{\frac{1}{4}Y^2,\frac{dY^2}{6m}\right\}$$

Combining the two cases, we get that a lower bound of the form

$$c \min \left\{ Y^2, B^2, \frac{dY^2}{m}, \frac{BY}{\sqrt{m}} \right\},$$

where c is a universal constant.

With Thm. 2 and Thm. 3 at hand, we now turn to prove our main result:

*Proof of Thm. 1.* Taking the maximum of Thm. 2 and Thm. 3, and using the fact that  $B \ge 2Y$ , we get a lower bound of

$$c \max \left\{ \min \left\{ Y^2, \frac{B^2}{m} \right\} \;,\; \min \left\{ Y^2, \frac{dY^2}{m}, \frac{BY}{\sqrt{m}} \right\} \right\}$$

for some constant c. If  $m \leq (B^2/Y^2)$ , this is at least  $Y^2$ , and otherwise it is

$$c\max\left\{\frac{B^2}{m}\;,\;\min\left\{\frac{dY^2}{m},\frac{BY}{\sqrt{m}}\right\}\right\}\;\geq\;\frac{c}{2}\left(\frac{B^2}{m}+\min\left\{\frac{dY^2}{m},\frac{BY}{\sqrt{m}}\right\}\right)\;\geq\;\frac{c}{2}\min\left\{\frac{B^2+dY^2}{m},\frac{BY}{\sqrt{m}}\right\}.$$

Combining the two cases, the result follows.

Acknowledgements: We thank Nati Srebro for helpful comments.

#### References

- [1] M. Anthony and P. Bartlett. *Neural network learning: Theoretical foundations*. Cambridge University Press, 1999.
- [2] K. Azoury and M. Warmuth. Relative loss bounds for on-line density estimation with the exponential family of distributions. *Machine Learning*, 43(3):211–246, 2001.
- [3] N. Cesa-Bianchi, A. Conconi, and C. Gentile. On the generalization ability of on-line learning algorithms. *Information Theory, IEEE Transactions on*, 50(9):2050–2057, 2004.
- [4] D. Hsu, S. M Kakade, and T. Zhang. Random design analysis of ridge regression. *Foundations of Computational Mathematics*, 14(3):569–600, 2014.
- [5] V. Koltchinskii. *Oracle Inequalities in Empirical Risk Minimization and Sparse Recovery Problems: Ecole dEté de Probabilités de Saint-Flour XXXVIII-2008*, volume 2033. Springer, 2011.
- [6] G. Lecué and S. Mendelson. Performance of empirical risk minimization in linear aggregation. *arXiv preprint arXiv:1402.5763*, 2014.
- [7] W. Lee, P. Bartlett, and R. Williamson. The importance of convexity in learning with squared loss. *Information Theory, IEEE Transactions on*, 44(5):1974–1980, 1998.
- [8] A. Singer, S. Kozat, and M. Feder. Universal linear least squares prediction: upper and lower bounds. *Information Theory, IEEE Transactions on*, 48(8):2354–2362, 2002.
- [9] N. Srebro, K. Sridharan, and A. Tewari. Smoothness, low noise and fast rates. In NIPS, pages 2199–2207, 2010.
- [10] A. Tsybakov. Optimal rates of aggregation. In *Learning Theory and Kernel Machines*, pages 303–313. Springer, 2003.
- [11] V. Vovk. Competitive on-line statistics. *International Statistical Review*, 69(2):213–248, 2001.