Problem solving in ID-logic with aggregates: some experiments

Bert Van Nuffelen and **Marc Denecker**

K.U.Leuven, Department of Computer Science
Celestijnenlaan 200A, 3001 Heverlee, Belgium
{marcd,bertv}@cs.kuleuven.ac.be

Abstract

The goal of the LP+ project at the K.U.Leuven is to design an expressive logic, suitable for declarative knowledge representation, and to develop intelligent systems based on Logic Programming technology for solving computational problems using the declarative specifications. The ID-logic is an integration of typed classical logic and a definition logic. Different abductive solvers for this language are being developed. This paper is a report of the integration of high order aggregates into ID-logic and the consequences on the solver SLDNFA.

Introduction

The goal of computational logic is to design logics for *knowledge representation* and to develop algorithms to solve computational problems using the *declarative specifications*. In principle, the declarative knowledge representation methodology in logic is based on a simple idea. To describe his knowledge, an expert designs the ontology of his problem domain: he defines the relevant types of objects, and the relevant relations and functions between these objects and chooses a logical alphabet to name them. In the next phase, the expert uses this alphabet to express his knowledge by a set of logical sentences that are *true* statements on the problem domain.

In the choice of the alphabet, the expert may be led by different and often non-compatible quality criteria: naturality of the representation, expressivity of the logic, efficiency of the representation, etc.. However, if the goal is to obtain a clear declarative representation, the expert will choose the alphabet as close as possible to what he views as the relevant objects, concepts and relationships in the problem domain. Alphabets matching more closely the natural ontology lead to more natural representations.

Choosing the alphabet in accordance with the natural ontology of the problem domain imposes high requirements both on the expressivity and on the problem solving capabilities of the used logic. With respect to the expressivity, higher order aggregates easily show up in declarative representations of practical problem domains. For example in the domain of university lecture scheduling one of the constraints will be that if a lecture l of course c takes place in a room r with a capacity of rc students, then the number n of students enrolled for course c must be less than rc. In a suitable logic, the constraint can be represented using the cardinality aggregate:

$$\forall\, c, l, r, rc, n.\ room\_of\_lecture(l, r) \wedge capacity(r, rc)$$
$$\wedge course(l, c) \wedge Card(\{st | enrolled(st, c)\}, n) \rightarrow rc \geq n$$

Other frequently occurring aggregates are *summation*, *minimum*, *maximum*, etc ... Therefore aggregates are extensively studied in logic programming and deductive databases (Kemp & Stuckey 1991; Van Gelder 1992; Ross & Sagiv 1992).

A natural choice of the alphabet also poses requirements on the level of problem solving capabilities of the used logic. Illustrated in the domain of university lecture scheduling, some basic concepts are *lectures*, *time slots*, *rooms*. Basic relations between these concepts describe when and where lectures take place. The natural choice to represent these relations are by (typed) predicates (e.g. *time_of_lecture*/2 and *room_of_lecture*/2). Now, consider the task of computing a schedule satisfying certain data and constraints (lectures to be given, rooms and lecturers available, no overlap, etc ...). This is a satisfiability problem or an abductive problem; tables of these predicates must be computed that satisfy the data and constraints imposed on correct schedules. Note that at this point, we have only chosen the alphabet; we have not formalized one single iota of knowledge in a logical formula. This shows that often, satisfiability checking or abduction is a natural companion and a natural consequence of applying the declarative knowledge representation methodology.

At the K.U.Leuven, the aim of the LP+ project is to use both semantical and implementational techniques of logic programming to develop a logic suitable for declarative knowledge representation and to implement efficient problem solvers for such a logic. In this paper, we study satisfiability problems in the context of a logic with higher order aggregates. The used logic is an extension of inductive definition logic (ID-logic) (Denecker 2000; Denecker 1998; Denecker 1995) developed within the project. ID-logic is a conservative extension

of classical logic with a generalized notion of non monotone inductive definitions. At the same time, the logic can be seen as a natural generalization of Abductive Logic Programming (Kakas, Kowalski, & Toni 1993) and Open Logic Programming (Denecker 1995).

The technology that is used so far for reasoning and problem solving in ID-logic is based on an integration of techniques of abductive logic programming (Kakas, Kowalski, & Toni 1993) and constraint logic programming techniques (Jaffar & Maher 1994). The first published presentation of such a procedure was in (Kakas & Michael 1995; Kakas, Michael, & Mourlas 2000). (Denecker & Van Nuffelen 1999) describes a similar integration of the abductive resolution method SLDNFA (Denecker & De Schreye 1998) with CLP techniques and presents some experiments in the context of ID-logic specifications of some typical CLP-problems (such as N-queens), and of scheduling and planning.

In this paper, we define an extension of ID-logic with aggregates, sketch an extension of the solver of (Denecker & Van Nuffelen 1999) to reason on aggregates and describe two computational experiments with the system for solving a sea-battle puzzle and a scheduling problem. At the logical level, our work is based on Van Gelders's work (Van Gelder 1992) on aggregates in the context of well-founded semantics. On the other hand the implementation and experiments extend previous work described in (Denecker *et al.* 1997; Seghers & Baeyens 1996). To the best of our knowledge, this paper is the first report on solving satisfiability and abductive problems in the context of (an extension of) classical logic with higher order aggregates. All the studies concerning aggregation we encountered, were in the context of querying systems.

Because an abductive system for ID-logic cannot be complete (computing satisfiability of a classical first order theory is a undecidable subproblem of the computation of an abductive solution of an ID-logic theory.), the current implementation is only valid for a restricted class of problems. But the experiments show the feasibility of reasoning on a useful class of specifications with aggregates.

ID-logic extended with aggregates

ID-logic

As mentioned, ID-logic is an extension of classical first order logic (FOL) with inductive definitions. The logic builds upon the earliest ideas on the declarative semantics of logic programs with negation as failure. The view of a logic program as a definition of its predicates is underlying both the least model semantics of van Emden and Kowalski (van Emden & Kowalski 1976) and Clark's completion semantics (Clark 1978). This idea was further explored in (Denecker, Marek, & Truszczynski 1998), where it was argued that the well-founded semantics for logic programming implements a generalized principle of non-monotone induction. A discussion of this is out of the scope of this paper.

An ID-logic theory T consists of a set of *definitions* and a set of classical logic sentences. A definition is an expression that *defines* a subset of predicates in terms of the other predicates. Formally, a definition D is a pair of a set $Defined(D)$ of predicates and a set $Rules(D)$ of rules that exhaustively enumerate the cases in which the predicates of $Defined(D)$ are true. A rule is of the form:

$$p(\bar{t}) \leftarrow F$$

where $p \in Defined(D)$ and F an arbitrary first order formula. The predicates in $Defined(D)$ are called *defined* by D, the others (not in $Defined(D)$) are called *open* in D.

The semantics of ID-logic integrates classical logic semantics and well-founded semantics. An interpretation M is a model of a definition D iff it is total (i.e. 2-valued) and the unique well-founded model of D extending some interpretation M_o of the functor and open predicate symbols of D. An interpretation M is a model of an ID-logic theory T iff it is a total model of its classical logic sentences and of its definitions.

Logical entailment is defined as usual: $T \models F$ iff F is true in all models of T.

ID-logic generalizes not only classical logic but also abductive logic programming (Kakas, Kowalski, & Toni 1993) and open logic programming (Denecker 1995). An abductive logic framework, consisting of a set of abducible predicates, a set of rules and a set of FOL *constraints* can be embedded in ID-logic as the theory consisting of the FOL constraints and one definition defining all non-abducible predicates. Formally, ID-logic extends ALP by allowing multiple definitions and generalized syntax. However, it can be shown that it is always possible to transform a set of definitions into one single definition. As will appear in the next section, current problem solvers for ID-logic are based on the technology of abductive logic programming.

In the current version of the SLDNFA-system a simple many-sorted version of ID-logic is used. A type inference system checks and completes a partial set of user-defined type declarations: this produces a small overhead compared to the benefits it brings: e.g. error checking, disambiguating expressions, etc. For more details, we refer to (De Mot *et al.* 1999). In the implementation, a Prolog-like style using capitals for variables, and ",", resp. ";" for conjunction, resp. disjunction is used. For example a definition is represented as

```
uncle(X,Y)<- ( exists(Z): parent(Y,Z),brother(X,Z);
               exists(A): aunt(A,Y),married(A,X) ).

aunt(X,Y) <- ( exists(Z): parent(Y,Z),sister(X,Z);
               exists(A): uncle(A,Y),married(A,X)).
```

This definition defines the two predicates *uncle* and *aunt* simultaneously. The other are open predicates.

FOL axioms are represented in the system in the same style but are prefixed by the key-word `fol`. Examples are:

```
fol forall(X,Y):
```

```
    uncle(X,Y), age(X,AgeX), ageY(Y,AgeY)
    => AgeX > AgeY.
fol aunt(mary,bob).
```

A special case. The definitions that appear in the experiments in this paper are of a simple kind; they do not contain recursion. The models of a definition without recursion are exactly the models of the completion of the definition (Clark 1978). Below, completed definitions of predicates will be denoted:

$$\forall(p(\overline{X}) \leftrightarrow B_p[\overline{X}])$$

We call from now on $B_p[\overline{X}]$ the *completion* of p.

Aggregates

As mentioned in the introduction, aggregates have been studied in logic programming and deductive databases. (Kemp & Stuckey 1991; Van Gelder 1992; Ross & Sagiv 1992) proposed extensions of logic programming with aggregates and showed how aggregate expressions can be transformed and reduced to recursive logic programs using certain schemata. These transformations define at the same time the declarative and procedural semantics for these extensions. The papers cover actually a subclass of ID-logic theories, because they don't have the notion of open predicates. In the context of ID-logic, where we deal with open predicates, the proposed schemata (inductive definitions) can be used to define the declarative semantics without extending ID-logic, but the current implementation of the solver cannot cope with the reduced programs due to their highly recursive nature. If the current implementation would be extended with a proper notion of tabling these definitions probably could be computed. Therefore we take another approach by introducing new language primitives, which allows us to use aggregation without the computation of recursive definitions.

We define first two basic concepts of our aggregate expressions.

- A **Set expression** is an expression of the form:

$$\{\overline{x}|F[\overline{x}]\}$$

where $\overline{x}$ is a tuple of variables and $F[\overline{x}]$ a first order formula. It denotes the set of tuples $\overline{x}$ that satisfy $F[\overline{x}]$. Three sorts of variables are distinguished in a set expression: local variables are quantified inside F; their scope is the quantifier. Parameter variables are the elements of $\overline{x}$; their scope is the set expression. The other variables are free variables.

In the following example the set expression denotes for a person Y the set of all aunts of Y which are older than 50.

```
set([X],(exists(AgeX):
         aunt(X,Y),age(X,AgeX), AgeX >50)))
```

- A **function expression** is of the form:

$$\lambda\overline{x}.y \ where \ F[\overline{x}, y]$$

where $\overline{x}$ is a tuple of variables, y is a variable and $F[\overline{x}, y]$ is a first order formula expressing a functional

relation between y and $\overline{x}$. I.e. it satisfies the constraint:

$$\forall\overline{x}\exists!y.F[\overline{x}, y]$$

This expression denotes the anonymous function mapping a tuple $\overline{x}$ to the object y for which $F[\overline{x}, y]$ is satisfied. As above, we distinguish between local variables, parameter variables and free variables of a function expression.

A pure arithmetical function

$$\lambda\overline{x}.t[\overline{x}]$$

where $t[\overline{x}]$ is an arithmetical term, is a shorthand for

$$\lambda\overline{x}.y \ where \ y = t[\overline{x}]$$

The next example represents the function which maps a person X to his age.

```
lambda([X] , Y where (exists(Z) : age(X,Z), Y=Z))
```

Set and function expressions are allowed to appear only in aggregate expressions. The following aggregate primitives have been implemented:

- minimum: $minimum(\Psi, n)$ means that n is the minimal element of the set denoted by Ψ. In the current implementation, the set expression must be of type integer.

- cardinality: $card(\Psi, n)$ means that the set represented by Ψ has n elements.

- summation: $sum/3$ has as arguments an n-ary set expression, an n-ary function expression of a number type, and a number as argument. $sum(\Psi, f, s)$ means that:

$$s = \sum_{\overline{x}\in\Psi} f(\overline{x})$$

- product: $product/3$ is analogous to $sum/3$ but expresses the product of a function over a set.

Semantics

In this section, we briefly explore how the semantics of ID-logic with aggregates can be defined. The semantics of ID-logic can be extended using the same transformational approach proposed by Van Gelder in (Van Gelder 1992). In this approach, aggregate expressions are transformed to recursive logic programs under well-founded semantics. Because the models of an ID-logic theory are well-founded models, the same approach applies.

In the following example we illustrate the approach of (Van Gelder 1992) in the case of the minimum-aggregate and show how it can be transformed into a definition. We illustrate the transformation in the context of open predicates. Consider the following theory T where a is an open predicate ranging over integers:

$$T = \{min(\{X|a(X)\}, 4)\}$$

This theory expresses that the minimum of the argument of the predicate a should be 4. Applying the transformation scheme in (Van Gelder 1992), we obtain:

```
fol mina(4).

mina(X)    <- a(X), not bettera(X).
bettera(X) <- a(Y), Y<X.
```

In this simple case, we obtain a non-recursive definition for the new predicates *mina* and *bettera*. Abductive solvers can be used to successfully reason on these theories. For example

- The query[1]*true* will lead to a successful derivation generating the abductive answer $\Delta = \{a(4)\}$. To solve the FOL constraint $mina(4)$, a(4) must be abduced.

- Likewise, the query $a(6)$ will succeed with abductive answer $\Delta = \{a(4), a(6)\}$.

- The query $a(1)$ will fail. During the derivation, the solver will abduce $a(1)$ and will attempt to solve $mina(4)$. This is impossible because $bettera(4)$ can be derived.

In the case of summation and cardinality, application of the schemata results in highly recursive logic programs. As shown in (Van Gelder 1992), the transformational approach poses no problem at the level of semantics; however SLDNFA loops on such programs. For this reason, we implemented aggregates in a different (more efficient) way (see further on).

An abductive problem solver

From here on, we will assume that an ID-logic theory T contains only one definition defining a number of predicates simultaneously. Recall from previous section that it is always possible to transform a set of definitions into this form.

Given an ID-logic theory T with definition D, an abductive problem for a given query Q consists of computing a set Δ of definitions of ground atoms for each open predicate of T and an answer substitution θ such that $D + \Delta$ is consistent and entails all FOL axioms in T and $\forall(\theta(Q))$. An abductive procedure computes tables for the open predicates that can be extended in a unique way to a well-founded model of the definition and a model of the FOL axioms and of the query. This way, an abductive answer can be seen as a compact space-efficient representation of a model.

SLDNFA (Denecker & De Schreye 1992; Denecker & De Schreye 1998) is an abductive procedure for normal logic programs. (Denecker & Van Nuffelen 1999) describes an extension of this procedure to deal with FOL axioms and generalized rules and queries, and describes an integration of this procedure with constraint solvers. This integration is in a similar spirit as ACLP (Kakas & Michael 1995; Kakas, Michael, & Mourlas 2000). Due to lack of space, we can only give the head-lines of this procedure. For more detailed description, we refer to (Denecker & Van Nuffelen 1999; Van Nuffelen 2000).

---

[1]Note that in contrast to Logic Programming conventions, queries are not denoted by denials.

A derivation for a query Q can be understood as a rewriting process of *states* $\mathcal{S}$, i.e. tuples $(\Theta, \Delta, \mathcal{CS})$ of a set Θ of FOL formulas and denials, a set Δ of abduced open atoms and a set $\mathcal{CS}$ of CLP expressions, called the constraint store. A denial is a formula of the form $\forall \overline{X}. \leftarrow F[\overline{X}, \overline{Y}]$, where $\leftarrow$ denotes negation. Denials are the only formulas that may contain universal quantifiers. Free variables in FOL formulas and denials represent objects of yet unknown identity.

The derivation starts with the initial state $(\Theta, \emptyset, \emptyset)$ where Θ consists of Q and the set of FOL axioms in T. The rewriting process proceeds by selecting an atom in a formula F from Θ and computing a new state depending on the sort of atom by applying the right rule. E.g. if F is an open atom, the atom is abduced; an atom interpreted in a CLP domain is added to the constraint store; defined predicates are substituted by their completion and the resulting formula is then simplified. Disjunctive goals are dealt with selecting alternative disjuncts using backtracking. Consistency of denials with a selected open atom is checked by matching this open atom against each abduced open atom in Δ. If $fail$ is derived, the computation backtracks. The computation ends in three possible ways:

- with a *floundering* error condition when universally quantified variables appear in the selected atom in a denial.

- with *failure*, if no solution is derived;

- with a *successful* derivation if a state $\mathcal{S}$ is derived where $\mathcal{CS}$ is a consistent constraint store, Δ a set of ground open atoms and Θ consists purely of denials that have been checked to be consistent with Δ.

Extending the implementation for aggregates

We extended SLDNFA in a heuristic manner to reason on aggregates. If an aggregate expression is selected during the derivation, the set expression is rewritten using the completion of the defined predicates and the table of abduced atoms. This process leads to a big disjunction enumerating potential values occurring in the set together with a CLP-constraint formula describing the logical conditions under which the potential value effectively belongs to the set.

When during the evaluation of a set expression open predicates are encountered, this partially evaluated expression is remembered and each time an atom is abduced later on, the procedure will check if it supports a new potential element of the set. Hence, abduction leads to new disjuncts in the set description.

The unfolded disjunction of potential elements and associated constraints can be used then to compute the value of the aggregate expressions. For example, assume that for some cardinality expression $Card(\{\overline{x}|F[\overline{x}]\}, N)$, $F[\overline{x}]$ could be reduced to a disjunction

$$\overline{x} = \overline{v_1} \wedge C_1 \vee .. \vee \overline{x} = \overline{v_n} \wedge C_n$$

in which v_i are distinct potential values and C_i is the associated constraint of v_i. In this case the value of N can be simply defined by the boolean sum

$$N = B_1 + .. + B_n$$

where B_i is a boolean variable defined by the constraint $B_i \iff C_i$. These kind of constraints are known as reified constraints (Carlsson, Ottosson, & Carlson). The above sum can be efficient computed using specialized library constraints of the finite domain constraint solver.

The same principle can be applied in the case of minimum, maximum and summation and product. For example, in the case of summation, each time a new potential value v_i is derived for the set, the function value f_i of this new potential element must be computed and the sum of the expression is computed as the sum of constraint expressions $B_i \times f_i$.

As can be seen above, the current implementation is strongly focussed on constructing a finite domain constraint store. This restricts, at this moment, the set expressions to have the property that all the variables which value is unknown during the unfolding of the set expression and which have influence on the membership of a value in the set should be finite domain variables. However this give us still the ability to reason on a large group of applications.

Another restriction stems from the fact that the procedure will only observe the state and maintain a complete disjunction w.r.t. Δ during the evaluation of aggregate expressions. The evaluation will not procedure new abduced atoms. In general, it is easy to find applications axioms containing aggregate expressions should lead to new abductions. For example, consider the following theory.

```
fol Card(set([X],a(X)),3).
```

To evaluate the query *true* with respect to this theory, it is necessary to abduce 3 *a*-atoms. The current implementation fails on this query.

It is a topic of future research to extend the current solver to deal with a broader class of problems. However, for an important class of practical applications, the solver works already fine. Namely when the open predicates appearing in the set expressions represent functions on some finite domain. In this case the solver will ultimately compute a complete table of abduced atoms of these predicates; consequently, the disjunctive representation of the set expressions will be complete as well. As illustrated in the experiments, many problems satisfy this condition.

Optimization functions

Often, the expert is not interested in an arbitrary solution of an abductive problem but in an optimal solution along some optimality criterion. In general it is an intractable problem. But in the context of Constraint Programming, one often recurs to the following pragmatic solution. Given a constraint program and an optimization function, an initial solution is computed and the value of the optimization function for this solution is recorded; then the system backtracks and tries to find other solutions; the value of the optimization function of the best solution so far, is used to prune the search. If the search stops, then it stops with an optimal solution; otherwise, the user may stop the system and extract the best solution so far. In a lot of cases this is satisfying for the user.

We extended the abductive solver with a similar facility. Together with the query, the user can specify an optimization function to be minimized or maximized. This is done by specifying either `minimize(V)` and `maximize(V)`, in which V is an expression which should be minimized or maximized. In practice, because the search can take a long time (or does not end), the best solution that can be computed within a given time is returned.

Experiments

We present two experiments: the first one is a solitaire puzzle based on the well-known battleship game. The second one is about scheduling the maintenances of units in power plants. The experimental results are obtained using an implementation of SLDNFA as a meta program on Sicstus Prolog 3.7 on a Sun Solaris machine.

The battleship puzzle

The objective of this puzzle (Gordon & Shenk 1998) is to find the locations of 10 ships hidden on a 10 by 10 board. There are ships from different sizes: one battleship, two cruisers, three destroyers and four submarines. The ships can be placed everywhere on the grid either horizontally either vertically oriented. They are not allowed to touch each other, therefore a ship is always surrounded by water (or the border of the grid). The data consist of a given set of known locations of boat pieces or water and the number of boat pieces on each row and column.

The formalization of this puzzle in ID-logic starts with the choice of the alphabet. The central concept in this puzzle is the location of a ship. There are two options: either one represents the location by one (X,Y) coordinate (e.g. the left upper one), the length and the orientation of the ship; or one defines the location by means of the locations of the different parts of a ship. We have chosen the last option.

Let us define the battle fleet by the following facts[2]:

```
ship(S) <- S in 1..10.

ship_type(1,battleship)<- true.
ship_type(S,cruiser)   <- S in 2..3.
ship_type(S,destroyer) <- S in 4..6.
ship_type(S,submarine) <- S in 7..10.

length(battleship,4)<-true.
length(cruiser,3)   <-true.
length(destroyer,2) <-true.
```

---

[2]In the specifications we omit the type information

```
length(submarine,1) <-true.

ship_length(S,L) <- ship_type(S,Type), length(Type,L).
```
ship_length is an auxiliary definition which defines for
a particular ship its length. Then depending on the
type, ships consist of different number of parts, each
connected to a location. We represent this by the open
predicate ship(S,P,X,Y), which denotes a part P from
a ship S located at coordinates (X,Y).

```
domx(X) <- X in 1..10.
domy(Y) <- X in 1..10.

fol forall(S,Length,Part) :
    ship(S),ship_length(S,Length), Part in 1..Length
    => ( exists(X,Y) :
        domx(X),domy(Y), ship(S,Part,X,Y) ).
```
Defining a ship by the locations of its parts introduces
one specific statement namely that the parts are con-
nected to each other and not separated by water or
other boats. The next statement ensures also that the
ships are either vertically or horizontally oriented.

```
fol forall(S,P1,P2,X1,X2,Y1,Y2):
    ship(S,P1,X1,Y1), ship(S,P2,X2,Y2), P1 \= P2
    => ( X1 - X2 = P1 - P2, Y1=Y2
       ; Y1 - Y2 = P1 - P2, X1=X2).
```
Another requirement is that two different ships do not
touch each other. Translated to the above chosen rep-
resentation of a ship, this means that the distance be-
tween two parts belonging to different ships is greater
than 1.

```
fol forall(S1,S2,P1,P2,X1,X2,Y1,Y2):
    ship(S1,P1,X1,Y1), ship(S2,P2,X2,Y2), S1 \= S2
    => (abs(X1-X2) > 1 ; abs(Y1-Y2) > 1).
```

The above statements describe the general knowledge
about how ships are located. In the context of a solitaire
puzzle, the data specify for a subset of locations whether
they contain water or a boat part. We represent these
data by a set of atomic fol axioms of the form:

```
fol water(i,j).          fol boat(i,j)
```
where water(i,j) (boat(i,j)) means that on coordi-
nate (i,j) there is water (a boat piece). As they exclude
each other, we can define water as a location that is
no occupied by a boat piece. A location occupied by a
boat piece is a location which is occupied by a part of
a ship (as defined above). Formally expressed

```
water(X,Y) <- not boat(X,Y).
boat(X,Y)  <- ship(S), ship_length(S,L),
              P in 1..L, ship(S,P,X,Y).
```

Up to now there was no need for aggregates. But the
puzzle gives also the number of ship parts located on a
certain row or column. As above we specify the data
by a set of atomic fol axioms.

```
fol row(i,n).          fol column(j,m).
```
These two predicates are defined using the cardinality
constraint as:

```
row(I,N)   <- card(set([S,P],
                       (exists(Y): ship(S,P,I,Y)),N).
column(J,M) <- card(set([S,P],
                       (exists(X): ship(S,P,X,J)),M).
```

We used the above specification to solve a number of
puzzles from the book (Gordon & Shenk 1998). The
abductive solver was able to reduce the above specifi-
cation to a finite domain constraint store. This phase
took about 1.5 second and is constant for all puzzles.
The time to find a solution of the constraint store var-
ied: we obtained from 18 seconds for easy puzzles till 3
minutes for most difficult ones. This means that most
of this time is spend in enumerating candidate solutions
by the CLP solver. The time could be improved if some
special search strategy as mentioned in the book, would
be followed, but in our declarative approach it is not
easy to specify a search strategy. This is an interesting
topic for future research.

Scheduling of maintenances

The next experiment is based on a real life problem
of a Belgian electricity provider. The company has a
network of power plants, distributed over different ar-
eas and each containing several power producing units.
These units need a fixed number of maintenances during
the year and the problem is to schedule these mainte-
nances so that the risk of power shortage (and hence,
import from France) is as low as possible. This ap-
proach extends earlier work described in (Seghers &
Baeyens 1996; Denecker *et al.* 1997).

The fact that a maintenance M lasts from week B till
week E, is represented by the predicate start(M,B,E).
This is the only open predicate in the specification.
Other predicates are either defined or are input data
and are defined by a table. We will introduce now the
constraints one by one.

- Maintenances (maint(M)) and their duration
  (duration(M,D)) are given by a table. All mainte-
  nances must be scheduled, thus for each maintenance
  there exists an according start relation.

  ```
  fol forall(M) : maint(M)
      => exists(B,E,D): week(B), week(E),
      duration(M,D), (E = B + D -1), start(M,B,E).

  week(W) <- W in 1..52.
  ```

- A table of prohibited(U,Bp,Ep) facts specify that
  maintenances M for unit U are not allowed during the
  period [Bp,Ep]:

  ```
  fol forall(U,Bp,Ep,M,B,E) :
      prohibited(U,Bp,Ep), maint_for_unit(M,U),
      start(M,B,E)
      => (E < Bp ; Ep < B).
  ```

- Some of the maintenances are not allowed to overlap.
  The table of non_simult_ maint(M1,M2,Pre,Post)
  facts describes this; Post and Pre represent the min-
  imum distance between the two maintenances.

  ```
  fol forall(M1,M2,Pre,Post,B1,E1,B2,E2) :
      non_simult_maint(M1,M2,Pre,Post),
      start(M1,B1,E1), start(M2,B2,E2)
      => (B2 > E1 + Post; B1 > E2 + Pre).
  ```

- Some maintenances should be done simultaneously,
  as defined by a table of simult_ maint(M1,M2)

atoms. Two maintenances are simultaneous if the period of one is included in the period of another.

```
fol forall(M1,M2,B1,E1,B2,E2) :
    simult_maint(M1,M2), start(M1,B1,E1),
    start(M2,B2,E2)
    => ( (B1 =< B2, E2 =< E1)
       ; (B2 =< B1, E1 =< E2)).
```

- Different maintenances for the same unit should not overlap:

```
fol forall(U,M1,M2,B1,E1,B2,E2) :
    unit(U),
    maint_for_unit(M1,U), maint_for_unit(M2,U),
    M1 \= M2,start(M1,B1,E1), start(M2,B2,E2)
    => (E1 < B2; E2 < B1).
```

- For each week the number of the units in maintenance belonging to a plant P should be less than a maximal number Max. A given table of plant_max(P,Max) atoms defines for each plant the maximal number of units in maintenance simultaneously.

```
fol forall(P,Max,We) :
    plant(P), plant_max(P,Max), week(We)
    => (exists(OnMaint):
           card(set([U], (unit(U),unit_in_plant(U,P),
                          in_maint(U,We))), Onmaint),
           OnMaint =< Max ).
```

We also define here a unit in maintenance, namely a unit is in maintenance during a certain week if there exists a maintenance M of this unit ongoing that week.

```
in_maint(U,W) <- exists(M,B,E) :
                   maint_for_unit(M,U),start(M,B,E),
                   B =<W, W=< E.
```

- The capacity of the units in maintenance belonging to a certain area should not exceed a given area maximum. To represent this, the summation aggregate is needed. A table of capacity(U,C) describes for each unit its capacity.

```
fol forall(A,Max,We,CapOnMaint):
    area(A),area_max(A,Max),week(We),
    sum(set([U],
            (unit(U),in_area(U,A),in_maint(U,We))),
        lambda([Un],
               C where capacity(Un,C)), CapOnMaint)
    => 0 =< CapOnMaint, CapOnMaint =< Max.

in_area(U,A) <- unit_in_plant(U,P),
                plant_in_area(P,A).
```

The above specification describes the problem properly. Given input data the solver comes up with a schedule for the maintenances. However, the problem is to find an optimal solution that keeps risk of power shortage low. To do this, an optimality function must be used.

This optimality function was proposed to us by the people of the company. In the past, the company has kept track of the electricity consumption during the year. These data can be used to compute an estimate for the peak load consumption during each week. Given a schedule, for each week one can compute the reserve capacity: the difference between available capacity (i.e. the sum of capacities of all units not in maintenance during this week) and the estimated peak load. The optimization function is to maximize the minimal reserve capacity over the year.

total_capacity(T) means that T is the sum of all capacities of all units. This is a constant value for the given problem. Peak loads are represented by a table of peakload(Week,Load) atoms. The predicate reserve(Week,R) can be defined as follows:

```
reserve(Week,R) <- exists(Load,T,InMaint) :
   peakload(Week,Load), total_capacity(T),
   sum(set([Unit], (unit(Unit), in_maint(Unit,Week))),
       lambda([U],C where capacity(U,C)), InMaint),
   R = T - Load - InMaint.
```

Using this predicate, an optimal solution for the scheduling problem can be searched for if we add the following to our query

```
minimum(set([R],(exists(W) : reserve(W,R))), M),
maximize(M).
```

This means that we are interested in a solution in which the minimal reserve for one year is as high as possible.

Except for the representation of the optimization function, the above representation is very similar to the one used in (Denecker *et al.* 1997; Seghers & Baeyens 1996). The actual problem, given by the company, consists of scheduling 56 maintenances for 46 units in one year. The original system needed for this 24 hours to setup the CLP constraint store. The bottleneck was the reduction of the aggregates. In the current implementation of SLDNFA the construction only takes 45 seconds. The huge difference comes from the fact that in the original system aggregates were implemented as large disjunctions, over which was backtracked until a consistent constraint store was found. The current implementation will reduce the aggregates to a large finite domain constraint at once, and will not backtrack within it. It is the CLP solver which backtracks in the constructed constraint store; which is much faster than the abductive solver can do. This difference also explains why the current implementation spends more time to find a good (optimal) solution (in 20 minutes we find a solution which is 94% away from the optimal which is 2 or 3 times slower a the original one): the current constructed constraint store will contain all the disjunctions over which the original system had been backtracking to construct a small one.

A comparison with a pure CLP solution and the above solution shows clearly the tradeoff between declarative representation and a very fast solution. The pure (optimized) CLP solution will setup its constraint store in several seconds (3 to 4 seconds), and find the same solution as the above specification within 2 minutes (compared to 20 minutes). But on the other hand the CLP solution is a long program (400 lines) developed in some weeks of time in which the constraints are hided within data structures, where as the above representation is a simple declarative representation of

11 logical formulae, written down after some hours of discussion.

We want to stress the advantages of moving to an (even) more declarative representation than in the CLP solution. Development time, mentioned above, is just one of these advantages. Adaptability, extendibility and maintenance are others. It is our experience (and also reported in the experiments conducted with ACLP (Kakas & Michael 1995; Kakas, Michael, & Mourlas 2000)) that changes in the specification of the problem may result in several hours of work to adapt the CLP program. In the ID-logic representation, the same changes typically require a few minutes of work. The same distinction arises on the level of extendibility and maintenance. Taking these advantages into account, we believe that the reductions in speed mentioned above are a very good buy.

Conclusion

In a lot of real world problems statements involving aggregates are very common. As such aggregates naturally show up in specifications. This paper gives some preliminary results on the feasibility of using declarative specifications with aggregates to solve nontrivial computational problems.

The extended abductive solver used in our experiments, is able to reduce the high level specification to a finite domain constraint store. As can be expected, the generated constraint store tends to be more complex and less tuned to a specific problem than the ones generated by hand-written CLP programs. This disadvantage is covered by advantages as development time, adaptability, extendibility and easier maintenance. However, at least in the context of the above experiments, reasonable efficiency could be obtained. A way to improve the performance will be to optimize the generated constraint store. Currently optimizations done in a hand-written CLP program aren't applied. We expect that an automated optimization will reduce the difference between both substantially.

Further research needs to be done to get a better characterization of the class of problems the implementation can handle. Together with a formal representation of the aggregate expression evaluation procedure we should be able to present a proof of correctness of the procedure for this class. Another topic is how we can broaden this class without loosing to much efficiency. Further we also expect that the introduction of tabling in the proof procedure will facilitate the treatment of aggregates considerably.

Acknowledgements

References

[Carlsson, Ottosson, & Carlson] Carlsson, M.; Ottosson, G.; and Carlson, B. An open-ended finite domain constraint solver. 191–206.

[Clark 1978] Clark, K. 1978. Negation as failure. In Gallaire, H., and Minker, J., eds., *Logic and Databases.* Plenum Press. 293–322.

[De Mot *et al.* 1999] De Mot, E.; Denecker, M.; De Schreye, D.; and Janssens, G. 1999. Prescriptive type inference as an application of abduction. In Etalle, S., ed., *Proceedings of the Eleventh Benelux Workshop on Logic Programming,* 10. Benelog.

[Denecker & De Schreye 1992] Denecker, M., and De Schreye, D. 1992. SLDNFA; an abductive procedure for normal abductive programs. In Apt, K., ed., *Proc. of the International Joint Conference and Symposium on Logic Programming,* 686–700. MIT Press.

[Denecker & De Schreye 1998] Denecker, M., and De Schreye, D. 1998. SLDNFA: an abductive procedure for abductive logic programs. *Journal of Logic Programming* 34(2):111–167.

[Denecker & Van Nuffelen 1999] Denecker, M., and Van Nuffelen, B. 1999. Experiments for integration CLP and abduction. In Apt, K. R.; Kakas, A. C.; Monfroy, E.; and Rossi, F., eds., *Workshop on Constraints,* 1–15. ERCIM/COMPULOG.

[Denecker *et al.* 1997] Denecker, M.; Vandecasteele, H.; Schreye, D. D.; Seghers, G.; and Baeyens, T. 1997. Scheduling by "Abductive Execution" of a Classical Logic Specification. Compulog Network Meeting on Constraint Programming.

[Denecker, Marek, & Truszczynski 1998] Denecker, M.; Marek, V.; and Truszczynski, M. 1998. Fixpoint 3-valued semantics for autoepistemic logic. In *NM'98, Seventh International Workshop on Nonmonotonic Reasoning.*

[Denecker 1995] Denecker, M. 1995. A Terminological Interpretation of (Abductive) Logic Programming. In Marek, V.; Nerode, A.; and Truszczynski, M., eds., *International Conference on Logic Programming and Nonmonotonic Reasoning,* 15–29. Springer, Lecture notes in Artificial Intelligence 928.

[Denecker 1998] Denecker, M. 1998. The well-founded semantics is the principle of inductive definition. In Dix, J.; Fariñas del Cerro, L.; and Furbach, U., eds., *Logics in Artificial Intelligence,* 1–16. Schloss Daghstull: Springer-Verlag, Lecture notes in Artificial Intelligence 1489.

[Denecker 2000] Denecker, M. 2000. Extending classical logic with inductive definitions. In Baral, C., and Truszczynski, M., eds., *Proceedings of NMR'2000.* 8th Intl. Workshop on Non-Monotonic Reasoning.

[Gordon & Shenk 1998] Gordon, P., and Shenk, M. 1998. *Solitaire Battleships, 108 Challenging Logic Puzzles.* Sterling Publishing Company, Inc.

[Jaffar & Maher 1994] Jaffar, J., and Maher, M. 1994. Constraint logic programming: A survey. *Journal of Logic Programming* 19–20:503–581.

[Kakas & Michael 1995] Kakas, A. C., and Michael, A. 1995. Integrating abductive and constraint logic programming. In Sterling, L., ed., *Proceedings of the 12th International Conference on Logic Programming*, 399–416. Cambridge: MIT Press.

[Kakas, Kowalski, & Toni 1993] Kakas, A. C.; Kowalski, R.; and Toni, F. 1993. Abductive Logic Programming. *Journal of Logic and Computation* 2(6):719–770.

[Kakas, Michael, & Mourlas 2000] Kakas, A.; Michael, A.; and Mourlas, C. 2000. ACLP: Abductive constraint logic programming. *To appear in the Journal of Logic Programming.*

[Kemp & Stuckey 1991] Kemp, D. B., and Stuckey, P. J. 1991. Semantics of logic programs with aggregates. In Saraswat, Vijay; Ueda, K., ed., *Proceedings of the 1991 International Symposium on Logic Programming (ISLP'91)*, 387–404. San Diego, CA: MIT Press.

[Ross & Sagiv 1992] Ross, K. A., and Sagiv, Y. 1992. Monotonic aggregation in deductive databases. In ACM., ed., *PODS '92. Proceedings of the Eleventh ACM SIGACT-SIGMOD-SIGART Symposium on Principles of Database Systems, June 2–4, 1992, San Diego, CA*, 114–126. New York, NY 10036, USA: ACM Press.

[Seghers & Baeyens 1996] Seghers, G., and Baeyens, T. 1996. Solving a combinatorial maintenance problem for electrical power plants, developed in OLP-FOL and CLP. Master's thesis, Department of Computing Science, K.U.Leuven.

[van Emden & Kowalski 1976] van Emden, M., and Kowalski, R. 1976. The semantics of Predicate Logic as a Programming Language. *Journal of the ACM* 4(4):733–742.

[Van Gelder 1992] Van Gelder, A. 1992. The well-founded semantics of aggregation. In ACM., ed., *PODS '92. Proceedings of the Eleventh ACM SIGACT-SIGMOD-SIGART Symposium on Principles of Database Systems, June 2–4, 1992, San Diego, CA*, 127–148. New York, NY 10036, USA: ACM Press.

[Van Nuffelen 2000] Van Nuffelen, B. 2000. SLDNFA-system. In Dix, J., ed., *accepted for the workshop NMR'2000*, 1–6. 8th Intl. Workshop on Non-Monotonic Reasoning, special track on System descriptions and demonstration.