

Distorted English Alphabet Identification: An application of Difference Boosting Algorithm.

Ninan Sajeeth Philip* and K. Babu Joseph†

Cochin University of Science and Technology, Kochi.

3rd June 2017

Abstract

The difference boosting algorithm is used on the letters dataset from the UCI repository to classify distorted raster images of English alphabets. In contrast to rather complex networks, the difference boosting is found to produce comparable or better classification efficiency on this complex problem. With a complete set of 16000 training examples and two chances for making the correct prediction, the network classified correctly in 98.35% instances of the complete 20,000 examples. The accuracy in the first chance was 94.1%.

Keywords: Difference Boosting, Neural Networks, Boosting.

*E-mail : nsp@stthom.ernet.in

†smanager@giasmd01.vsnl.net.in

1 Introduction

Learning is a process that involves identification and distinction of objects. A child learns by examples. In the early days of learning, a child identifies the proximity of his mother by smell or sound or body temperature. Thus it is often not easy for him to distinguish between his mother, grandmother or mother's sister in the first few months of his learning process. But as he grows and his brain develops, he starts identifying the differences in each of his observations. He learns to differentiate colours, between flowers and leaves, between friends and enemies. This is reflected in his responses too. From the binary expressions of smiling and weeping, different expressions appear in his reactions. He learns to read the expressions on the faces of others and also develops the skill to manage it. Difference boosting is an attempt to implement these concepts of child learning into machine code. In the first round it picks up the unique features in each observation using the naive Bayes' theorem, something that is shown to be very similar to the functions of animal brain in object identification. In the second step it attempts to boost the differences in these features that enables one to differentiate almost similar objects. This is analogous to the contention of the child that something that looks like a human but having a tail is a monkey. The feature 'having a tail' gets boosted even if all the other features are strikingly similar. Or, this is how the mother is able to effortlessly distinguish her identical twin children or the artist is able to alter the expressions on the face of his portrait with a few strokes of disjoint line segments. Classic examples to these include the lighting effects on the facial expressions of

a sculpture. In all these examples, the brain does not appear to keep the details but only those unique differences required to identify the objects.

In complex problems, such as the distorted English alphabet detection example discussed in this paper, the selection of a complete training set is not trivial. We thus devise a set of rules to identify the examples from a dataset that form the complete training set for the network. The rules we follow in the said example are:

1. If the network classify an object incorrectly, but with a high probability of around 90% or above, the example could be a new possibility and should be included in the training set.

2. The network is allowed to produce a second guess on the possible class of an object when it fails in the first prediction. If this is a correct guess and if the difference between the degree of confidence between the first and the second guess is greater than 90% or less than 2%, again the example is assumed to be a new possibility or is in the vicinity of boostable examples and is added to the training set.

The underlying logic in these rules are simple. Psychologists point out that the human expert also develops his skill primarily based on past experiences rather than on logical deduction or symbolic reasoning[1]. We also expect a similar situation in the learning process and assume that if an example is incorrectly classified with a confidence higher than 90%, there could be two possible reasons for this. One possibility is that the network is unaware of the existence of that example in the stated class. The other possibility is that the features used are identical to that of an object from another class. In this case, classification of the object into its actual class is difficult without

additional information. Assuming that the reason for the misclassification is that such an sample is not known in the training set, we add that sample also to the training set. In the second rule, we take the difference of the confidence levels since we want to identify new examples and to tackle the border problem that makes it difficult for the network to identify the exact class based on the limited information content in the given features. The first condition picks up the new examples while the second condition picks up the border examples. These are the so-called 'difficult problems' in the learning process.

One word of caution here is that the purpose of these rules are just to pickup a complete set of examples in the training set, which is a pre-requisite of any probability dependent classification problem. Once this dataset is generated, the training process is done on this dataset and the testing is done on the entire dataset and also on the independent test set. A good classification is when the classification accuracy in both these cases are more or less the same.

2 Naive Bayesian learning

Each object has some characteristic features that enables the human brain to identify and characterize them. These feature values or attributes might be either associated to the object by a logical AND or a logical OR relation. The total probability of a system with feature values associated by the logical OR relation is the sum of the individual probabilities. Naive Bayesian classifiers handles this situation by assigning probability distribution values for each

attribute separately. If on the otherhand, the attributes are associated by a logical AND relation, meaning that each attribute value should be somewhere around a stipulated value simultaneously, then the total probability is given by the product of the individual probabilities of the attribute values.

Now, the naive Bayesian classifier assumes that it is possible to assign some degree of confidence to each attribute value of an example while attempting to classify an object. Assume that the training set is complete with K different known discrete classes. Then a statistical analysis should assign a maximal value of the conditional probability $P(C_k | U)$ for the actual class C_k of the example. By Bayes' rule this probability may be computed as :

$$P(C_k | U) = \frac{P(U | C_k) P(C_k)}{\sum_K P(U | C_k) P(C_k)}$$

$P(C_k)$ is also known as the background probability. $P(U | C_k)$ is given by the product of the probabilities due to individual attributes. That is:

$$P(U | C_k) = \prod_m P(U_m | C_k)$$

Following the axioms of set theory, one can compute $P(U_m | C_k)$ as $P(U_m \cap C_k)$. This is nothing but the ratio of the total count of the attribute value U_m in class C_k to the number of examples in the entire training set. Thus naive Bayesian classifiers complete a training cycle much faster than perceptrons or feed-forward neural networks.

3 Difference Boosting

Boosting is an iterative process by which the network upweights misclassified examples in a training set until it is correctly classified. The Adaptive Boosting (AdaBoost) algorithm of Freund and Schapire [2, 3] attempts the same thing. In this paper, we present a rather simple algorithm for boosting. The structure of our network is identical to AdaBoost in that it also modifies a weight function. Instead of computing the error in the classification as the total error produced in the training set, we take each misclassified example and apply a correction to its weight based on its own error. Also, instead of upweighting an example, our network upweights the weight associated to the probability $P(U_m | C_k)$ of each attribute of the example. Thus the modified weight will affect all the examples that have the same attribute value even if its other attributes are different. During the training cycle, there is a competitive update of attribute weights to reduce the error produced by each example. It is expected that at the end of the training epoch the weights associated to the probability function of each attribute will stabilize to some value that produces the minimum error in the entire training set. Identical feature values compete with each other and the differences get boosted up. Thus the classification becomes more and more dependent on the differences rather than on similarities. This is analogous to the way in which the human brain differentiates between almost similar objects by sight, like for example, rotten tomatoes from a pile of good ones.

Let us consider a misclassified example in which P_k represent the computed probability for the actual class k and P_k^* that for the wrongly repre-

sented class. Our aim is to push the computed probability P_k to some value greater than P_k^* . In our network, this is done by modifying the weight associated to each $P(U_m | C_k)$ of the misclassified item by the negative gradient of the error, i.e. $\Delta W_m = \alpha \left[1 - \frac{P_k}{P_k^*} \right]$. Here α is a constant which determines the rate at which the weight changes. The process is repeated until all items are classified correctly or a predefined number of rounds completes.

4 The classifier network.

Assuming that the occurrences of the classes are equally probable, we start with a flat prior distribution of the classes ,i.e. $P(C_k) = \frac{1}{N}$. This might appear unrealistic, since this is almost certain to be unequal in most practical cases. The justification is that since $P(C_K)$ is also a weighting function, we expect this difference also to be taken care of by the connection weights during the boosting process. The advantage on the otherhand is that it avoids any assumptions on the training set regarding the prior estimation. Now, the network presented in this paper may be divided into three units. The first unit computes the Bayes' probability for each of the training examples. If there are M number of attributes with values ranging from m_{min} to m_{max} and belonging to one of the K discrete classes, we first construct a grid of equal sized bins for each k with columns representing the attributes and rows their values. Thus a training example S_i belonging to a class k and having one of its attributes l with a value m will fall into the bin B_{klm} for which the Euclidean distance between the center of the bin and the attribute value is a minimum. The number of bins in each row should cover the range of the

attributes from m_{min} to m_{max} . It is observed that there exist an optimum number of bins that produce the maximum classification efficiency for a given problem. For the time being, it is computed by trial and error. Once this is set, the training process is simply to distribute the examples in the training sets into their respective bins. After this, the number of attributes in each bin i for each class k is counted and this gives the probability $P(U_m | C_k)$ of the attribute m with value $U_m \equiv i$ for the given $C_k = k$. The basic difference of this new formalism with that of the popular gradient descent backpropagation algorithm and similar Neural Networks is that, here the distance function is the distance between the probabilities, rather than the feature magnitudes. Thus the new formalism can isolate overlapping regions of the feature space more efficiently than standard algorithms.

The naive Bayesian learning fails when the data set represent an XOR like feature. To overcome this, associated to each row of bins of the attribute values we put a tag that holds the minimum and maximum values of the other attributes in the data example. This tag acts as a level threshold window function. In our example, if an attribute value in the example happens to be outside the range specified in the tag, then the computed $P(U_m | C_k)$ of that attribute is reduced to one-fourth of its actual value (gain of 0.25). Applying such a simple window enabled the network to handle the XOR kind of problems efficiently.

The second unit in the network is the gradient descent boosting algorithm. To do this, each of the probability components $P(U_m | C_k)$ is amplified by a connection weight before computing $P(U | C_k)$. Initially all the weights are set to unity. For a correctly classified example, $P(U | C_k)$ will be

a maximum for the class specified in the training set. For the misclassified items, we increment its weight by a fraction ΔW_m . The training set is read repeatedly for a few rounds and in each round the connection weights of the misclassified items are incremented by $\Delta W_m = \alpha \left[1 - \frac{P_k}{P_k^*} \right]$ as explained in section 3, until the item is classified correctly.

The third unit computes $P(C_k | U)$ as :

$$P(C_k | U) = \frac{\prod_m P(U_m | C_k) W_m}{\sum_K \prod_m P(U_m | C_k) W_m}$$

If this is a maximum for the class given in the training set, the network is said to have learned correctly. The wrongly classified items are re-submitted to the boosting algorithm in the second unit.

5 Results on the letters dataset

The letters dataset consists of 20,000 unique letter images generated randomly distorting pixel images of the 26 uppercase letters from 20 different commercial fonts. Details of these dataset may be found in [4]. The parent font represented a full range of character types including script, italic, serif and Gothic. The features of each of the 20,000 characters were summarized in terms of 16 primitive numerical attributes. The attributes are[4]:

1. The horizontal position, counting pixels from the left edge of the image, of the center of the smallest rectangular box that can be drawn with all "on" pixels inside the box.
2. The vertical position, counting pixels from the bottom, of the box.

3. The width, in pixels, of the box.
4. The height, in pixels, of the box.
5. The total number of "on" pixels in the character image.
6. The mean horizontal position of all "on" pixels relative to the center of the box and divided by the width of the box. This feature has a negative value if the image is "left heavy" as would be the case for the letter L.
7. The mean vertical position of all "on" pixels relative to the center of the box and divided by the height of the box.
8. The mean squared value of the horizontal pixel distances as measured in 6 above. This attribute will have a higher value for images whose pixels are more widely separated in the horizontal direction as would be the case for the letters W and M.
9. The mean squared value of the vertical pixel distances as measured in 7 above.
10. The mean product of the horizontal and vertical distances for each "on" pixel as measured in 6 and 7 above. This attribute has a positive value for diagonal lines that run from bottom left to top right and a negative value for diagonal lines from top left to bottom right.
11. The mean value of the squared horizontal distance times the vertical distance for each "on" pixel. This measures the correlation of the horizontal variance with the vertical position.

12. The mean value of the squared vertical distance times the horizontal distance for each "on" pixel. This measures the correlation of the vertical variance with the horizontal position.
13. The mean number of edges (an "on" pixel immediately to the right of either an "off" pixel or the image boundary) encountered when making systematic scans from left to right at all vertical positions within the box. This measure distinguishes between letters like "W" or "M" and letters like "I" or "L".
14. The sum of the vertical positions of edges encountered as measured in 13 above. This feature will give a higher value if there are more edges at the top of the box, as in the letter "Y".
15. The mean number of edges (an "on" pixel immediately above either an "off" pixel or the image boundary) encountered when making systematic scans of the image from bottom to top over all horizontal positions within the box.
16. The sum of horizontal positions of edges encountered as measured in 15 above.

Using a Holland-style adaptive classifier and a training set of 16,000 examples, the classifier accuracy reported on this dataset[4] is a little over 80%. The naive Bayesian classifier[5] produces an error rate of 25.26% while when boosted with AdaBoost reduces the error to 24.12%. Using AdaBoost on the C4.5 algorithm[6] could reduce the error to 3.1% on the testset. However the computational power required over 100 machines to generate the tree

structure[7] for its effectuation. A fully connected MLP with 16-70-50-26 topology[7] gave an error of 2.0% with AdaBoost and required 20 machines to implement the system.

The proposed algorithm on a single Celeron processor running at 300MHz and Linux 6.0 attained an error rate of 14.2 % on the independent testset of 4000 examples in less than 15 minutes. Applying the rules mentioned in section one resulted in an overall error of 5.9 % on 20,000 examples. With two chances, the error went as low as 1.65 %. The result is promising taking into consideration the low computational power required by the system. Since bare character recognition of this rate can be improved with other techniques such as grammer and cross-word lookup methods, we expect a near 100 % recognition rate on such systems. Further work is to look into this aspect in detail.

6 Conclusion

Bayes' rule on how the degree of belief should change on the basis of evidences is one of the most popular formalism for brain modeling. In most implementations, the degree of belief is computed in terms of the degree of agreement to some known criteria. However, this has the disadvantage that some of the minor differences might be left unnoticed by the classifier. We thus device a classifier that pays more attention to differences rather than similarities in identifying the classes from a dataset. In the training epoch, the network identifies the apparent differences and magnify them to separate out classes. We applied the classifier on many practical problems and found

that this makes sense. The application of the method on the letter dataset produced an error as low as 1.65 % when two chances were given to make the prediction. Further study is to look into the application of other language recognition techniques in conjunction with the network.

Acknowledgment

Ninan Sajeeth Philip would like to thank David J Slate (Pattern Recognition Group, Odesta Corporation, Evanston) for sending him the original of their paper and also for the useful comments and discussions.

References

- [1] Charness, N., Aging and skilled problem-solving, *Journal of experimental psychology:General*, 110, 21-38, 1981.
- [2] Yoav Freund and Robert Schapire, E., A decision-theoretic generalization of on-line learning and an application to boosting. in *Proceedings of Second European Conference on Computational Learning Theory*, 23-37,1995
- [3] Yoav Freund and Robert Schapire, E., A decision-theoretic generalization of on-line learning and an application to boosting. *Journal of Computer and System Sciences*, 55(1), 119-139, 1997.
- [4] Peter W Frey and David J Slate, Letter Recognition using Holland-style adaptive classifiers. *Machine Learning*, 6,161-182 ,1991.
- [5] Kai Ming Ting and Zijian Zheng, *Proceedings of PAKDD'99*, Springer-Verlag, 1999.

- [6] Robert E. Schapire, Yoav Freund, Peter Bartlett and Wee Sun Lee.
Boosting the margin: A new explanation for the effectiveness of voting
methods, Machine Learning: Proceedings of the fourteenth International
Conference, 1997.

- [7] Holger Schwenk and Yoshua Bengio, Adaptive Boosting of Neural Net-
works for Character Recognition, Technical report #1072, Department
d' Informatique et Recherche Opérationnelle, Université de Montréal,
Montreal, Qc H3C-3J7, Canada.