THE PERFORMANCE OF THE BATCH LEARNING ALGORITHM

IGOR RIVIN

Abstract. We analyze completely the convergence speed of the *batch learning algorithm*, and compare its speed to that of the memoryless learning algorithm and of learning with memory (as analyzed in [KR2001b]). We show that the batch learning algorithm is never worse than the memoryless learning algorithm (at least asymptotically). Its performance *vis-a-vis* learning with full memory is less clearcut, and depends on certain probabilistic assumptions.

Introduction

The original motivation for the work in this paper was provided by research in learning theory, specifically in various models of language acquisition (see, for example, [KNN2001, NKN2001, KN2001]). In the paper [KR2001b], we had studied the speed of convergence of the *memoryless learner algorithm*, and also of *learning with full memory*. Since the *batch learning algorithm* is both widely known, and believed to have superior speed (at the cost of memory) to both of the above methods by learning theorists, it seemed natural to analyze its behavior under the same set of assumptions, in order to bring the analysis in [KR2001a] and [KR2001b] to a sort of closure. It should be noted that the detailed analysis of the batch learning algorithm is performed under the assumption of *independence*, which was not explicitly present in our previous work. For the impatient reader we state our main result (Theorem 6.1) immediately (the reader can compare it with the results on the memoryless learning algorithm and learning with full memory, as summarized in Theorem 2.1):

Theorem A. *Let N_Δ be the number of steps it takes for the student to have probability $1-\Delta$ of learning the concept using the batch learner algorithm. Then we have the following estimates for N_Δ:*

1991 *Mathematics Subject Classification.* 60E07, 60F15, 60J20, 91E40, 26C10.

Key words and phrases. learning theory, zeta functions, asymptotics.

The author would like to think the EPSRC and the NSF for support, and Natalia Komarova and Ilan Vardi for useful conversations.

- *if the distribution of overlaps is* uniform*, or more generally, the density function* $f(1-x)$ *at* 0 *has the form* $f(x) = c + O(x^\delta)$, $\delta, c > 0$, *then there exist positive constants* C_1, C_2 *such that*

$$\lim_{n\to\infty} \mathbf{P}\left(C_1 < \frac{N_\Delta}{(1-\Delta)^2 n} < C_2\right) = 1$$

- *if the probability density function* $f(1-x)$ *is asymptotic to* $cx^\beta + O(x^{\beta+\delta}), \quad \delta, \beta > 0$, *as* x *approaches* 0, *then*

$$\lim_{n\to\infty} \mathbf{P}\left(c_1 < \frac{N_\Delta}{|\log \Delta| n^{\frac{1}{1+\beta}}} < c_2\right) = 1,$$

  for some positive constants c_1, c_2*;*
- *if the asymptotic behavior is as above, but* $-1 < \beta < 0$, *then*

$$\lim_{x\to\infty} \mathbf{P}\left(\frac{1}{x} < \frac{N_\Delta}{|\log \Delta| n^{1/(1+\beta)}} < x\right) = 1$$

The plan of the paper is as follows: in this Introduction we recall the learning algorithms we study; in Section 1 we define our mathematical model; in Section 2 we recall our previous results, in Section 3 we begin the analysis of the batch learning algorithm, and introduce some of the necessary mathematical concepts; in Sections 4-6 we analyze the three cases stated in Theorem A, and we summarize our findings in Section 7.

Memoryless Learning and Learning with Full Memory. The general setup is as follows: There is a collection of concepts $R_0, \ldots, R_n$ and words which refer to these concepts, sometimes ambiguously. The teacher generates a stream of words, referring to the concept R_0. This is not known to the student, but he must learn by, at each step, guessing some concept R_i and checking for consistency with the teacher's input. The *memoryless learner algorithm* consists of picking a concept R_i at random, and sticking by this choice, until it is proven wrong. At this point another concept is picked randomly, and the procedure repeats. *Learning with full memory* follows the same general process with the important difference that once a concept is rejected, the student never goes back to it. It is clear (for both algorithms) that once the student hits on the right answer R_0, this will be his final answer. We would like to estimate the probability of having guessed the right answer is after k steps, and also the expected number of steps before the student settles on the right answer.

Batch Learning. The batch learning situation is similar to the above, but here the student records the words $w_1, \ldots, w_k, \ldots$ he gets from the teacher. For each word w_i , we assume that the student can find (in his textbook, for example) a list L_i of concepts referred to by the word. If we define

$$\mathcal{L}_k = \bigcap_{i=1}^{k} L_i,$$

then we are interested in the smallest value of k such that $\mathcal{L}_k = \{R_0\}$. This value k_0 is the time it has taken the student to learn the concept R_0. We think of k_0 as a random variable, and we wish to estimate its expectation.

1. The mathematical model

We think of the words referring to the concept R_0 as a probability space $\mathcal{P}$. The probability that one of these words also refer to the concept R_i shall be denoted by p_i; the probability that a word refers to concepts $R_{i_1}, \ldots, R_{i_k}$ shall be denoted by $p_{i_1 \ldots i_k}$. All the results described below (obviously) depend in a crucial way on the $p_1, \ldots, p_n$ and (in the case of the batch learning algorithm) also on the joint probabilities. Since there is no *a priori* reason to assume specific values for the probabilities, we shall assume that all of the p_i are themselves *independent, identically distributed random variables.* We shall refer to their common distribution as $\mathcal{F}$, and to the density as f. It turns out that the convergence properties of the various learning algorithms depend on the local analytic properties of the distribution $\mathcal{F}$ at 1 – some moments reflection will convince the reader that this is not really so surprising.

Sharper analysis of the batch learning algorithm, depends on the *independence hypothesis*:

$$p_{i_1 \ldots i_k} = p_{i_1} \ldots p_{i_k}.$$

It is again not too surprising that some such assumption on correlations ought to be required for precise asymptotic results, though it is obviously the subject of a (non-mathematical) debate as to whether assuming that the various concepts are truly independent is reasonable from a cognitive science point of view.

2. Previous results

In previous work [KR2001a] and [KR2001b] we obtained the following result.

Theorem 2.1. *Let N_Δ be the number of steps it takes for the student to have probability $1 - \Delta$ of learning the concept. Then we have the following estimates for N_Δ:*

- *if the distribution of overlaps is* uniform, *or more generally, the density function $f(1-x)$ at 0 has the form $f(x) = c + O(x^\delta)$, $\delta, c > 0$, then there exist positive constants C_1, C_2, C_1', C_2' such that*
$$\lim_{n\to\infty} \mathbf{P}\left(C_1 < \frac{N_\Delta}{|\log \Delta| n \log n} < C_2\right) = 1$$
for the memoryless algorithm and
$$\lim_{n\to\infty} \mathbf{P}\left(C_1' < \frac{N_\Delta}{(1-\Delta)^2 n \log n} < C_2'\right) = 1$$
when learning with full memory;
- *if the probability density function $f(1-x)$ is asymptotic to $cx^\beta + O(x^{\beta+\delta})$, $\quad \delta, \beta > 0$, as x approaches 0, then for the two algorithms we have respectively*
$$\lim_{n\to\infty} \mathbf{P}\left(c_1 < \frac{N_\Delta}{|\log \Delta| n} < c_2\right) = 1,$$
and
$$\lim_{n\to\infty} \mathbf{P}\left(c_1' < \frac{N_\Delta}{(1-\Delta)^2 n} < c_2'\right) = 1$$
for some positive constants c_1, c_2, c_1', c_2';
- *if the asymptotic behavior is as above, but $-1 < \beta < 0$, then*
$$\lim_{x\to\infty} \mathbf{P}\left(\frac{1}{x} < \frac{N_\Delta}{|\log \Delta| n^{1/(1+\beta)}} < x\right) = 1$$
for the memoryless learning algorithm, and similarly
$$\lim_{x\to\infty} \mathbf{P}\left(\frac{1}{x} < \frac{N_\Delta}{(1-\Delta)^2 n^{1/(1+\beta)}} < x\right) = 1$$
for learning with full memory.

Recall that $f(x) = \Theta(g(x))$ means that for sufficiently large x, the ratio $f(x)/g(x)$ is bounded between two strictly positive constants. The distribution of overlaps referred to above is simply the distribution $\mathcal{F}$. Notice that the theorem says nothing about the situation when $\mathcal{F}$ is supported in some interval $[0, a]$, for $a < 1$. That case is (presumably) of scientific interest, but mathematically it is relatively trivial: we replace the arguments of all the Θs above by 1, though, of course, we are thereby hiding the dependence on a.

3. General bounds on the batch learner algorithm

Consider a set of words $w_1, \ldots, w_k$. The probability that they all refer to the concept R_i is, obviously p_i^k.

Lemma 3.1. *The probability q_k that we still have not learned the concept R_0 after k steps is bounded above by $\sum_{i=1}^n p_i^k$, and below by $\max_i p_i^k$.*

Proof. Immediate. □

We will first use these upper and lower bounds to get corresponding bounds on the convergence speed of the batch learner algorithm, and then invoke the independence hypothesis to sharpen these bounds in many cases.

We begin with a trivial but useful lemma.

Lemma 3.2. *Let G be a game where the probability of success (respectively failure) after at most k steps is s_k (respectively $f_k = 1 - s_k$). Then the expected number of steps until success is*

$$\sum_{k=1}^{\infty} k(s_k - s_{k-1}) = \sum_{k=1}^{\infty} s_k = 1 - \sum_{k=1}^{\infty} f_k,$$

if the corresponding sum converges.

Proof. The proof is immediate from the definition of expectation and the possibility of rearrangment of terms of positive series. □

We can combine Lemma 3.2 and Lemma 3.1 to obtain:

Theorem 3.3. *The expected time T of convergence of the batch learner algorithm is bounded as follows:*

$$\sum_{i=1}^{n} \frac{1}{1-p_i} \geq T \geq \max_{1 \leq i \leq n} \frac{1}{1-p_i}. \tag{1}$$

The leftmost term in equation (1) has been studied at length in [KR2001a]. We state a version of the results of [KR2001a] below:

Theorem 3.4. *Let $S = \sum_{i=1}^n \frac{1}{1-p_i}$, where the p_i are independently identically distributed random variables with values in $[0,1]$, with probability density f, such that $f(1-x) = x^\beta + O(x^{\beta+\delta})$, $\quad \delta > 0$ for $x \to 0$. Then If $\beta > 0$, then there exists a mean m, such that $\lim_{n\to\infty} \mathbb{P}(|S/n - m| > \epsilon) = 0$, for any $\epsilon > 0$. If $\beta = 0$, then $\lim_{n\to\infty} \mathbb{P}(|S/(n \log n) - 1| > \epsilon) = 0$). Finally, if $-1 \leq \beta < 0$, then $\lim_{n\to\infty} \mathbb{P}(S/n^{1/\beta+1} - C > a) = g(a)$, where $\lim_{a\to\infty} g(a) = 0$, and C is an arbitrary (but fixed) constant, and likewise*

$$\mathbb{P}(S/n^{1/(\beta+1)} < b) = h(b),$$

where $\lim_{a\to 0} h(a) = 0$,

The right hand side of Eq. (1) is easier to understand. Indeed, let $p_1, \ldots, p_n$ be distributed as usual (and as in the statement of Theorem 3.4). Then

Theorem 3.5.

$$\lim_{n\to\infty} n^{\frac{1}{1+\beta}} \mathbf{E}\left(1 - \max_{1\leq i\leq n} p_i\right) = C,$$

for some positive constant C.

Proof. First, we change variables to $q_i = 1 - p_i$. Obviously, the statement of the Theorem is equivalent to the statement that

$$E = \mathbf{E}(\min_{1\leq i\leq n} q_i) = Cn^{-1/1+\beta}.$$

We also write $h(x) = f(1-x)$, and let H be the distribution function whose density is h, so that $H(x) = 1 - F(1-x)$. Now, the probability of that all of the q_i are greater than t equals $1 - (1 - H(t))^n$, so that

$$E = \int_0^1 t\ d\left[1 - (1 - H(t))^n\right] = \int_0^1 (1 - H(t))^n dt.$$

We change variables $t = u/n^{1/(1+\beta)}$, to obtain

$$E = \frac{1}{n^{1+\beta}} \int_0^{n^{\frac{1}{1+\beta}}} \left(1 - H\left(\frac{u}{n^{1/(1+\beta)}}\right)\right)^n du. \tag{2}$$

Let us write $E = E_1(n) + E_2(n)$, where

$$E_1(n) = \int_0^{n^{\frac{1}{3(\beta+1)}}} \left[1 - H\left(\frac{u}{n^{1/(1+\beta)}}\right)\right]^n du, \tag{3}$$

$$E_2(n) = \int_{n^{\frac{1}{3(\beta+1)}}}^{n^{\frac{1}{1+\beta}}} \left[1 - H\left(\frac{u}{n^{1/(1+\beta)}}\right)\right]^n du, \tag{4}$$

Recall that

$$H(x) = cx^{\beta+1} + O(x^{\beta+\delta+1}). \tag{5}$$

Let

$$\mathcal{I} = \int_0^\infty \exp\left(cx^{1+\beta}\right) dx. \tag{6}$$

We now show:

$$\lim_{n\to\infty} E_1(n) = \mathcal{I}. \tag{7}$$

This is an immediate consequence of Lemma 3.7 and Eq. (5). Also,

$$\lim_{n\to\infty} E_2(n) = 0. \tag{8}$$

Since H is a monotonically increasing function, it is sufficient to show that

$$\lim_{n\to\infty} n^{\frac{1}{1+\beta}} \left[1 - H\left(n^{\frac{2}{3(1+\beta)}}\right)\right]^n = 0.$$

This is immediate from Eq. (5) and Lemma 3.7. □

Remark 3.6. *The argument shows that $C = \mathcal{I}$, where C is the constant in the statement of lemma, and $\mathcal{I}$ is the integral introduced in Eq. (6).*

Lemma 3.7. *Let $f_n(x) = (1 - x/n)^n$, and let $0 \leq z < 1/2$.*

$$f_n(x) = \exp(-x)\left[1 - \frac{x^2}{2n} + O\left(\frac{x^3}{n^2}\right)\right].$$

Proof. Note that

$$\log f_n(x) = n \log(1 - x/n) = -x - \sum_{k=2}^{\infty} \frac{x^k}{kn^{k-1}}.$$

The assertion of the lemma follows by exponentiating the two sides of the above equation. □

We need one final observation:

Theorem 3.8. *The variable $n^{1/(1+\beta)} \min_{i=1}^n q_i$ has a limiting distribution with distribution function $G(x) = 1 - \exp(-x^{1+\beta})$.*

Proof. Immediate from the proof of Theorem 3.5. □

We can now put together all of the above results as follows.

Theorem 3.9. *Let $p_1, \ldots, p_k$ be independently distributed with common density function f, such that $f(1 - x) = cx^\beta + O(x^{\beta+\delta})$, $\delta > 0$. Let T be the expected time of the convergence of the batch learning algorithm with overlaps $p_1, \ldots, p_k$. Then, if $\beta > 0$, then there exist C_1, C_2, such that $C_1 n^{1/(1+\beta)} \leq T \leq C_2 n$, with probability tending to 1 as n tends to ∞. If $\beta = 0$, then there exist C_1, C_2, such that $C_1 n \leq T \leq C_2 n \log n$, with probability tending to one as n tends to ∞. If $\beta > 0$, then $C^{-1} n^{1/(\beta+1)} \leq T \leq C n^{1/(\beta+1)}$ with probability tending to 0 as C goes to infinity.*

The reader will remark that in the case that $\beta > 0$, the upper and lower bounds have the same order of magnitude as functions of n.

4. Independent concepts

We now invoke the independence hypothesis, whereby an application of the inclusion-exclusion principle gives us:

Lemma 4.1. *The probability l_k that we have learned the concept R_0 after k steps is given by*

$$l_k = \prod_{i=1}^{n}(1-p_i^k).$$

Note that the probability s_k of winning the game *on the k-th step* is given by $s_k = l_k - l_{k-1} = (1-l_{k-1}) - (1-l_k)$. Since the expected number of steps T to learn the concept is given by

$$T = \sum_{k=1}^{\infty} k s_k,$$

we immediately have

$$T = \sum_{k=1}^{\infty}(1-l_k)$$

Lemma 4.2. *The expected time T of learning the concept R_0 is given by*

$$T = \sum_{k=1}^{\infty}\left(1 - \prod_{i=1}^{n}\left(1-p_i^k\right)\right). \tag{9}$$

Since the sum above is absolutely convergent, we can expand the products and interchange the order of summation to get the following formula for T:

Notation. Below, we identify subsets of $\{1,\dots,n\}$ with multindexes (in the obvious way), and if $s = \{i_1,\dots,i_l\}$, then

$$p_s \stackrel{\text{def}}{=} p_{i_1}\cdots p_{i_l}.$$

Lemma 4.3. *The expression Eq. (9) can be rewritten as:*

$$T = \sum_{s\subseteq\{1,\dots,n\}} (-1)^{|s|-1}\left(\frac{1}{1-p_s} - 1\right), \tag{10}$$

Proof. With notation as above,

$$\prod_{i=1}^{m}\left(1-p_i^k\right) = \sum_{s\subseteq\{1,\dots,n\}} (-1)^{|s|} p_s^k,$$

so

$$
\begin{aligned}
T &= \sum_{k=1}^{\infty}\left(1-\prod_{i=1}^{n}\left(1-p_i^k\right)\right)\\
&= \sum_{k=1}^{\infty}\left(1-\sum_{s\subseteq\{1,\dots,n\}}(-1)^{|s|}p_s^k\right)\\
&= \sum_{s\subseteq\{1,\dots,n\}}(-1)^{|s|-1}\sum_{k=1}^{\infty}p_s^k\\
&= \sum_{s\subseteq\{1,\dots,n\}}(-1)^{|s|-1}\left(\frac{1}{1-p_s}-1\right),
\end{aligned}
$$

where the change in the order of summation is permissible since all sums converge absolutely. □

Formula (10) is useful in and of itself, but we now use it to analyse the statistical properties of the time of success T under our distribution and independence assumptions. For this we shall need to study the *moment zeta function* of a probability distribution, introduced below. Its detailed properties are investigated in my paper [Rivin2002], where Theorems 4.9, 4.10 and 4.11 below are proved. Below we summarize the definitions and the results.

4.1. Moment zeta function.

Definition 4.4. *Let $\mathcal{F}$ be a probability distribution on a (possibly infinite) interval I, and let $m_k(\mathcal{F}) = \int_I x^k \mathcal{F}(dx)$ be the k-th moment of $\mathcal{F}$. Then the* moment zeta function of $\mathcal{F}$ *is defined to be*

$$\zeta_{\mathcal{F}}(s) = \sum_{k=1}^{\infty} m_k^s(\mathcal{F}),$$

whenever the sum is defined.

The definition is, in a way, motivated by the following:

Lemma 4.5. *Let $\mathcal{F}$ be a probability distribution as above, and let $x_1, \dots, x_n$ be independent random variables with common distribution $\mathcal{F}$. Then*

$$\mathbb{E}\left(\frac{1}{1-x_1\dots x_n}\right) = \zeta_{\mathcal{F}}(n). \tag{11}$$

In particular, the expectation is undefined whenever the zeta function is undefined.

Proof. Expand the fraction in a geometric series and apply Fubini's theorem. ☐

Example 4.6. *For $\mathcal{F}$ the uniform distribution on $[0,1]$, $\zeta_{\mathcal{F}}$ is the familiar Riemann zeta function.*

Using standard techinques of asymptotic analysis, the following can be shown (see [Rivin2002]):

Theorem 4.7. *Let $\mathcal{F}$ be a continuous distribution supported in $[0,1]$, let f be the density of the distribution $\mathcal{F}$, and suppose that $f(1-x) = cx^\beta + O(x^{\beta+\delta})$, for some $\delta > 0$. Then the k-th moment of $\mathcal{F}$ is asymptotic to $Ck^{-(1+\beta)}$, for $C = c\Gamma(\beta)$.*

Corollary 4.8. *Under the assumptions of Theorem 4.7, $\zeta_{\mathcal{F}}(s)$ is defined for $s > 1/(1+\beta)$.*

The moment zeta function can be used to two of the three situations occuring in the study of the batch learner algorithm: In the sequel, we set $\alpha = \beta + 1$.

4.2. $\alpha > 1$**.** In this case, we use our assumptions to rewrite Eq. (10) as

$$\mathbb{E}(T) = -\sum_{k=1}^{n} \binom{n}{k} (-1)^k \zeta_{\mathcal{F}}(k). \tag{12}$$

This, in turn, can be rewritten (by expanding the definition of zeta) as

$$\mathbb{E}(T) = -\sum_{j=1}^{\infty} [(1 - m_j(\mathcal{F}))^n - 1] = \sum_{j=1}^{\infty} [1 - (1 - m_j(\mathcal{F}))^n] \tag{13}$$

Using the moment zeta function we can show:

Theorem 4.9. *Let $\mathcal{F}$ be a continuous distribution supported on $[0,1]$, and let f be the density of $\mathcal{F}$. Suppose further that*

$$\lim_{x\to 1} \frac{f(x)}{(1-x)^\beta} = c,$$

for $\beta, c > 0$. Then,

$$\lim_{n\to\infty} n^{-\frac{1}{1+\beta}} \left[\sum_{k=1}^{n} \binom{n}{k} (-1)^k \zeta_{\mathcal{F}}(k)\right]$$
$$= -\int_0^\infty \frac{1 - \exp\left(-c\Gamma(\beta+1)u^{1+\beta}\right)}{u^2} du$$
$$= -\left(c\Gamma(\beta+1)\right)^{\frac{1}{\beta+1}} \Gamma\left(\frac{\beta}{\beta+1}\right).$$

4.3. $\alpha = 1$**.** In this case,

$$f(x) = L + o(1) \tag{14}$$

as x approaches 1, and so Theorem 4.7 tells us that

$$\lim_{j\to\infty} jm_j(\mathcal{F}) = L. \tag{15}$$

It is not hard to see that $\zeta_{\mathcal{F}}(n)$ is defined for $n \geq 2$. We break up the expression in Eq. (10) as

$$T = \sum_{j=1}^{n} \frac{1}{1-p_j} - 1 + \sum_{s\subseteq\{1,\dots,n\},\quad |s|>1} (-1)^{|s|-1}\left(\frac{1}{1-p_s} - 1\right). \tag{16}$$

Let

$$T_1 = \sum_{j=1}^{n} \frac{1}{1-p_j} - 1,$$

$$T_2 = \sum_{s\subseteq\{1,\dots,n\},\quad |s|>1} (-1)^{|s|-1}\left(\frac{1}{1-p_s} - 1\right).$$

The first sum T_1 has no expectation, however T_1/n does have have a stable distribution centered on $c\log n + c_2$. We will keep this in mind, but now let us look at the second sum T_2. It can be rewritten as

$$T_2(n) = -\sum_{j=1}^{\infty} \left[(1 - m_j(\mathcal{F}))^n - 1 + nm_j(\mathcal{F})\right]. \tag{17}$$

We can again use the moment zeta function to analyse the properties of T_2, to get:

Theorem 4.10. *Let $\mathcal{F}$ be a continuous distribution supported on $[0,1]$, and let f be the density of $\mathcal{F}$. Suppose further that*

$$\lim_{x\to 1} \frac{f(x)}{(1-x)} = c > 0.$$

Then,

$$\sum_{k=2}^{n} \binom{n}{k} (-1)^k \zeta_{\mathcal{F}}(k) \sim cn\log n.$$

To get error estimates, we need stronger assumption on the function f than the weakest possible assumption made in Theorem 4.10.

Theorem 4.11. *Let $\mathcal{F}$ be a continuous distribution supported on $[0,1]$, and let f be the density of $\mathcal{F}$. Suppose further that*

$$f(x) \sim c(1-x) + O\left((1-x)^{\delta}\right),$$

where $\delta > 0$. *Then,*

$$\sum_{k=2}^{n} \binom{n}{k} (-1)^k \zeta_{\mathcal{F}}(k) \sim cn \log n + O(n).$$

The conclusion differs somewhat from that of section 4.2 in that we get an additional term of $cn \log n$, where $c = \lim_{x\to 1} f(x) = \lim_{j\to\infty} jm_j$. This term is equal (with opposing sign) to the center of the stable law satisfied by T_1, so in case $\alpha = 1$, we see that T has no expectation but satisfies a *law of large numbers*, of the

Theorem 4.12 (Law of large numbers)**.** *There exists a constant* C *such that* $\lim_{y\to\infty} \mathbf{P}(|T/n - C| > y) = 0$.

5. $\alpha < 1$

In this case the analysis goes through as in the preceding section when $\alpha > 1/2$, but then runs into considerable difficulties. However, in this case we note that Theorem 3.9 actually gives us tight bounds.

6. The inevitable comparison

We are now in a position to compare the performance of the batch learning algorithm with that of the memoryless learning algorithm and of learning with full memory, as summarized in Theorem 2.1. We combine our computations above with the observation that the batch learner algorithm converges geometrically (Lemma 4.1), to get:

Theorem 6.1. *Let* N_Δ *be the number of steps it takes for the student to have probability* $1 - \Delta$ *of learning the concept using the batch learner algorithm. Then we have the following estimates for* N_Δ*:*

- *if the distribution of overlaps is* uniform*, or more generally, the density function* $f(1-x)$ *at* 0 *has the form* $f(x) = c + O(x^\delta)$, $\delta, c > 0$, *then there exist positive constants* C_1, C_2 *such that*

$$\lim_{n\to\infty} \mathbf{P}\left(C_1 < \frac{N_\Delta}{(1-\Delta)^2 n} < C_2\right) = 1$$

- *if the probability density function* $f(1-x)$ *is asymptotic to* $cx^\beta + O(x^{\beta+\delta})$, $\quad \delta, \beta > 0$, *as* x *approaches* 0, *then*

$$\lim_{n\to\infty} \mathbf{P}\left(c_1 < \frac{N_\Delta}{|\log \Delta| n^{\frac{1}{1+\beta}}} < c_2\right) = 1,$$

for some positive constants c_1, c_2*;*

- *if the asymptotic behavior is as above, but* $-1 < \beta < 0$, *then*
$$\lim_{x \to \infty} \mathbf{P}\left(\frac{1}{x} < \frac{N_\Delta}{|\log \Delta| n^{1/(1+\beta)}} < x\right) = 1$$

Comparing Theorems 2.1 and 6.1, we see that batch learning algorithm is uniformly superior for $\beta \geq 0$, and the only one of the three to achieve *sublinear* performance whenever $\beta > 0$ (the other two *never* do better than linearly, unless the distribution $\mathcal{F}$ is supported away from 1.) On the other hand, for $\beta < 0$, the batch learning algorithm performs comparably to the memoryless learner algorithm, and worse than learning with full memory.

References

[BenOrsz] C. M. Bender and S. Orszag (1999) *Advanced mathematical methods for scientists and engineers, I,* Springer-Verlag, New York.

[KNN2001] Komarova, N. L., Niyogi, P. and Nowak, M. A. (2001) The evolutionary dynamics of grammar acquisition, *J. Theor. Biology*, **209**(1), pp. 43-59.

[KN2001] Komarova, N. L. and Nowak, M. A. (2001) Natural selection of the critical period for grammar acquisition, *Proc. Royal Soc. B*, to appear.

[KR2001a] Komarova, N. L. and Rivin, I. (2001) Harmonic mean, random polynomials and stochastic matrices, *preprint.*

[KR2001b] Komarova, N. L. and Rivin, I. (2001) On the mathematics of learning.

[Niyogi1998] Niyogi, P. (1998). *The Informational Complexity of Learning.* Boston: Kluwer.

[NKN2001] Nowak, M. A., Komarova, N. L., Niyogi, P. (2001) Evolution of universal grammar, *Science* **291**, 114-118.

[Rivin2002] Igor Rivin (2002). The moment zeta function and applications, arxiv.org preprint NT/0201109.

Mathematics department, University of Manchester, Oxford Road, Manchester M13 9PL, UK

Mathematics Department, Temple University, Philadelphia, PA 19122

Mathematics Department, Princeton University, Princeton, NJ 08544
E-mail address: irivin@math.princeton.edu