

Parametric external predicates for the DLV System

G. Ianni, F. Calimeri, A. Pietramala, M.C. Santoro

October 14, 2013

Abstract

This document describes syntax, semantics and implementation guidelines in order to enrich the DLV system with the possibility to make external C function calls. This feature is realized by the introduction of "parametric" external predicates, whose extension is not specified through a logic program but implicitly computed through external code.

1 Intuition

It is very usual the necessity to encode and employ external functions in dlv. For instance, dlv lacks of mathematical functions and string functions such as:

`div(X,Y,Z)`, or `sqr(X,Y)`.
...

2 Syntax

We introduce in dlv programs a new type of predicate which we will call *external* predicate. The name of an external predicate is preceded by the conventional `#` symbol. For instance,

`#fatt(X,N)`, `#sum(1,2,3)`, are (external) atoms with external predicate name. An external atom can appear inside bodies and constraints only. An external predicate cannot be true negated. It can be negated with default negation, and, in this case, its variables must be safe in ordinary way.

3 Semantics

Roughly speaking, the formal semantics of an external atom is the following. Given an external predicate `#p` of arity n , it is associated to it an "oracle" boolean function p' taking n constant arguments. A ground atom `#p(a1, ..., an)` is true iff $p'(a1, ..., an)$ is true.

Example: consider the external predicate `#fatt` implementing the external computation of the factorial function. Its oracle function $fatt'$ is such that $fatt'(a,b)$ is true if b is the factorial of a .

4 Operational semantics

From the operational point of view, external predicates must be implemented at grounding level. It is given a rule

$$H(\dots) :- A1(\dots), \dots, \#P(X1, \dots, XN), \dots, AN(\dots).$$

containing the external atom $\#P$ associated to its oracle P' . If $X1, \dots, Xn$ are safe, external predicates are easily implemented at grounding level by checking if $P'(a1, \dots, an)$ is true for a given set of constants. In this case, it suffices to reorder rules and put external atoms at the very right edge of any rules where they appear. It is very important instead to introduce an operational semantics taking into account the possibility that some of the $X1, \dots, XN$ are not known.

For instance, assume it is given the external predicate $\#sqr(N, S)$, computing the square S of a given number N , we may want the following rule to be possible

$$squares(S) :- number(N), \#sqr(N, S).$$

where S is not bounded. *Note that, as a important collateral effect, S is not bounded at all, neither by an $\#int$ atom. This works around the current need to ground and represent all the integer constants. New integer constants are introduced "on demand". Furthermore, this should allow to deal with floating point values and strings in an efficient way. Only constants actually necessary will be computed.*

In the above case, the call to oracle function sqr' is not feasible, since we do not know a priori which values of S are needed in order to call it.

In order to provide this further feature, we introduce the concept of "talkative oracle".

It is given an external predicate $\#p$, of arity n and its oracle function p' . A *pattern* is a list of a 's and A 's. A a will represent a placeholder for a constant, whereas an A will be a placeholder for a variable. Given a list of terms, the corresponding pattern will be given by replacing each constant with a , and each variable with A . For instance the pattern for the list of terms X, b, Y is A, a, A . Let pat be a pattern of length n having k placeholders a (which we will call input positions), and $n - k$ placeholders of A type (which we will call output positions). A talkative oracle p^{pat} for the pattern pat , associated to the external predicate $\#p$, is a function taking k constant arguments, returning a tuple of arity $n - k$. p^{pat} is such that $p^{pat}(a_1, \dots, a_k) = b_1, \dots, b_{n-k}$ iff $p'(X^{pat})$ is true, where X^{pat} is a list of terms built from pat by in each input position a a_i , and each output position with a b_i .

Given an external predicate $\#p$, it may be associated with one or more "consistent" talkative oracles. For instance, consider the $\#fatt$ external predicate. We associate to it two talkative oracles, $fatt^{a,A}$ and $fatt^{A,a}$. For instance,

$$fatt^{a,A}(3) = 6 \tag{1}$$

$$fatt^{A,a}(6) = 3 \tag{2}$$

consistently with the fact that $fatt'(3, 6)$ is true.

From now on, given an external predicate, we will assume it comes equipped with its oracle and a set of consistent talkative oracles.

5 Body reordering and safeness constraints

At this point we are able to relax the constraint on safeness of external atoms. A variable Xi , belonging to an external atom $p(X1, \dots, Xn)$ may be kept free if p has a talkative oracle associated to a pattern where Xi is in the same position of a A symbol. For instance, assume $\#sqr$ comes with the $sqr^{a,A}$ oracle, and $\#fatt$ comes with the $fatt^{a,A}$ oracle.

The following rules are safe:

$H(S) :- \text{number}(N), \#sqr(N,S).$

$H(S1) :- \text{number}(N), \#fatt(N,S), \#sqr(S,S1).$

The following are, viceversa, unsafe:

$H(S) :- \text{number}(S), \#sqr(N,S).$

$H(S1) :- \text{number}(S), \#fatt(N,S), \#sqr(S,S1).$

Safeness rules A completely defined rule (or constraints), is a rule where the order of instantiation has been completely fixed, from left to right, and for each external atom $\#p$ an unique talkative oracle has been chosen.

We consider three notion of safety:

- Usual safeness. Each variable must appear in at least a positive atom within the body.
- Weak safeness. This notion applies only to a completely defined rule. A variable is weakly safe if it is either,
 - usually safe, or
 - it appears in an external atom in output position (with respect to the corresponding talkative oracle), and all the variables in input position are weakly safe.
- Strong safeness. The strong safeness applies only in those specific cases where recursion may induce an infinite Herbrand universe. A rule will require strong safety if in the rules' dependency graph (the one where rules appears as nodes) it appears in some cycle. A variable X is strongly safe if it is weakly safe. In case X appears within the head of the rule, then X must be usually safe.

From now on we will use the notion of weak safety as default safety constraint. Weak safety coincides with usual safety in case external predicates are not defined. We require strong safety whenever needed.

Strong safeness constraints There are cases where the above "weak" safety rules may lead to the generation of an infinite Herbrand universe when employed in recursive rules. Imagine to define an oracle for the $\#succ$ such that $\#succ(a,b)$ is true if and only if a and b are integers and $b = a + 1$. The grounding algorithm launched on the rule

$\text{int}(X) :- \text{int}(Y), \#succ(X,Y).$

never reaches a fixed point. To avoid such cases, we introduce the "strong" safety criterium which has to be applied only when recursion is detected.

Body reordering In general, given a rule, the following tasks have to be performed:

1. choice a suitable talkative oracle for each external atom in it, such that the rule is weakly safe, and , if necessary, strongly safe; intuitively, this problem is, in its general setting, NP-complete (reduction from covering), but in practice there will be very few talkative oracles for each external atom, so the search space will be actually very small. Please note that there could be more than one solution. We do not discuss at the moment the problem of making the best choice. Just pick a suitable talkative oracle for each external atom.
2. put each external atom in a proper position such that:
 - a. the rule rests safe;
 - b. the atoms ordering is efficient as much as possible.

The two tasks are actually strictly related. Let's try to briefly explain a possible algorithm to solve the problem.

Since the external built-in atoms are, in a sense, functionally conceived (i.e. given a set of terms as "input" the built-in will have return an unique value for the output positions) they can be computed "on-the-fly", each time the grounding module reaches them processing a body of some rule; and so it is preferable to put them as soon as possible. This can also give a way to choose the proper oracle per each atom: each time an "ordinary" atom is placed, it is enough to check per each external atom whether there are some oracles whose "calling patterns" are subsets of the variables which are currently bounded. If so, the external atom can be placed.

Imagine to being ordering the following rule:

$p(X) :- q(X, Y), s(Y, T), m(Z), n(Z, T), \#r(Y, Z, T).$

and let the "r" atom have the only oracle:

$\#r: X, Y \rightarrow Z.$

Now let's be in the partial situation:

$p(X) :- m(Z), q(X, Y), \dots ??? \dots$

The currently bounded variables are $\{Z, X, Y\}$. There is an oracle for $\#r$ with calling signature requiring $\{Y, Z\}$: so the atom $\#r(Y, Z, T)$ can be put just after those currently placed.

Please note as the previous example is trivial. In general, two kinds of problem may arise:

- a. more than one oracle are suitable for an extern atom;
- b. more than one atom are placeable at a given step.

So some criteria in order to choose between oracles has to be depicted; and well as some criteria for choosing the order between external atoms.

Grounding stage modification tips I suggest to introduce a derived class of atoms having the possibility to call the selected talkative oracle whenever it is necessary to discover new values. This atom may follow a caching schema in order not to call the oracle when some call has already been computed.

A supporting class for this new type of atoms could be:

```

class externalpredicate {

    CollectionOfOracles o;
    CollectionOfGroundAtoms g; // Caches already computed ground
                               // atoms. Corresponds to the
                               // current predicate extension.
    CollectionOfOracleCalls c; // Stores already called patterns.
    friend operator< (call v1, call v2);
    CollectionOfGroundAtoms callanoracle( oracle o, call v )
    {
        if (there is no call v1 <= v in c)
            push o.invoke(v) in g;
    }
}

```

a call is simply a list of terms. There is a partial order between calls, given by their generality. For instance the call (a,b,X) precedes (a,b,c), whereas (c,d,X) and (a,d,X) cannot be compared. Each talkative oracle is able to respond to a given pattern, so each call must follow such a pattern. An external atom in a rule has ONLY one talkative oracle which can be called. It is anyway possible that the same external predicate appears in the same rule or in another point of the program and ANOTHER talkative oracle has been associated to it. The class `externalpredicate` concentrates the oracle calls in a central data structure, and performs oracle calls only if necessary. An oracle call must be performed in the `InstantiateRule` stage on demand.

Here: modifications on the `instantiaterule`.

6 External predicate oracle specification language

First of all, we need to provide a public header file (`extpred.h`) containing some public classes, coupled with a binary dynamic library `extpred.so`. The external predicates programmer should include this header in order to program his/her own external predicates.

`extpred.h` should contain:

- The definition of a public `CONSTANT` class. It should be a minimal version of the `TERM` internal class. The `TERM` class should be extended with constructors and methods in order to convert to and from the `CONSTANT` class. It is not necessary that `CONSTANT` is aware of `TERM` at all.

FIXME Another possibility is to do not rely on the `TERM` class, but only have a *data* field and proper conversion methods in order to convert from/to internal dlw types. The latter are essentially only numbers (integers) and strings (`char*`). We will talk about that a little bit later.

The custom code should contain a set of functions taking as argument an `int` representing the arity of the built-in predicate and an array of `CONSTANT`s (each one of these is linked to a term of the predicate).

Each function name will encode the name of the associated external predicate and the pattern to which it responds to. Some of the `CONSTANT`s on the array

constitute the "input" and some others the "output", accordingly to the specific pattern. A `bool` is returned specifying whether the "unification" has succeeded or not.

To do that we suggest to introduce the preprocessor directive

```
#define BUILTIN(name,pattern)
    bool __name__pattern(CONSTANT *argv[])
```

where `name` is the external predicate name, and `pattern` is the pattern of the oracle to be implemented. Pattern is essentially a string constituted by "i" (lowercase) and "O" (uppercase); the succession is positional, and the meaning is that at j-th position of the pattern string will appear an "i" if for the oracle being defined the j-th term is in "input"; and as it can be easily imagined an "O" means that the j-th term is in output. Let's here recall that this will mean that each term for which an "i" is specified has to be bound when the oracle is invoked (and we think to proper reorder the body of the rule in order to ensure that).

Example:

```
// sketch of a external predicate definition source code
# include "extpred.h"

BUILTIN(fatt,iO) // oracle for the pattern (a,A)
{
// We know we want just ONE parameter as input
assert (argc == 2);

// Convert the CONSTANT data to what we do have to compute
int x = argv[1].data;

// Just perform the computation...
assert ( x >= 0);
int fatt = 1;
if ( x > 0)
    for (int i = 1; i <= x; i++)
        fatt *= i;

// Convert the result into the proper CONSTANT data
CONSTANT result (fatt);

// Save the result onto the output vector
output.push_back(result);
}

BUILTIN(fatt,Xx) // oracle for the pattern (A,a)
{
// We know we want just ONE parameter as input
assert (input.size == 1);
```

```

// Convert the CONSTANT data to what we do have to compute
int x = int( (input.begin().data()) );

// Just perform the computation...
assert ( x >= 0);

int base = 1, y = x;
int carry = 0;

bool ok = true;

while ( y > 1 && (carry == 0) )
{
    y = x / base;
    carry = x % base;
}

CONSTANT result;

// FIXME: DEFINE HOW TO REPRESENT FAILING (no unification)
if ( carry == 0 )
    result = base
else
    result = "FAIL";

output.push_back(result)

// Save the result onto the output vector
output.push_back(result);
}

```

Some notes:

- . we do have to finely declare how to treat input and output (i.e.: should the size have to equal the number of "input arguments" of the oracle? If not, the oracle may take the first arguments, or the latest ones...
- . the same for output vector: can we simply care of inserting it from back?
- . how to manage the data contained in each CONSTANT? We should have to precisely project the class, since the types have to be *at least* real constants or numbers...
- . how to model class CONSTANT in order to represent even the failing case (i.e., the corresponding missing unification for a traditional predicate);
- . what else?

7 Using Externally Defined Built-ins into dlV Programs

Once the user has her own libraries and wants to use one or more built-ins into a logic program, she has simply to tell dlV where to find oracle definitions (i.e., where are the (compiled) libraries). The syntax is quite the same as for object oriented languages, and requires the use of the `#include` directive. The directive can ask to import a complete "package" (i.e., everything is in a folder) or a single library file:

```
#include mylib.strings.*
```

imports every lib file contained into strings package, while

```
#include mylib.strings.compare <- FIXME: which extension?
```

imports only the compare library.

From then, the program can exploit all built-ins defined into the imported libraries. For instance, let the `#contains` built-in be defined in some imported library:

```
p(X) :- q(_,X), #contains(X,"stripes").
```

NOTE: the default path should be `./lib`. A command line option should be provided in order to specify a different path (folders separated by semicolons):

```
.... -path=/usr/local/lib/dlv;/home/myuser/lib;
```

7.1 Some Cares

- . In case of predicate name conflict, the last import directive takes precedence, but it is suggested to output a warning in this case. The wanted version of a given predicate may be explicitly selected using an external atom in the form `#packagename.predicatename(Vars)` inside the logic program.
- . Please note as more than one `#import` directive can be given, but all of them have to be put at the **beginning of the file** containing the logic program.
- . The user can freely define "standard" predicates having the same name of some built-in. For instance the same program may contain `p(...)` and `#p(...)`, and they will be distinguished.
- . On the other hand, the user cannot override some identifiers; think about predefined aggregates (`#sum`, `#count`, etc.) or keywords like `#template`.

8 Tips on dynamic linking of external built-ins

We suggest two ways of linking new built-ins: either statically (need recompiling DLV), or dynamically (recompiling is not needed). The second way allows third parties to develop and introduce new built-in libraries for dl原因, and it is recommended.

9 Notes on wrapping of the older builtins (arithmetics and so on)

10 Some Notes on Implementation

1. During rules' reordering (see *orderedBody* in *grounding-body.h*) the function *isAdmissible*¹ is invoked. It is a boolean function receiving a built-in predicate and the list of bounded variables; it checks whether there exists a suitable oracle for the current built-in, and if so it returns true. Please note that (in case of success) it also sets a pointer (or an index, let's see) to the found oracle, after a safety check.
2. A brand new class called `EXTENDED_BUILTIN` will be built. It will have to be able to manage built-ins replacing the existing ones ("standard") as well as the new ("extended"). A possible (rough) structure may be the following:

```
class EXTENDED_BUILTIN
{
    vector<ORACLE> *oracle;

    .....

    bool isAdmissible(const char* name, const TERMS *bounded)
};
```

The `oracle` vector stores all existing oracles related to the single built-in (i.e. all user-defined oracles in the dynamic library).

3. Essentially an instance of this class has to be built per each built-in, once and for all built-in atoms related to it. All these, will refer to this.
All instances of `EXTENDED_BUILTIN` will be stored in an (hash?)map having the name as key.
4. Each instance of a built-in atom should have a pointer to the actual oracle.
5. The `CONSTANT` class have to provide suitable methods in order to cast to standard types, such as `toString()`, `toInt()`, etc.
Ex.: `int base = argv[2].toInt()`
6. In order to call the proper oracle we can use a `switch` for rapid prototyping, and then use a more efficient array of pointers to functions.

¹Or "lookup", we'll see.

7. During the grounding phase it is possible that some oracle generates *new* constants (i.e., not yet present in the Herbrand Universe): think about some "output" constant. If so, we do have to maintain up-to-date the `CONSTANT_NAMES` collection (i.e., the "register" of existing constants). Maybe we will need proper methods.

NOTE: In such a case, a problem about generation of an infinite Universe arises. We could:

- Check under a safety criterion (see above).
 - Make the user responsible about that.
8. Does the `#maxint` setting have to be forced in case of extended built-ins usage?
 9. In a successfully grounded rule a built-in atom will be always true, so it can be deleted. Check whether it may be automatically removed or a new procedure is needed.
 10. It is mandatory to define per each built-in the "base" oracle (i.e., the one with all parameters bounded – iii [...] i).

11 Ambiguity on choosing the oracle

During the body reordering the built-ins are placed and a proper oracle has to be chosen among all admissible (see above). In case more than one oracle are admissible, a way to choose has to be found. Since built-ins are conceived under "functional" semantics we could take any of the oracles: the result will be always the same. For the sake of efficiency, we think that the less are the output `CONSTANTS`, the more is preferable the oracle: in fact, this will allow to avoid more checks with already bounded variables.

NOTE: should it be necessary to ask the user for a partial order between oracles?