

Toward the Implementation of Functions in the DLV System (Preliminary Technical Report)

Francesco Calimeri, Nicola Leone, and others

Department of Mathematics, University of Calabria

Abstract. This document describes the functions as they are treated in the **DLV** system. We give first the language, then specify the main implementation issues.

1 Language Specifications

In this section we briefly provide a formal definition of the language.

1.1 Syntax

A rule is in the form

$$head : - body.$$

where

head is a disjunction of *atoms*;
body is a conjunction of *literals*.

So we have

$$\begin{aligned} head &::= atom \mid [\vee head] \\ body &::= literal \mid [, body] \\ literal &::= [\text{not}] atom \\ atom &::= predicate_symbol \mid predicate_symbol "(" terms ")" \\ terms &::= term \mid [, terms] \end{aligned}$$

and finally

$$term ::= constant \mid variable \mid f "(" t_1 \dots t_n ")"$$

where $t_1 \dots t_n$ are terms and f is a function symbol of arity n .

2 Implementation Issues

Functions implementation can be done involving only the grounding module. Basically, we rewrite the rules containing some function symbol in a specific manner, and set an order in their body; then we properly control the match during the grounding. The resulting ground program is passed to the Model Generator module, that does not need to be changed.

2.1 Rewriting Rules

All rules that do not contain functions are treated as usual, while others are rewritten as follows. *Functions* are “*flattened*” in special body predicates, let’s call them “*function predicates*”, that are *built-in* predicates. They have the same arity as the original functions, plus one. All terms are the same as those in them, but one: the first argument (“*id-argument*”) is an identifier representing the new term; it is special, and it allows to “link” the new predicate symbol to the ones that contain the original function.

Example 1. Let’s consider the rule

$$r : p(s(X)) :- a(X, f(Y,Z)).$$

It contains two function symbols, s and f . We create two new predicates, s^* and f^* , and so rewrite the rule as

$$r' : p(S) :- a(X, F), f^*(F, Y, Z), s^*(S, X).$$

The last arguments in the new predicates are the same as in the original functions; the first one (in this case S and F for s^* and f^* respectively) is their “*id*”, and it appear where necessary instead of the old function symbol (in this case in the p and the a atoms).

Optimization Note If the same function appears in more than an atom, we should create only one new function symbol. But it have to appear once or more than once depending on the arguments.

Example 2. The rule

$$r : p(s(X)) :- q(s(X), Y).$$

will become

$$r : p(S) :- q(S, Y), s^*(S, X).$$

But

$$r : p(s(X)) :- q(s(Y), X).$$

will become

$$r : p(S) :- q(S, X), s^*(S, X), s^*(S, Y).$$

Implementation Note The parser module should continue doing only its standard job, so it should have only to recognize the *function predicates* as new kind of tokens; in particular, as arguments of some predicate. Simply, we have to remember that they are functions and lately rewrite them: the job will be done by the rewriting module.

Some Cares

- A. (*Negations*) No particular attentions should be paid to the rules containing negations. The translation described above works properly.

Example 3. The rule

$$a(X) :- p(X), \textbf{not } ab(s(X)).$$

is rewritten as

$$a(X) :- p(X), s^*(S, X), \textbf{not } ab(S).$$

- B. (*Aggregates*) The “flattening” of rules containing aggregates, on the other hand, requires some care. It can be done in two ways:

1. it is applied *after* the rewriting for the aggregates;
2. it is applied at the beginning; but, in this case, it has to be applied to the conjunctions inside the aggregates.

We think that the way (1.) is easier and more clean.

Example 4. Consider the rule

$$r : a(X) :- X = \#count(Y: p(s(Y)), q(Y)).$$

It can be flattened choosing one of the following.

1. Wait that is rewritten (during the parsing) to

$$\begin{aligned} r_1 : a(X) &:- X = \#count(Y: aux(Y)). \\ r_2 : aux(Y) &:- p(s(Y)), q(Y). \end{aligned}$$

Then flatten only r_2 as

$$r'_2 : aux(Y) :- p(S), q(Y), s^*(S, Y).$$

2. Flatten r as

$$r' : a(X) :- X = \#count(Y: p(S), q(Y), s^*(S, Y)).$$

Then let it be rewritten to

$$\begin{aligned} r'_1 : a(X) &:- X = \#count(Y: aux(Y)). \\ r'_2 : aux(Y) &:- p(S), q(Y), s^*(S, Y). \end{aligned}$$

- C. (*Dependency Graph*) Function predicates are neglected in the dependency graph, in the same way it is already done with built-in predicates: they do not generate any arc in the *DG*.

2.2 Reordering Bodies

In the body-reordering, a function predicate should be inserted in the first position such that:

- a - either the *id* argument is bound, **or**
- b - all previous arguments are bound, excepting at most the *id*.

Example 5. Let's consider the predicate $f^*(Y, Z, F)$ in the rule r' of example 1. Here f^* and s^* should be put so that either F (S , respectively) is bound, or both Y and Z (X , respectively) are bound:

$$r' : p(S) :- a(X, F), f^*(F, Y, Z), s^*(S, X).$$

Note as the atom $a(X, F)$ is considered as first. So if the match does not fail (in the case it does f^* and s^* are not considered at all), X and F become both bound (or they already are, as constants). So for f^* we have the “id-argument” (F , in this case) bound; and for s^* we have all arguments different from the *id* (X , in this case) bound.

Example 6. The rewritten rule

$$r'_1 : m(X, Y) :- s^*(S, X, Y), k(S, T), p(W, Z, T).$$

has the body not properly reordered: s^* is considered as first, when neither the *id* argument nor all others are bound. A correct rewriting is the following:

$$r'_1 : m(X, Y) :- k(S, T), s^*(S, X, Y), p(W, Z, T).$$

Here the function predicate s^* is in the first position such that (at least) one of the two conditions explained above is satisfied: in this case the “id-argument” S is always bound, if the matching process reaches s^* .

Actually, even if we place all function predicates at the end of each body, everything should work; but we risk to loose performance advantages due to the way the algorithm cut the search space of ground rules. So we should put any function predicate in the *place* such that at least one of the two previously exposed conditions holds.

To obtain what described, we can let the old algorithm work as usual, but adding further checks to properly place function predicates.

We can put somewhere the “functional” atoms, for instance at the very end of the body. Then we let the existing algorithm work as usual to place one “standard” atom per time. Between two standard atoms' placements, as soon as one took its place, we check, for each unplaced functional atom, whether (at least) one of the two conditions holds; if so, we put it after the standard atom just placed. It should be not too difficult to check per each argument whether it occurs in at least a placed atom or not. At the end, we are sure that all atoms were properly placed.

Example 7. Let's examine the rule

$$r : p(X) :- q(S, X), t(Y), s^*(S, Y).$$

We put initially s^* at the end. The existing algorithm chooses the first atom, say it q . Then we check s^* : it contains Y and its *id* S . Because the *id* appears in an already placed atom, we can place also s^* in the current position, removing it from the last place. Then t is placed, too, and

we stop. Of course, if the first atom placed was t , we could have placed s^* anyway, because all its arguments but the id were ok.

Basically it is possible to use the same trick we did for the built-ins: we extend the heuristic function computing the “weight” of each atom to deal with “function atoms”. If one of those conditions is verified, then the function atom has weight $-maxint$; otherwise its weight is $+maxint$.

2.3 Grounding

As already said, we would like to make functions transparent as much as possible. So everything could stay the same (actually function predicates can be viewed more or less like standard predicates), excepting for a case during the matching algorithm.

With function predicates, it is possible to add “tuples” to some table even if its symbol appears in the body, but not also in the head (in fact, because of our “flattening”, function predicates can not appear in the head of a rule). So their “tables” can be filled only by the facts (example: $f(s(1))$) or by some body occurrences. For instance, if all arguments but the id are bound as in $f^*(F, a, b)$, then the match function should insert a new tuple (say $\langle a, b, @_2 \rangle$) in f^* , with a new identifier ($@_2$, in this case) to which the id should be bound.

This possibility is due to the fact that after flattening no function predicates appear in the heads; but maybe they do in the original rules. In these cases, we *have* to derive anyway.

Example 8. Let’s have a look at the rule

$$r : p(s(X)) :- q(X).$$

Here if X matches with some value in the table of q , we should derive what is in the head. But the flattened rule appear as

$$r' : p(S) :- q(X), s^*(S, X).$$

So if we match X for some value, if we do not find an appropriate entry in s^* we have to create it and succeed. Of course *nothing* should be done if a matching tuple is present in s^* .

But if the function symbol did not appear in the head of the original rule, we *should not* add new entries.

Example 9. With the rule

$$r : p(X) :- q(X, s(Y)), t(Y).$$

if no matches are found for q , we should fail anyway. The rewritten rule should appear as

$$r' : p(X) :- t(Y), s^*(S, Y), q(X, S).$$

So if we have Y bound, say $Y = b$, and in s^* no tuples of kind $\langle b, _ \rangle$ are found, we have to fail and come back to the choice of Y .

In practice, if the function predicate is not an “head-predicate” (in the sense just discussed), it has to be treated exactly as a “standard” predicate. No changes at all. If it is an “head-predicate”, on the other hand, then 3 cases are possible, when the matching function reaches it (these cases are always the same, even if it is not an head predicate, but then nothing changes anyway):

1 - it has id argument bound, and not all the others are so;

- 2 - it has all arguments bound, including the *id*;
- 3 - it has all arguments bound but the *id*.

It is easy to see as first two cases are the same, in our situation. The function predicate has to be treated exactly like a “standard” one. Otherwise, in case 1, we have to add a new tuple having the values fixed to the bound values, and a *new* identifier. Of course, this is a *try*: only if we reach the last body atom with success we can, let’s say, “commit” the new values. It should work even if the same function with the same arguments appears in more than one predicate, for instance one in the head and one in the body. This will show us also the way we have to behave with respect to the backtracking.

Example 10. Suppose that the body reordering in tail to the flattening gives for the rule

$$r : p(s(X)) :- t(X), q(s(X), Y).$$

the one

$$r' : p(S) :- t(X), s^*(S, X), q(S, Y).$$

Here if we reach s^* , X is bound (say to $X=b$) and S is not; if no tuples of kind $j b, _$ are found, we try to prepare a new tuple $j b, \text{next_valid_value_for_the_identifier_of_} s^* _$, binding S and passing to q . We will fail for sure because no such a value for S is in q (it is new!!!). So we should backtrack, but it is useless to try again with s^* : we have to jump it and come back to t . Note as this seems to be consistent: we should have the same result if the rule was ordered as

$$r'' : p(S) :- q(S, Y), s^*(S, X), t(Y).$$

In this case we succeed only if they match with an existing value in q .

Since a “function atom” always has a key argument which is bound at the time when it is instantiated (either the *id* or all others), then the functional dependency guarantees that it cannot “produce” more than one match. Therefore, during the backward phase (i.e., if the match of a subsequent body atom failed) we should not retry to match a function atom; but we should skip it and go to the previous body atom.

Resuming: no changes if the predicate is not an “head-predicate”; pay attention only if it is *and* the *id* is not bound. About backtracking, the function predicates have to be jumped. It is worth to remember that even in the previous case (3), when the match succeed because of an existing tuple, everything remain the same (no new tuple creation), excepting the backtracking policy (do backjumping). But it is necessary to provide an efficient way to know whether a function atom comes from the head or from the body.

All these stuffs make sense with respect to the way to manage facts containing functions. An example will clarify.

Example 11. Consider $p(s(1))$. Writing a sort of “temporary rule”, only for meaning, we could see $p(S) :- s^*(1, S)$. We should “ground” somehow that S . Following a kind of “high-level” semantics, we should add a value in the table of p , and a value in the table of s^* . We could look-up the table for s^* : if a tuple of kind $j 1, _$ already there exists (say $j 1, @_1 _$) we add $j @_1 _$ to p and do nothing else (think about a previous fact like $q(s(1))$). If not, we create a new identifier (say $@_3$), put it in p , and in addition we add the tuple $j 1, @_3 _$ to s^* .

This choice will grant a single identifier for a single value of the function. It should work even in case of arity > 1 .

Implementation Note With respect to the choice about the parser's role, we can shortly describe the behaviour of the rewriting module. It can add the tuples for facts containing functions, while should simply rewrite the rules in which some functions appear to be processed lately by the grounding module.

Odd Stuffs Some considerations.

A function predicate born from a “function atom” appearing both in the head predicate and in the body, has to be considered as a full “body-predicate”: it can not produce new values.

Since the matching of a function predicate “coming from head” always succeed (as exposed above), then we could avoid to match it during the instantiation of the body (it seems useless); rather we match it adding the corresponding value only when we “commit” the grounding instantiation of the rule. A possibility is to put those kind of predicates at the very end of the body.

Once we've inserted the new values in the function tables, they are true. When the program is simplified, removing all true atom from the bodies after the grounding, the function atoms should be deleted. If we reach the aim to treat the function predicates almost as standard ones, we should not have to do a thing, because of the existing simplifying methods.

Processing nested functions is almost transparent, and can be done inside-to-outside or outside-to-inside. For instance, $p(s(f(1, a)), 2)$ can be first rewritten as $p(s(F), 2) :- f^*(F, 1, a)$. and then as $p(S, 2) :- f^*(F, 1, a), s^*(F, S)$.

About the problem of no-termination due to the nesting, we can give the user the possibility to guarantee the termination. This can be done by telling the system when it should stop the generation of function symbols. That is, the user can tell: please don't generate functions having nesting level greater than k . The nesting level is the maximum number of nested “subterms” which are in a term. This mean 0 for a constant, and so on: for instance 1 for $s(a,b)$, 3 for $f(s(t,w(a)), f(b,c))$, etc. The option will be something like $-maxNesting = k$. In order to be able to support this option, we need to store the nesting level of each term. At the beginning, we set it to 0 for the constants. Whenever we create a new function id (i.e., we insert a new element in a function table), we look at its arguments, we check what is the maximum nesting level “1” among them, and we set to “1+1” the nesting level of the new id . If $maxNesting$ has been set to k , and the new term has nesting $> k$, then the insertion fails. This means that even the set predicates coming from the flattening of the head terms may fail. If the option is not set, no checks are performed, so the termination is not guarantee and it is under the responsibility of the programmer.

This leads us to process the functions inside-to-outside. For instance $f(q(X, (w(Y), g(Y))))$ gives first g and w at level 1, then q at level 2, and at last f at level 3 (supposing that Y is bound to a constant).