Modelling the costs and benefits of Honeynets

Maximillian Dornseif[*]   Sascha A. May[†]

February 1, 2008

For many IT-security measures exact costs and benefits are not known. This makes it difficult to allocate resources optimally to different security measures. We present a model for costs and benefits of so called Honeynets. This can foster informed reasoning about the deployment of Honeynet technology.

1 Introduction

Honeynets are collections networked of computer systems which are intended to be attacked and broken into in an observed fashion, keeping track of any (mis-)use. Similar to other IT-security technologies there is a lot of gospel on the benefits of Honeynets, while there is little analysis on the exact gain offered by them and the associated cost. We are presenting a model helpful in understanding the economic aspects of Honeynet deployment.

In section 2 we present an overview of the technical aspects of Honeynets. Sections 3 and 4 collect benefits and respective costs of Honeynets. Following that section 5 contains our model of Honeynet deployment followed by section 6 where we summarize our findings.

2 Honeynet Technology

Honeynets are a term used to describe one or more computers destined as being penetrated and supportive technology designed to capture activity on the Honeynet and to decrease the risk imposed by the Honeynet to other systems [2].

The usual setup consists of the hosts to be attacked connected via a transparent firewall to the Internet. This firewall termed "Honeywall" is responsible for logging all network traffic entering and leaving the Honeynet. It also tries to suppress grave attacks from compromised hosts in the Honeynet by rate limiting outgoing traffic radically and by rewriting outgoing traffic on the fly to stop known exploits initiated from the Honeynet.

---

[*]Laboratory for Dependable Distributed Systems, RWTH Aachen University – dornseif@informatik.rwth-aachen.de

[†]sascha@may.de

The hosts to be attacked are instrumented in a way aimed at allowing the operators of the Honeynet to capture encrypted network traffic on the host where it exists in decrypted form while users of this hosts. This is either archived by using trojaned binaries like SSH which log data processed by them in disguise or by modifying the kernel to log certain system calls resulting in a crude keystroke logging facility. The captured data is covertly transferred to the Honeywall for logging.

3 Benefit of Honeynets

Deployment of Honeynets results in information gathered and possibly an increased security for the operator of the Honeynet.

3.1 Possible Information Gain on Attacks by Honeynets

Honeynets can gain information on the attacks against them. We assume that a Honeynet can basically gather two different qualities of Information: After starting his attack at t_a the attacker is unaware of the fact that he is attacking a Honeynet the data gathered shows the attacker's typical actions against the class of system the Honeynet is emulating. At a certain point in time the attacker realizes that he is confronted with a Honeynet. At this point labeled t_d the attacker's motivation shifts which should also result in a change of behavior. t_d can be even before t_a if an attacker is able to gather information about the Honeynet out of band and attacks with the knowledge that he is attacking a Honeynet. t_d also can be in infinitive future if the attacker isn't willing or able to find out that he is attacking an Honeynet.

It is safe to assume that after t_d the attacker will be more reluctant to act in a way which will allow the observer to gather further information. The attacker usually will completely stop the attack and vanish. But we also know of one instance where attackers using the Honeynet as an IRC proxy just ignored the fact that they where observed.

While attacking the attacker will try to escalate his privileges. He will increase his privileges in zero or more steps. The higher he was able to escalate his privileges the more likely he is to find out the true nature of the host he is attacking which results in t_d moving into future.

It is therefore safe to assume that sophisticated attackers t_d is relatively early. [1] A sophisticated attacker will be able to escalate his privileges relatively fast increasing his chances of detection. For attackers with full local privileges detecting a Honeynet is trivial.

Honeynets can not collect informations on all kind of attackers equally. Honeynets are be able to gather representative data on attackers which choose their targets more or less randomly like autonomous malware and very unsophisticated attackers do. Gathering data more on focused attackers can be only done for attackers actively choosing to attack the operators systems.

An attacker not penetrating systems in a random fashion must be tricked into attacking a Honeynet by making it look like a worthwhile target. It can be assumed that the more sophisticated the attacker is the less likely he will fouled by such deceptions.

So while Honeynets might be able to gather relatively much Information about unsophisticated attackers or autonomous malware like worms, with the same investment much less Information can be gathered about sophisticated attackers.

3.2 Possible increased security by using Honeynets as an decoy

It is claimed that Honeynets can increase the search space for finding valuable systems in a network and thus increasing security by luring attackers into spending effort attacking the Honeynets instead of the real thing. This claim has to be evaluated against different adversary scenarios.

Attackers attacking random hosts in your network have a bigger search space. But only extremely unsophisticated attackers like autonomous malware can be assumed to attack completely random hosts. Also these attackers can only be significantly slowed down when a significant percentage of a network are Honeynets which is unlikely.

More sophisticated attackers will choose their target based on their objectives and on a systems perceived value to complete this objectives. Simply by their existence Honeynets will slow down the attackers target selection process. To foul the attacker in attacking the Honeynet the Honeynet has to look more attractive than the target the attacker is aiming for or the "real" system has to be hidden in a way that the attacker will not be able to detect it.

3.3 Possible increased security by using aggressive Honeynets for redirection

There are also attempts to deploy honeypots as part of active network security. It is tried to reroute attackers from a production server to a Honeynet distracting the attacker and allowing further gathering of data [6].

The detection of the attack triggering the rerouting is a non trivial problem. Also the Honeynet must mirror very closely the production host to make rerouting seamless and less detectable. Due to this unsolved problems we will exclude aggressive Honeynets from further investigation.

4 Costs of Honeynets

Honeynets come with costs. Costs can be separated into costs for deploying and operating the Honeynet and costs due to increased risk imposed by the Honeynet to the operators network.

4.1 Cost of deploying

The initial costs of deploying a Honeynet at t_0 consists of the hardware, including computers, network devices and wiring, housing, personal cost for setup and cost of fitting the Honeynet into the policy framework of the organization. Since Honeynets are a relatively little understood and very new technology, personal costs for setup are likely to be exceptionally high. Also since Honeynets are something relatively new, explaining them to all stakeholders in the organization, evaluation of policy implications and weighing risks against benefits is likely to consume considerable resources.

4.2 Cost of operation

Operational costs consist of maintenance costs consist of fixed costs for housing, power, basic monitoring and software maintenance. Considerable maintenance effort has to go into maintain-

ing the Honeywall to ensure the security of itself and capability to minimize the risk of being used as a stepping stone to attack further systems.

Between t_a and t_d variable costs arise for IP-traffic, log space, forensics and active avoidance of being used as a stepping stone to attack further systems. At t_d variable costs for damage repair arise.

Since in the Internet there is climate of permanent attack, t_0 and t_a are likely to coincide. This results in permanent incurrence of the variable costs.

4.3 Cost of increased risk to own network

By deploying a Honeynet the complexity of a network is increased [3]. Also less-than state of the art protected systems are added to the network. This decreases network security and legal risks to the organization running a Honeynet.

4.3.1 Risk to your Security

The obvious effect of adding complexity and not state of the art protected systems is decreased security resulting in increased risk to the network attached to the Honeynet. Possible scenarios include the Honeynet attracting additional attackers, unexpected interaction of the Honeynet with other network components, use of the Honeynet as an attack platform against others despite countermeasures implemented in the Honeywall or the own network or use of the Honeynet for gathering intelligence about the attacked organization and it's methods.

4.3.2 Legal Liability Risk

The main liability risk of Honeynets consist of the risk that an attacker uses the Honeynet to attack systems of a third party and that this party seeks damages against the operator of the Honeynet. Legal issues raised by Honeynets have seen up to now no in-depth analysis by experts. While the whole "downstream liability" issue, that is liability of organizations whose IT-systems are penetrated and used to attack others, has seen no consensus or even satisfying discussion by the legal community, this discussion is often use to argue that there is no liability to operators of Honeynets [4]. Such argument misses the important fact that the "downstream liability" is about possible negligence in securing your systems against misuse by third parties, while with Honeynets you are actually willingly and knowingly facilitate your network to be misused by making it less secure.

This carries a risk of legal liability against the operators of a Honeynet not only in relation to possible victims of an attack but also possibly liability of management in relation to the shareholders of the organization operating the Honeynet. Besides civil liability there might also be criminal liability depending very much on the jurisdiction applied.

5 Modelling

Based on this observations we can build a microeconomic model of Honeynets:

In trying to build a model of Honeynet operation we assume that during the attack between t_a and t_d the attacker does a move every unit of time. He manages zero or more times to escalate his privileges. Regardless of that every unit of time the Honeynet operator earns the same value of information on the attacker.

5.1 Honeynet Operator

We model the Honeynet operator according to the following rules:

1. The operator of the Honeynet is not interested in observing attacks per se, but only on attacks specifically aimed at his systems, since information on less focused attackers can be bought on the marketplace. The attacker which is of interest to the operator will be called qualified attacker and probably has the profile of a professional spy. We assume that information on qualified attackers is not available in the marketplace.

2. We assume that only a an extremely small percentage of attacks are committed by qualified attackers.

3. We assume an attacker stops the attack after he discovers the nature of the Honeynet. So after t_d there will be no additional attacker activity.

4. At t_0 the Honeynet is deployed and generates considerable fixed startup costs.

5. Every unit of time the Honeynet generates constant costs for housing, energy routine maintenance and updates.

6. Between t_a and t_d costs for additional resources like bandwidth and logfile storage, monitoring, increased risk, restoring of the Honeypot and forensic analysis arise. While some of this costs are per unit of time during an attack, others occur only at the end of an attack.

7. By investing in the Honeynet the Operator can make the true nature of the System harder to detect thus moving t_d to the right.

8. For every unit of time the attack by an qualified attacker persists, the operator gains information of a certain constant value. The information gained while attacks by unqualified attackers is worthless.

9. On the Internet there is a climate of permanent aggression. This means that $t_0 = t_a$ and that the constant stream of attacks can be superimposed to build a constant attack pressure.

10. The Operator is not interested in prosecuting attackers [5]. Experience shows that companies are extremely reluctant to prosecute computer crime. Also the circumstances under which evidence is gathered in Honeynets suggests that using such evidence in court would be difficult.

We feel that this rules can realistically model most potential operators of a Honeynet. A prominent exception are institutions whose mission includes information security or criminology research. To these organizations qualitative and quantitative data not only on qualitative attackers but on any attack might be of value.

5.2 Attacker

We model the attacker according to the following rules:

1. The attacker has fixed costs like office space and Internet access.

2. The attacker has fixed costs per unit of time during the attack starting at t_a.

3. The possible duration of the attack is unlimited and can converge to infinity.

4. The attacker learns at t_d that the attacked host is a Honeynet.

5. The Honeynet can not be perverted to remove collected information on the attacker. A qualified attacker will not use it as an attack proxy to third party systems.

6. Since the attacker has minimal interest on fake data, to the attacker the Honeynet suddenly turns out as worthless. The Attack is being ended.

These rules we use for modeling the attacker are a profound restriction on the types of behavior in the real world by assuming that the Honeynet can not be perverted. Real world experience shows that a seasoned attacker might be able to pervert all components of a Honeynet and use them to his will.

Also the costs to the attacker are hard to estimate. One could speculate that an attacker will use the cheapest forms of attack first and resort to more costly attacks gradually when failing to reach it goals with the cheaper attacks. Than again the attacker might set priority to avoiding detection and use his best, most expensive tools attacks at first to do so. This could be modeled by including the specific risk of detection in the cost of an attack. We argue that most attacks are in a the same magnitude and that the attacker usually does not base its decision on cost of a specific attack but on his actual level of intelligence and perception which attack might be most successful. Therefore we assume a fixed costs per unit of time during the attack.

6 Composition

Building on those assumptions we can construct two curves describing the cost and the utility associated with a Honeynet:

1. The cost of a Honeynet is expressed by $c(t) = S + Mt$ where S are the startup cost for deploying the Honeynet and M are the maintenance costs per unit of time. Maintenance costs are bound to be greater than zero.

2. The utility of a Honeynet is expressed by $u(t) = Pt\frac{M}{I}$ whereas P is the value of information gained by a single attack and I is the factor by which higher investments in the maintenance influence the likelihood of being attacked.

The use of a Honeynet is profitable if the curve $u(t)$ supersedes $c(t)$. If we vary M while keeping S, P and I constant we can find the optimum investment in maintenance for organizations using Honeynets.

7 Conclusion

While the start-up cost S seems relatively easy to determine, P and I are problematic. P refers to the value of the Information gained by observing a qualified attacker. Since this information is not available on the marketplace it's technical value is extremely high. But we doubt that the actual utility of this information to the operator is that high in all cases. Sometimes actions by qualified attackers might only qualify as boring.

I expresses the relation between maintenance and frequency of qualified attacks. We are not aware of means to determine real-world values for I but we suspect that the general frequency of attacks is very low and can be only slightly increased by investing more.

References

[1] Maximillian Dornseif, Thorsten Holz, and Christian Klein. NoSEBrEaK - attacking honeynets. In *Procceedings of the 5th Annual IEEE Information Assurance Workshop*, 2004. forthcomming.

[2] Honeynet Project. *Know Your Enemy*. Addison-Wesley, 2001.

[3] Lance Spitzner. Honeypot farms, August 2003.

[4] Lance Spitzner. Honeypots: Are they illegal?, June 2003.

[5] The honeynet project: Frequently asked questions.

[6] Jack Whitsitt, Jason Larsen, and Albert Gonzalez. The bait and switch honeypot.