

Log Analysis Case Study Using LoGS

Dmitry Mogilevsky
National Center for Supercomputer Applications
dmogilev@ncsa.uiuc.edu

7th February 2008

Abstract

A very useful technique a network administrator can use to identify problematic network behavior is careful analysis of logs of incoming and outgoing network flows. The challenge one faces when attempting to undertake this course of action, though, is that large networks tend to generate an extremely large quantity of network traffic in a very short period of time, resulting in very large traffic logs which must be analyzed post-generation with an eye for contextual information which may reveal symptoms of problematic traffic¹. A better technique is to perform real-time log analysis using a real-time context-generating tool such as LoGS.

1 Introduction

One of the simplest and most comprehensive intrusion detection methods available to a network administrator is analysis of firewall or network manager (such as Cisco's NetFlow) logs[Ranum04]. Even logs containing very basic information, such as destination and source IP address-port pairs can under proper scrutiny reveal important information about attempted attacks on the system and internal problems. Firewall or another central point for network traffic is an ideal spot for collecting this information, since

- All network traffic passes through it.
- Most firewalls generate excellent activity logs[Cid04].

The usefulness of this technique is diminished by the fact that on large networks, any interesting messages will be interspersed among a large number of non-malicious traffic. Combined with the fact that large networks send and receive on millions of connections per second, doing intrusion detection becomes a problem of finding a needle in a hay stack. For example, consider the following fictional log entries:

¹Such as non-obvious port scans and traffic indicative of worm activity

```
Nov 5 14:03:33,*.*.*.10:3434,1.2.3.5:12346
Nov 5 14:15:13,*.*.*.10:3434,1.2.3.6:12346
Nov 5 14:28:32,*.*.*.10:3434,1.2.3.7:12346
Nov 5 14:40:11,*.*.*.10:3434,1.2.3.8:12346
```

This example shows a host on address *.*.*.10 scanning our network for open NetBus ports[ISS98]. However, since this is a very slow scan (individual connections are approximately 12 minutes apart), there may be several million non-informative log entries separating these interesting entries. Depending on the way the log is audited by the administrator, it is possible for this trend to be missed. Instead of analyzing large network logs in search of tell-tale signs of attempted (or successful) intrusion, the network administrator may opt for using a real-time log analysis tool to analyze log messages as they are generated and detect problems based on context created by previously seen messages.

2 The Problem

For this case study, anonymized NetFlow log generated by connections to and from the network of a large university was examined. The log is over 1 GB in size. Chronologically, it spans a little over one hour. It contains 13.6 million individual entries. Currently, these logs are collected, but not analyzed, despite containing a wealth of information. It is possible to write a rudimentary script which will process such logs post-creation and search for interesting information. In adopting this strategy a number of questions must be considered.

1. How do I subdivide the log into separate files?
2. How do I collect and store contextual information that will help me detect port scans?
3. Can I handle potential context overflow from one log to the next?
4. Do I process each log file after its created, or do I collect some number of log files and process them as a batch?

Appendix 1 shows a sample script to process such logs. While it is fairly concise, it does a poor job of addressing contextual issues. Specifically, it only stores state for the last connection, which makes it impossible for this script to detect any port scans where the connections are not immediately subsequent. While it is possible to increase the depth of stored context, doing so would necessitate creating a store-and-search infrastructure for context, which, when done improperly can greatly reduce the performance of the script. Additionally, this script does not address the problem of context overflowing from one log file to another. These questions can be bypassed altogether, on the other hand, by real-time scanning with LoGS.

3 LoGS

LoGS, currently under development by James Prewett at the Center for High Performance Computing at University of New Mexico is a highly customizable and extensible real-time log analysis engine written in Common Lisp. It is quite efficient, able to process as many as 72,000 messages per second [Pre05]. Because it uses Common Lisp in its rule definition, it can be programmed to create and store states and run complex scripts whenever a message is matched [Pre04].

The use of Common Lisp makes LoGS unique from other log analysis tools. Common Lisp was chosen for its flexibility, availability and ease of use. Lisp also has a fast regular expression engine[Wei03], which allows LoGS to achieve its high message processing speed. Also, since Common Lisp is used to configure the rulesets as well, LoGS is user-extendable.

LoGS consists of five components - Messages, Rules, Rulesets, Actions and Contexts. Rules associate Messages (input from the analyzed log) with Actions (Lisp or external scripts). Rulesets extend Rules by grouping them together into related sets. Contexts collect related messages together to be processed as a group. Because LoGS actions can create new rules, as well as update existing rules, LoGS is run-time configurable. [Pre04]

4 Analyzing Firewall logs with LoGS

The context-oriented design of LoGS fits perfectly for the goal of real-time firewall log analysis. Firewall logs differ from, for example, system logs in that you cannot implement artificial ignorance[Ranum04], as every message may potentially contain interesting information when taken in a context of other messages. Every incoming connection must be examined against existing contents, and either update them if necessary or create new ones.

4.1 Rules

To effectively detect suspicious behavior in connection logs, LoGS must be configured to match every incoming message, isolating three crucial pieces of information - local address, remote address, and local port. The local port is then checked against a list of vulnerable ports to identify a possible vulnerability scan. Next, the remote address, local address and local port are checked against the existing contexts for matches on remote address and local address (indicating a potential vertical port scan) or remote address and local port (indicating a potential horizontal port scan)[Lis]. If an existing context is found, the message is added to that context. Otherwise a new context is created for the new message.

4.2 Contexts

Each log entry is entered into a context, either a pre-existing one, or a new one, depending on whether a similar message has been previously seen. Each

context that exists has a timeout, which is incremented each time a new message is added to the context. When context timeout occurs, a context action is triggered and the context is removed from the system. By increasing the size of the timeout value, the administrator can detect slower port scans (at the price of performance).

4.3 Actions

When a context times out and certain conditions are met (for example, number of accumulated messages in the context. It should be at least two or more), an action will be triggered. Since LoGS allows to define actions with arbitrary Lisp programs and even external scripts, its possible for contexts to trigger very complex series of actions. At the very least, all the messages in the context should be written to a separate port scans file. A real-time alert should also be displayed to the console whenever there is an open context which has accumulated two or more messages.

4.4 Defining real-time log analyzer in LoGS

Using the mentioned elements of LoGS it is fairly easy to construct a small ruleset to do detailed real-time log analysis. The following code is a sample of the LoGS code used to accomplish such a task, written in Common Lisp.

```
;Firewall analysis ruleset to spot vertical
;and horizontal port scan
(make-instance
  'rule
  :match
  (lambda (message)
    (multiple-value-bind (matches sub-matches)
      (cl-ppcre::scan-to-strings
        "([0-9]+),([0-9]+),([0-9]+:[0-9]+:[0-9]+),
        (TCP|UDP|ICMP),([0-9]+.[0-9]+.[0-9]+.[0-9]+):([0-9]+|--),
        ([0-9]+.[0-9]+.[0-9]+.[0-9]+):([0-9]+|--),
        [0-9]+,([0-9]+)"
        (message message)))
      (when matches
        (values
          t
          '((sub-matches ,sub-matches))))))
  :actions
  (
    (make-instance 'rule
      :match
      (lambda (message)
        (multiple-value-bind (matches sub-matches)
```

```

        (cl-ppcre::scan-to-strings
          "([0-9]+),([0-9]+),
           ([0-9]+:[0-9]+:[0-9]+),
           (TCP|UDP|ICMP),
           (aref sub-matches 4):([0-9]+|--),
           (aref sub-matches 6):([0-9]+|--),
           [0-9]+,([0-9]+)"
          (message message))
      (when matches
        (values
          T
          '((sourceip ,(aref sub-matches 4))
            (sourceport,(aref sub-matches 5))
            (destip, (aref sub-matches 6))
            (destport, (aref sub-matches 7))
            (time, (aref sub-matches 2))
          )))))
:actions
(list
  (lambda (message)
    (declare (ignore message))
    (ensure-context
      :name (format () "vertical scan from ~A" sourceip)
      :timeout (+ get_universal_time timeout_value)
      :actions
      (list
        (lambda (message)
          (add-to-context
            (format () "Vertical scan: ~A:~A to ~A:~A at ~A"
              sourceip sourceport destip destport time)
            message))))))
)
(make-instance 'rule
  :match
  (lambda (message)
    (multiple-value-bind (matches sub-matches)
      (cl-ppcre::scan-to-strings
        "([0-9]+),([0-9]+),
         ([0-9]+:[0-9]+:[0-9]+),
         (TCP|UDP|ICMP),
         (aref sub-matches 4):([0-9]+|--),
         ([0-9]+.[0-9]+.[0-9]+.[0-9]+):
         (aref sub-matches 7),
         [0-9]+,([0-9]+)"
        (message message))
      (when matches

```

```

(values
  T
  '((sourceip ,(aref sub-matches 4))
    (sourceport,(aref sub-matches 5))
    (destip, (aref sub-matches 6))
    (destport, (aref sub-matches 7))
    (time, (aref sub-matches 2))
  ))))
:actions
(list
  (lambda (message)
    (declare (ignore message))
    (ensure-context
      :name (format () "horizontal scan from ~A" sourceip)
      :timeout (+ get_universal_time timeout_value)
      :actions
      (list
        (lambda (context)
          (if (>= (get_universal_time) (timeout context))
              (report_context)
              )
          )
        (lambda (message)
          (add-to-context
            (format () "Horizontal scan: ~A:~A to ~A:~A at ~A"
              sourceip sourceport destip destport time)
            message))))
    )
  )
)

```

5 Conclusions

Real time firewall log analysis in LoGS offers a flexible and extensible alternative to batch offline analysis. Its capability for contextual message parsing is ideally suited for the task of detecting port scans, as these cannot be detected from any single connection and can only be inferred by looking at the context of previous connections. LoGS provides the infrastructure for contextual data collection, and the capability to trigger arbitrarily sophisticated response. This makes log analysis a very powerful tool for successful (and potentially real-time) intrusion detection.

6 Appendix 1 - Sample Perl code

```

#!/usr/bin/perl
print "Input? ";

```

```

$infile = <STDIN>;
open(INPUT,$infile)||die "Could not open $infile\n";
print "Outgoing connections? ";
$outgoingfile = <STDIN>;
print "Incoming connections? ";
$incomingfile = <STDIN>;
print "File date (yyyy/mm/dd)? ";
$date = <STDIN>;
print "Vulnerabilities file? ";
$vulfile = <STDIN>;
chop($date);
$portscans = 0;
$ongoingps = 0;
open(VULNERABILITIES, "$vulfile")
||die "Could not open $vulfile\n";
open(OUTGOING,">$outgoingfile")
||die "Could not open $outgoingfile\n";
open(INCOMING,">$incomingfile")
||die "Could not open $incomingfile\n";
open(PORTSCANS,">portscans_$infile")
||die "Could not create portscans file\n";
open(VULSCANS,">vulscans_$infile")
||die "Could not create vulnerabilities scan file\n";
while(<VULNERABILITIES>) {
    m/^[a-zA-Z0-9 ]+,[0-9]+,(in|out)$/;
    if ($3 == 'in'){
        $invul{$2} = $1 ;
    }
    else{
        $outvul{$2} = $1 ;
    }
}
$counter = 0;
$sourceip = "";
$sourceport = "";
$destip = "";
$destport = "";
$time = "";
while(<INPUT>) {
/([0-9]+),([0-9]+),([0-9]+:[0-9]+:[0-9]+),
(TCP|UDP|ICMP),([0-9]+.[0-9]+.[0-9]+.[0-9]+)
:([0-9]+|--),([0-9]+.[0-9]+.[0-9]+.[0-9]+)
:([0-9]+|--),[0-9]+,([0-9]+)/;
$lastsourceip = $sourceip;
$lastsourceport = $sourceport;
$lastdestip = $destip;

```

```

$lastdestport = $destport ;
$lasttime = $time ;
$time = $3;
$protocol = $4;
$sourceip = $5;
$sourceport = $6;
$destip = $7;
$destport = $8;
$packets = $9 ;
if ($destport == "--")
    {$destport = "0";}
if ($sourceport == "--")
    {$sourceport = "0";}
if ($sourceip =~ /10.[0-9]+.[0-9]+.[0-9]+/){
    print OUTGOING "FWOUT,$date,$time -5:00
        GMT,$sourceip:$sourceport,$destip:
        $destport,$protocol\n";
    if (exists($outvul{$sourceport})) {
        print VULSCANS "Potential Vulnerability:
            $outvul{$sourceport}.\n $time:
            $sourceip:$sourceport -> $destip:
            $destport\n" ;
    }
}
else{
    print INCOMING "FWIN,$date,$time -5:00
        GMT,$sourceip:$sourceport,$destip:
        $destport,$protocol\n";
    if ($sourceip == $lastsourceip &&
        (($destip == $lastdestip &&
            $destport != $lastdestport)||
            ($destport == $lastdestport &&
            $destip != $lastdestip))) {
        if (exists($invul{$destport})) {
            print VULSCANS "Potential Vulnerability:
                $invul{$destport}.\n $time:
                $sourceip:$sourceport -> $destip:
                $destport\n" ;
        }
    }
    if ($ongoingps == 0){
        $ongoingps = 1;
        $portscans = $portscans + 1;
        if ($destip == $lastdestip){
            print PORTSCANS "Potential vertical portscan from
                $lastsourceip at $lasttime\n";
        }
    }
}

```



```

else {
    print PORTSCANS "Potential horizontal portscan from
        $lastsourceip at $lasttime\n";
}
print PORTSCANS "$lastsourceip:$lastsourceport ->
    $lastdestip:$lastdestport\n";
print PORTSCANS "$sourceip:$sourceport ->
    $destip:$destport\n";
}
else {
    print PORTSCANS "$sourceip:$sourceport ->
        $destip:$destport\n";
}
}
else {
    $ongoingps = 0;
}
}
$counter = $counter + 1 ;
if(($counter % 10000) == 0)
{
    print "$counter\n";
}
}
print "$portscans Portscans detected
    and written to portscans_${infile}\n";
close (INPUT);
close (OUTGOING);
close (INCOMING);
close (VULSCANS);
close (PORTSCANS);
close (VULNERABILITIES);

```

References

- [Pre04] James Prewett. "Listening to Your Cluster With LoGS". LCI's 5th International Conference, Linux Clusters: The HPC Revolution 2004. Available on WWW: http://www.linuxclustersinstitute.org/Linux-HPC-Revolution/Archive/PDF04/05-Prewett_J.pdf
- [Pre04-2] James Prewett. "LoGS". January 2004. Available on WWW: <http://www.hpc.unm.edu/download/LoGS.pdf>
- [Pre05] James Prewett. "Latest real-world benchmark". Email communication. January 22, 2005.

- [Wei03] Edi Weitz, “CL-PPCRE Benchmarking”. December 2003. Available on WWW: <http://www.weitz.de/cl-ppcre/#bench>.
- [Ranum04] Marcus J. Ranum. “Untapped Riches”. June 2004. Information Security Magazine. Available on WWW: http://infosecuritymag.techtarget.com/ss/0,295796,sid6_iss407_art816,00.html
- [Cid04] Daniel B. Cid. “Detecting Intrusions with your Firewall Log and OsHids”. Relevant Technologies, January 22nd 2004. Available on WWW: http://www.relevanttechnologies.com/detint_220104.asp
- [ISS98] “ISS Vulnerability Alert”. Internet Security Systems. September 10th 1998. Available on WWW: <http://xforce.iss.net/xforce/alerts/id/advised>
- [Lis] Kevin Liston. “Intrusion Detection Faq”. SANS. Available on WWW: http://www.sans.org/resources/idfaq/anomaly_detection.php